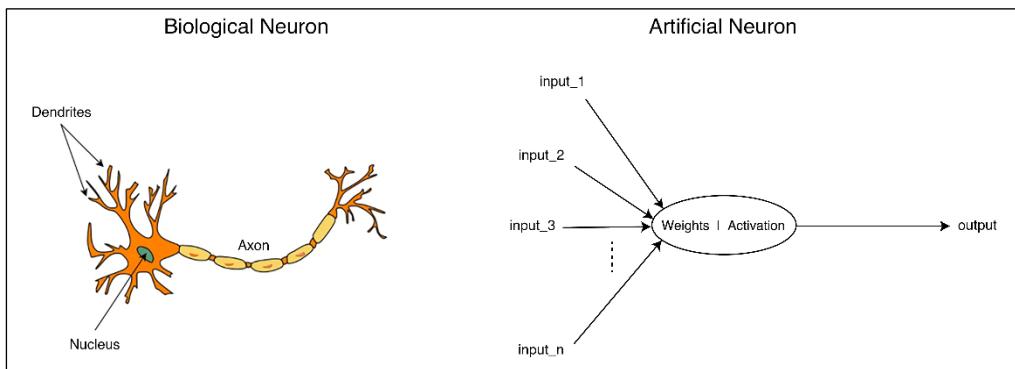
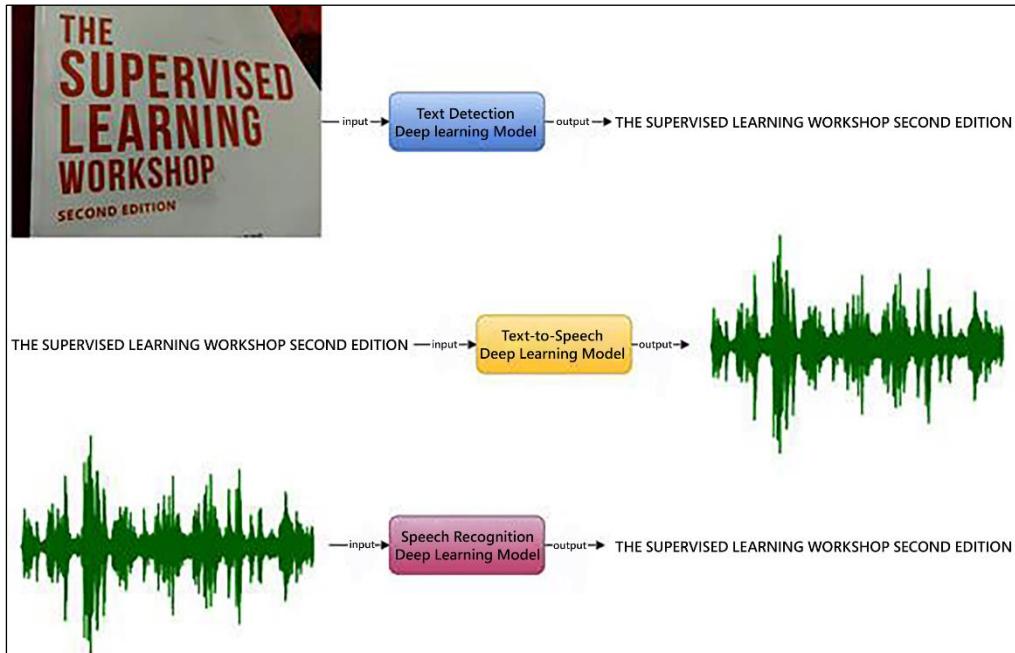
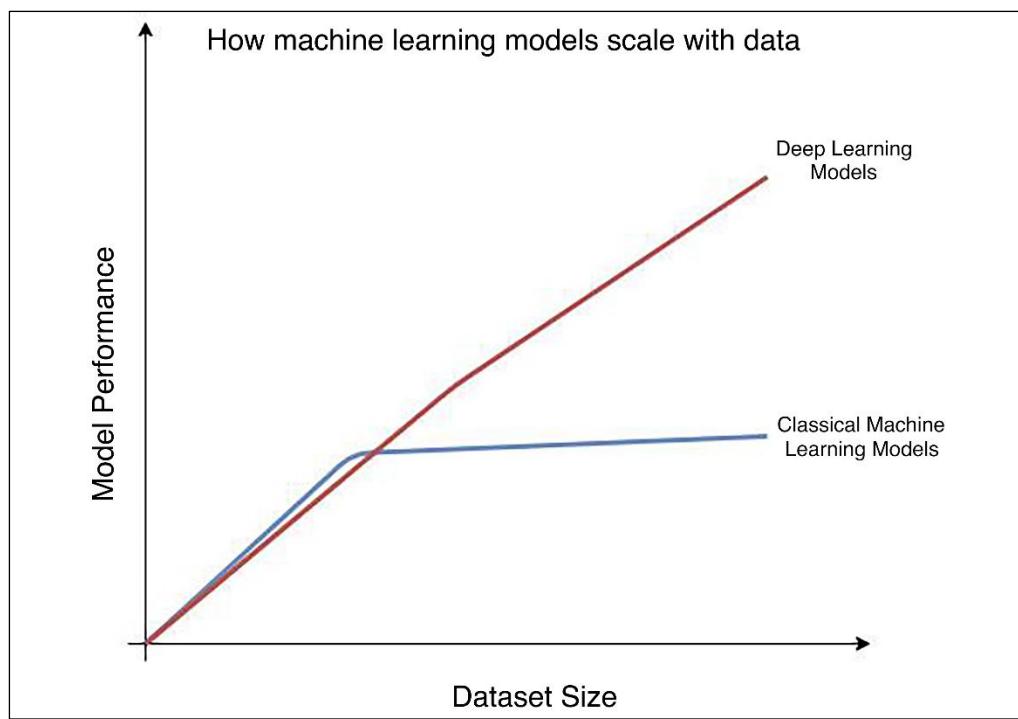
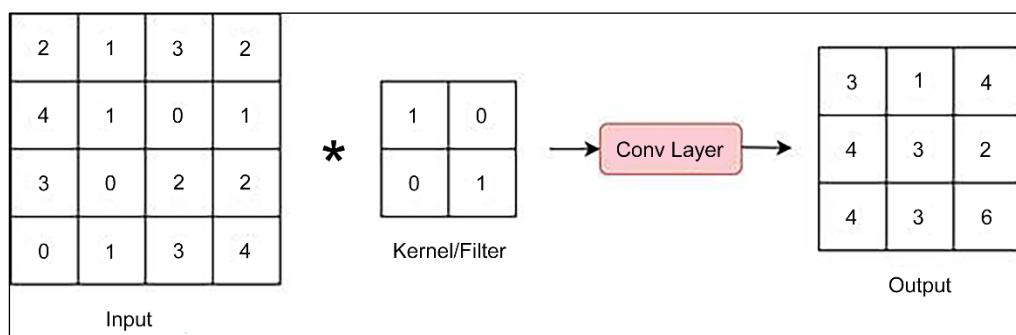
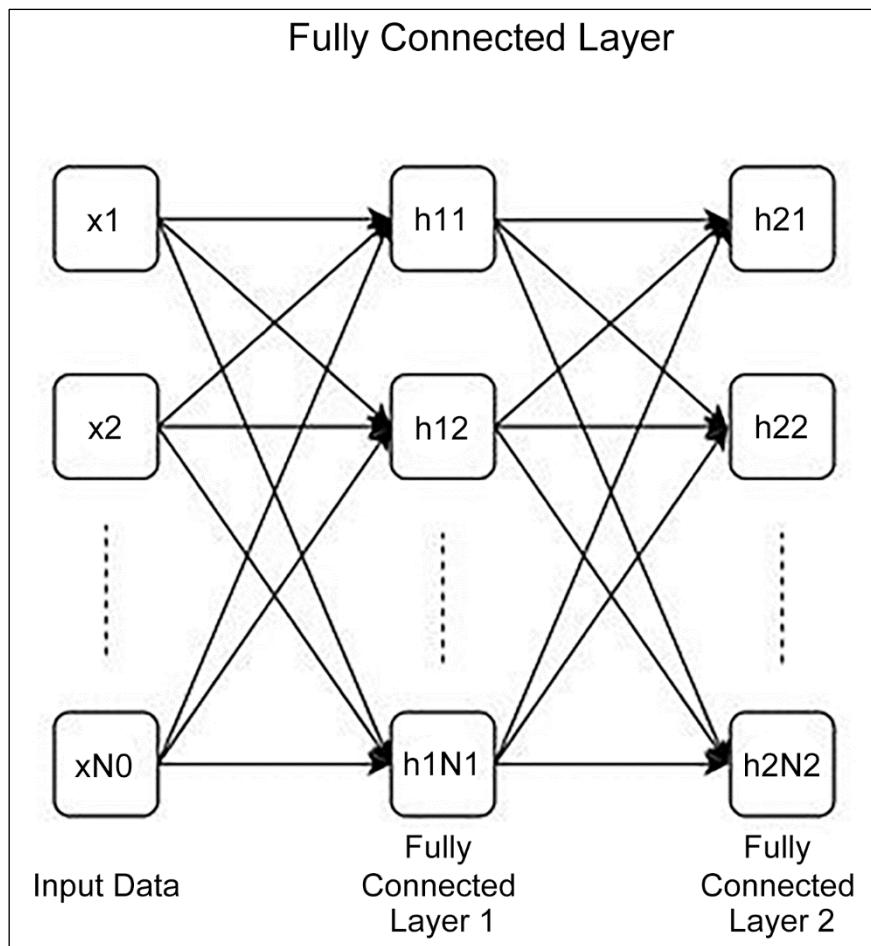
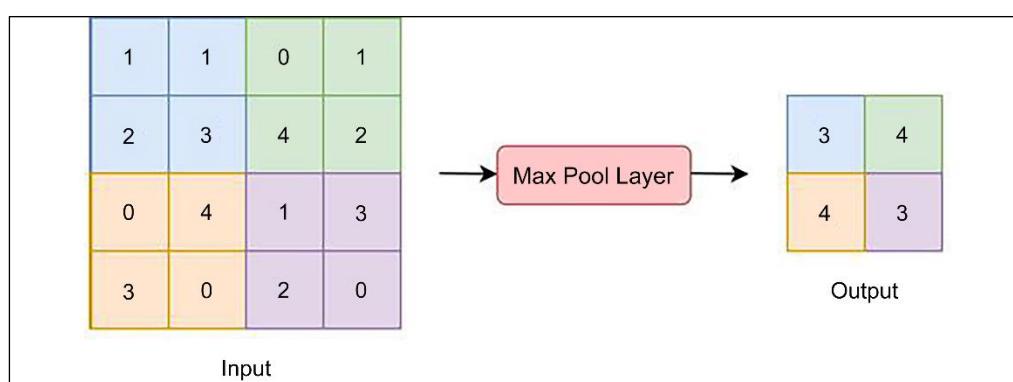
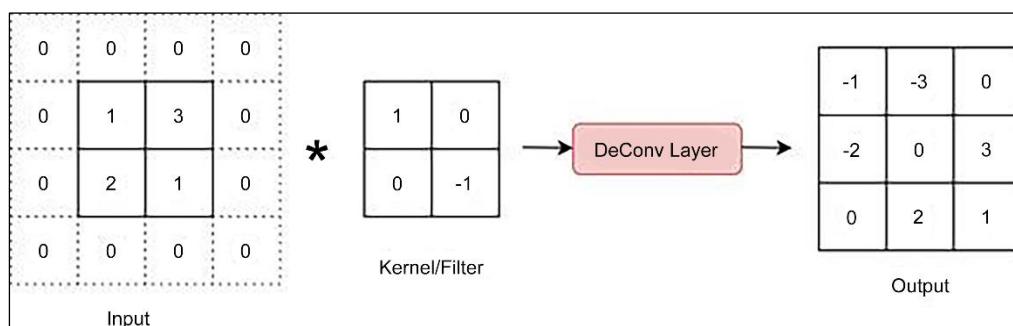
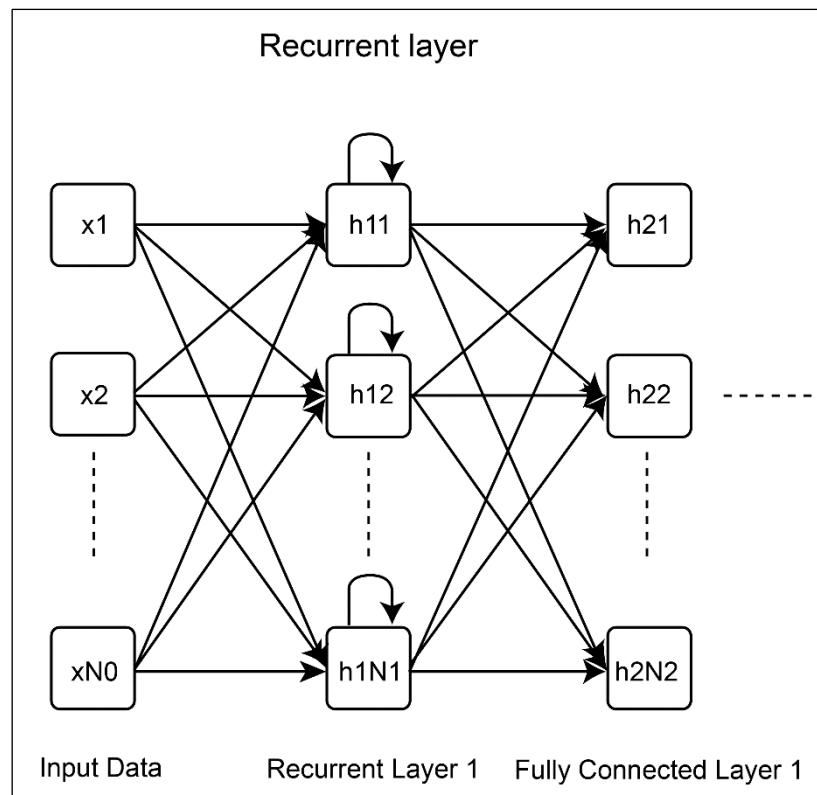


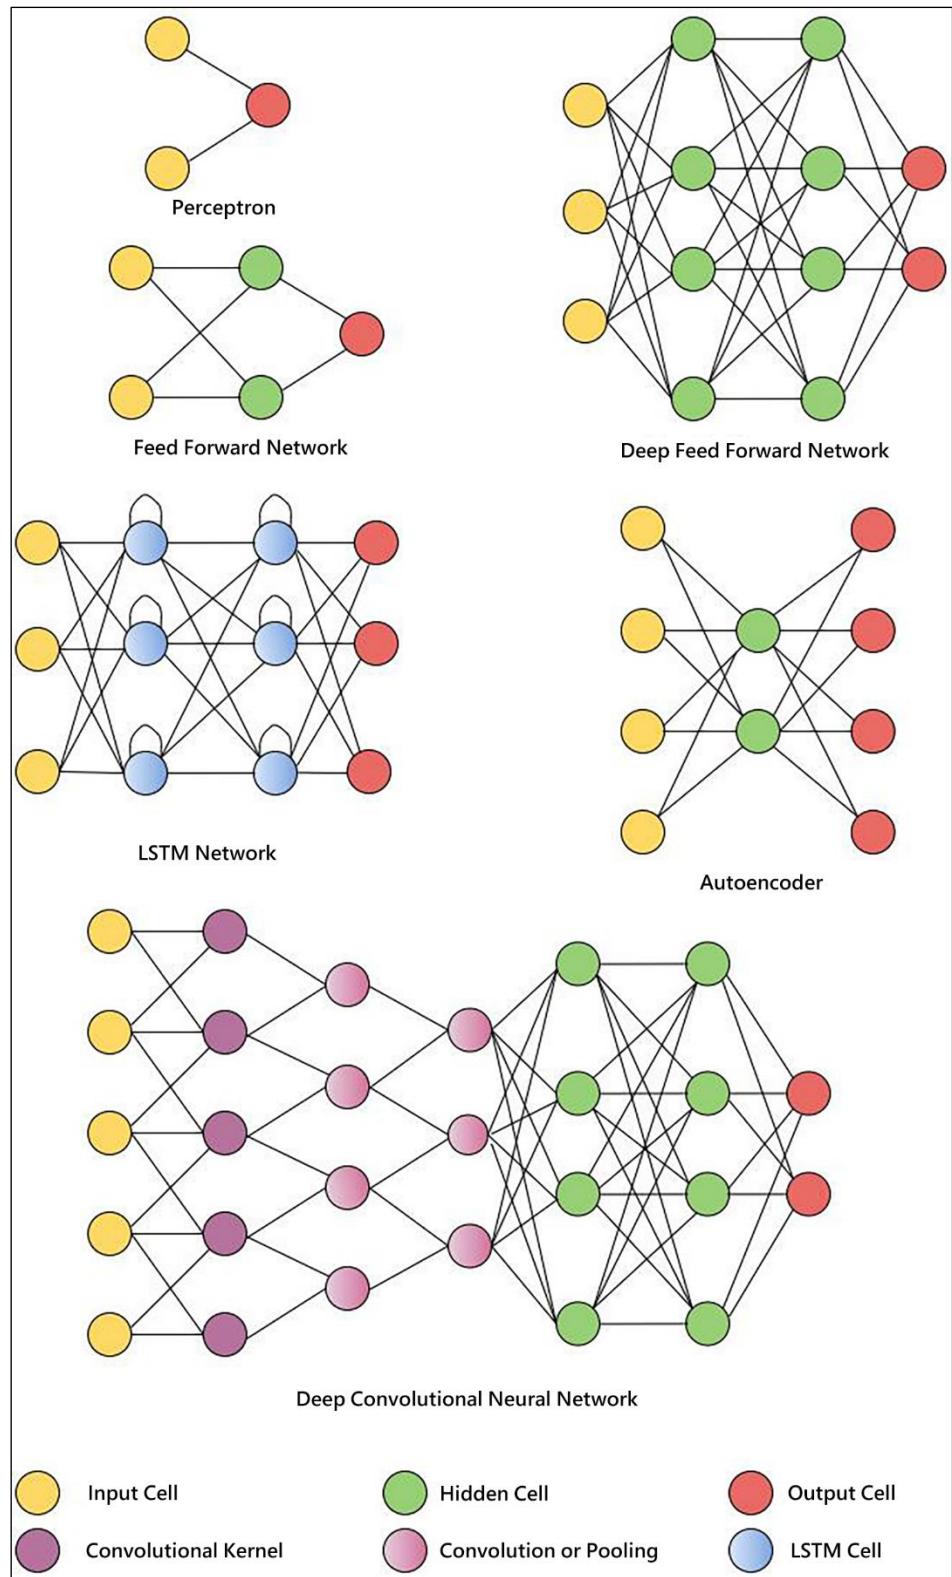
Chapter 1: Overview of Deep Learning Using PyTorch

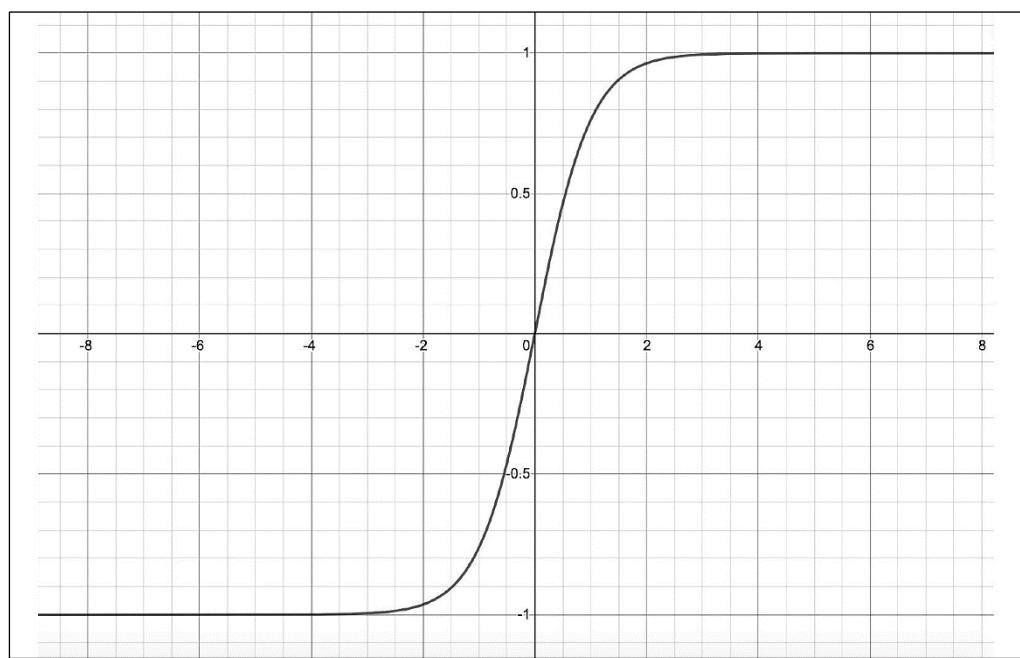
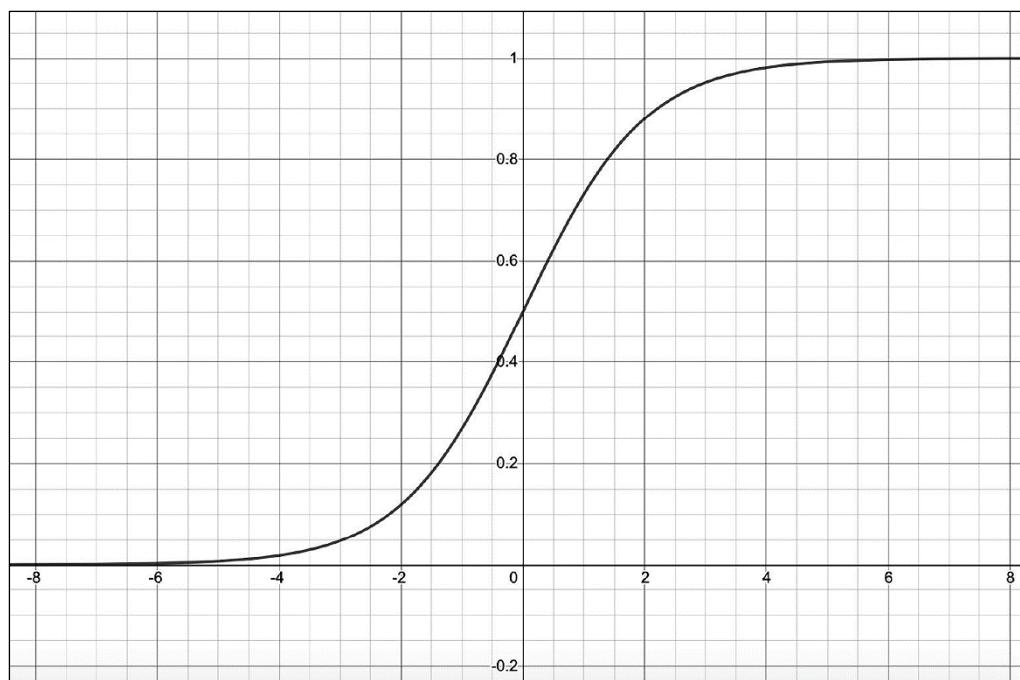


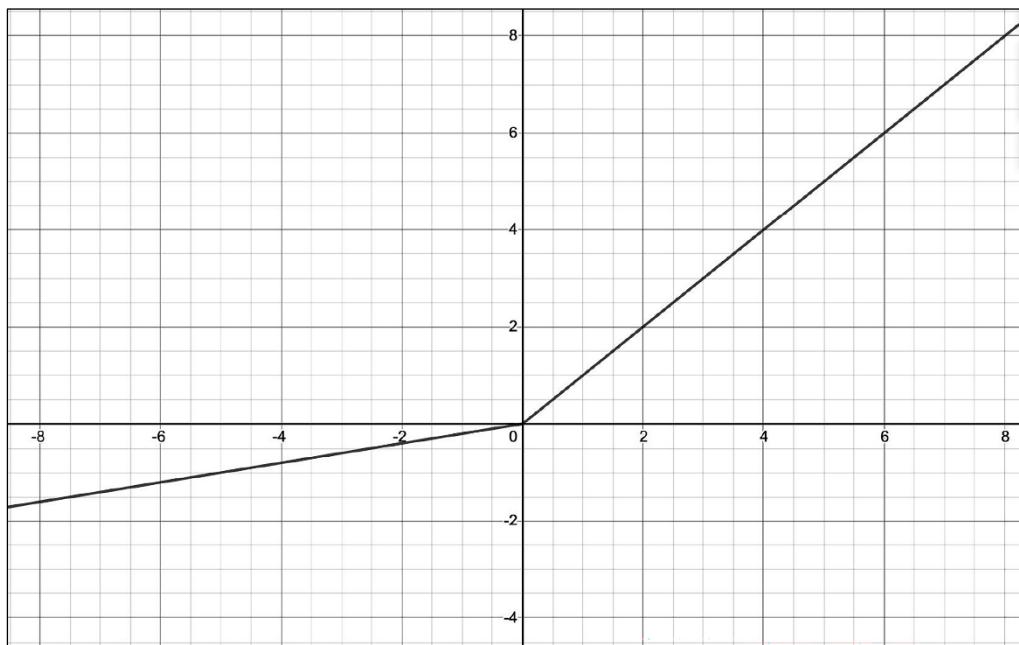
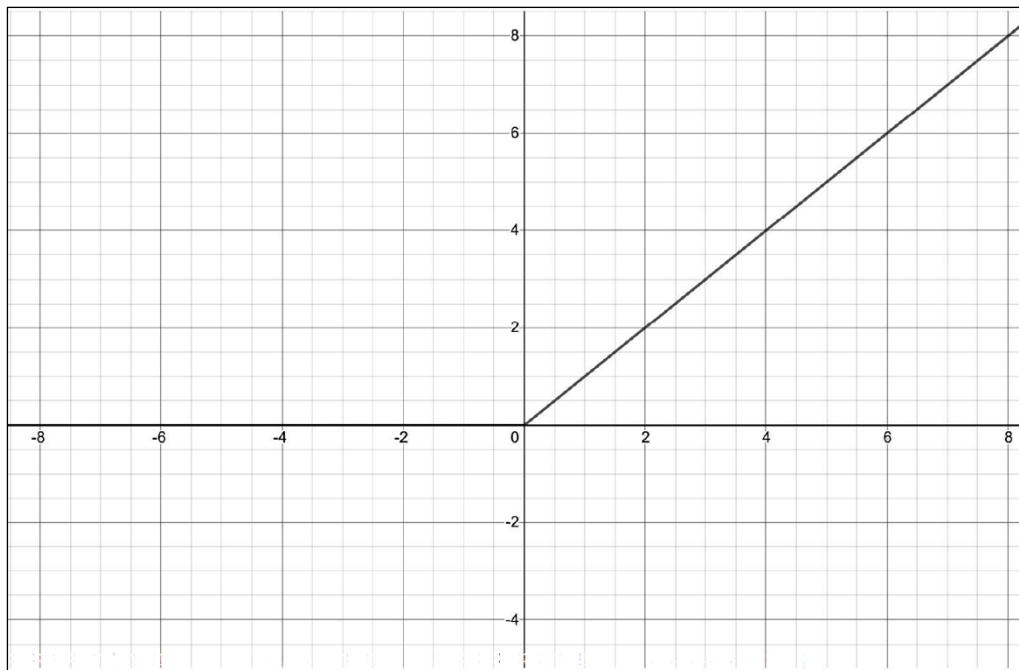


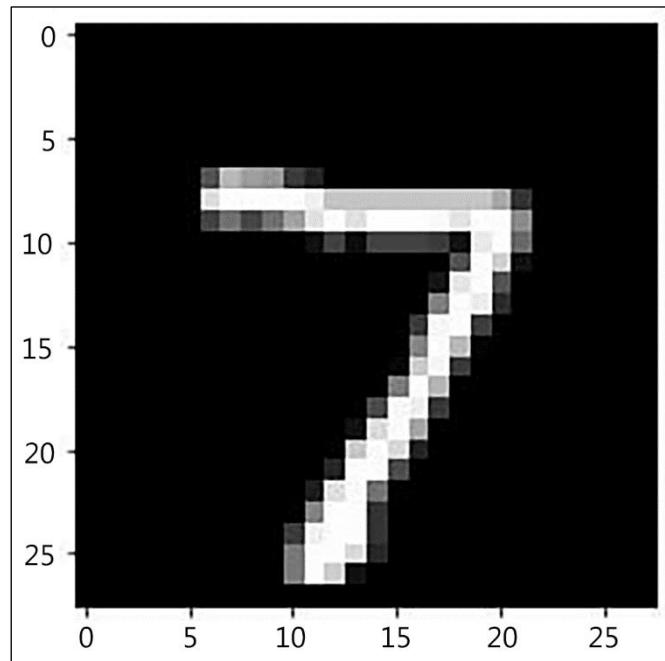
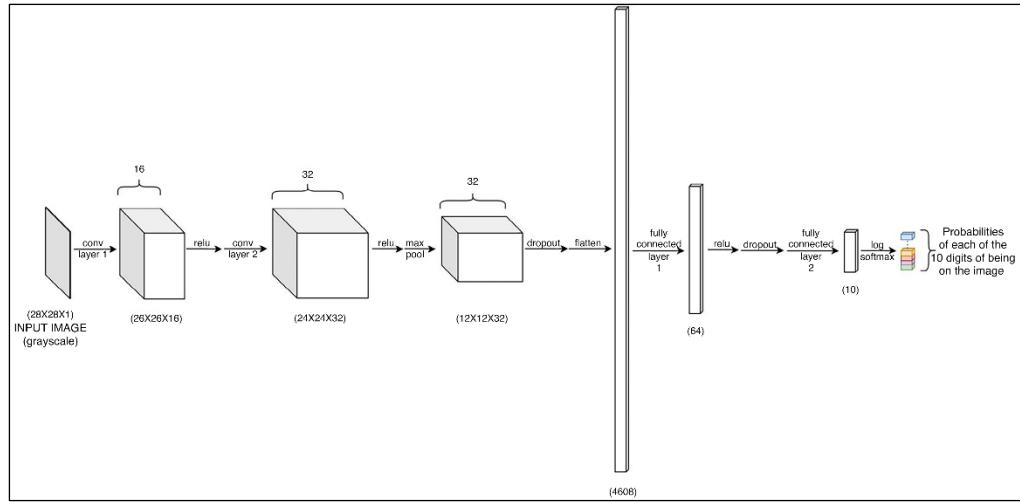




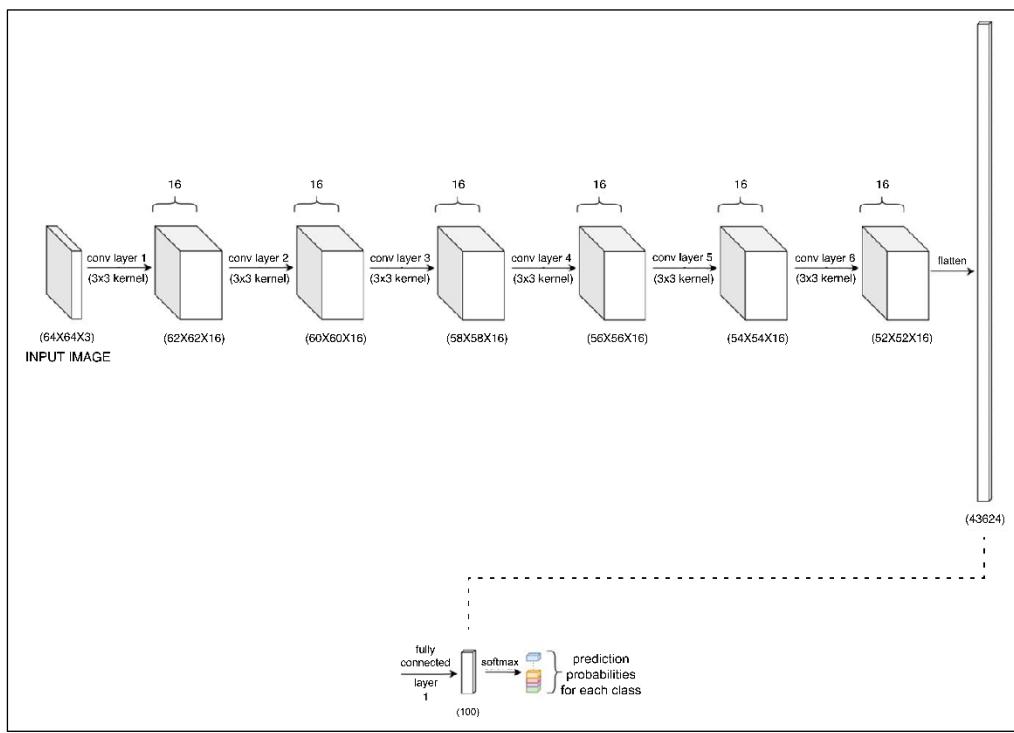
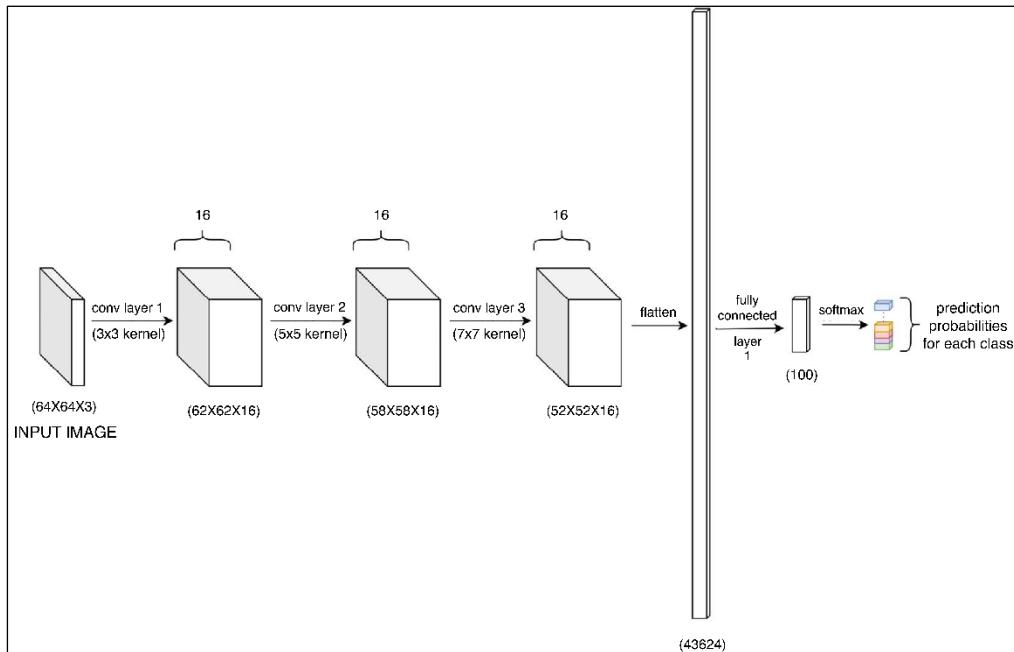


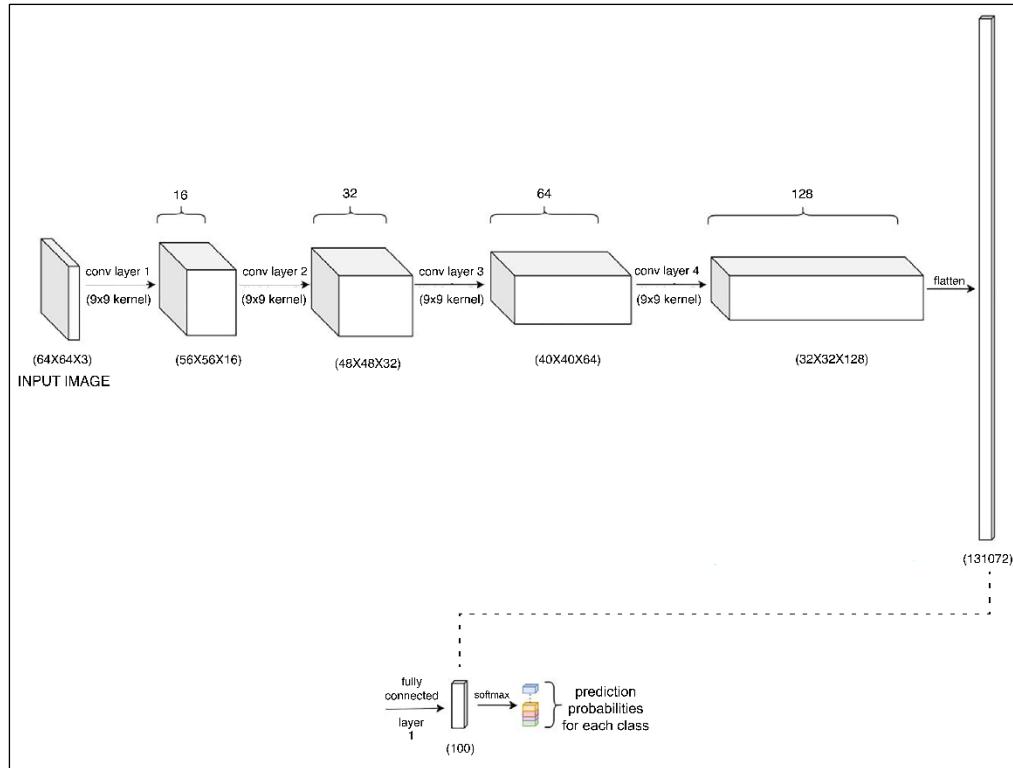


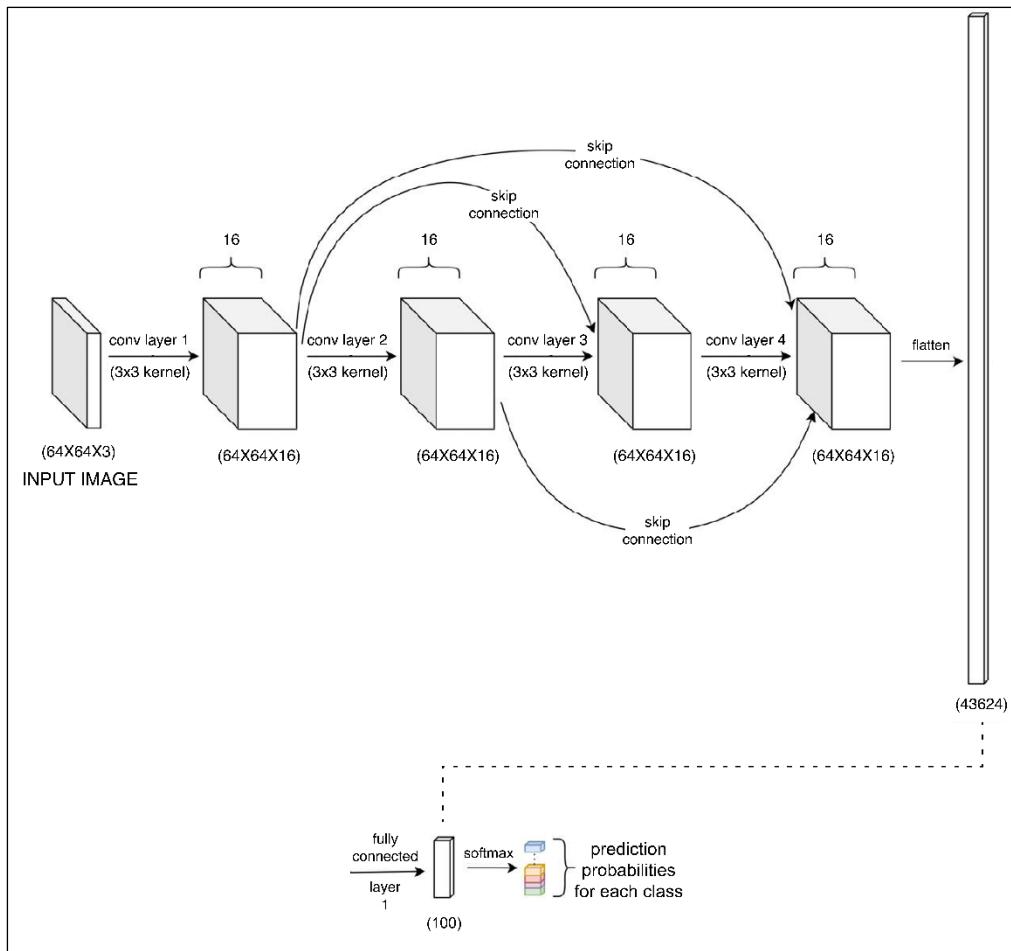


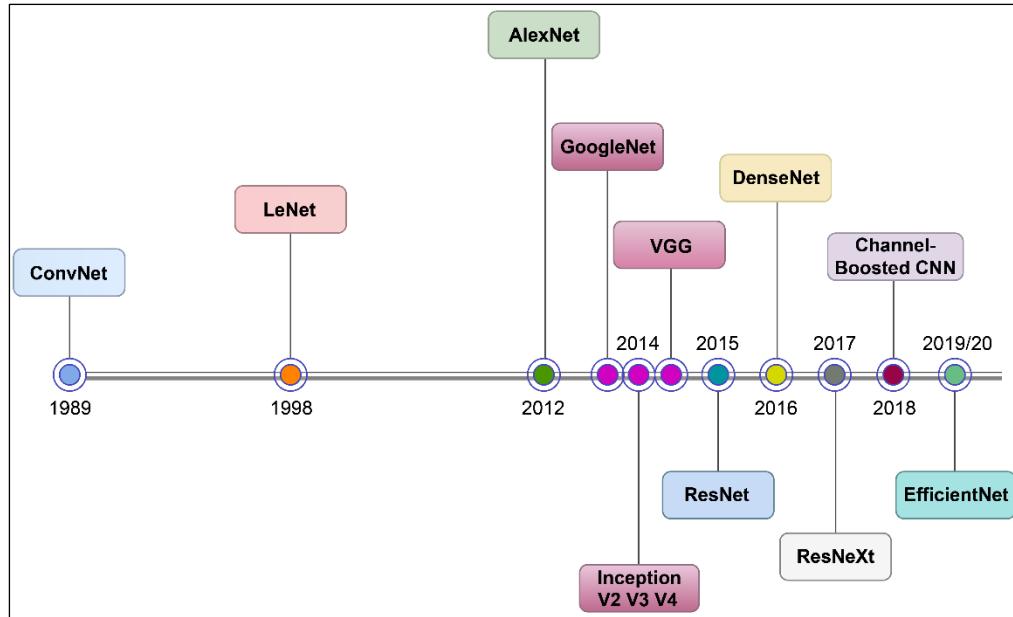


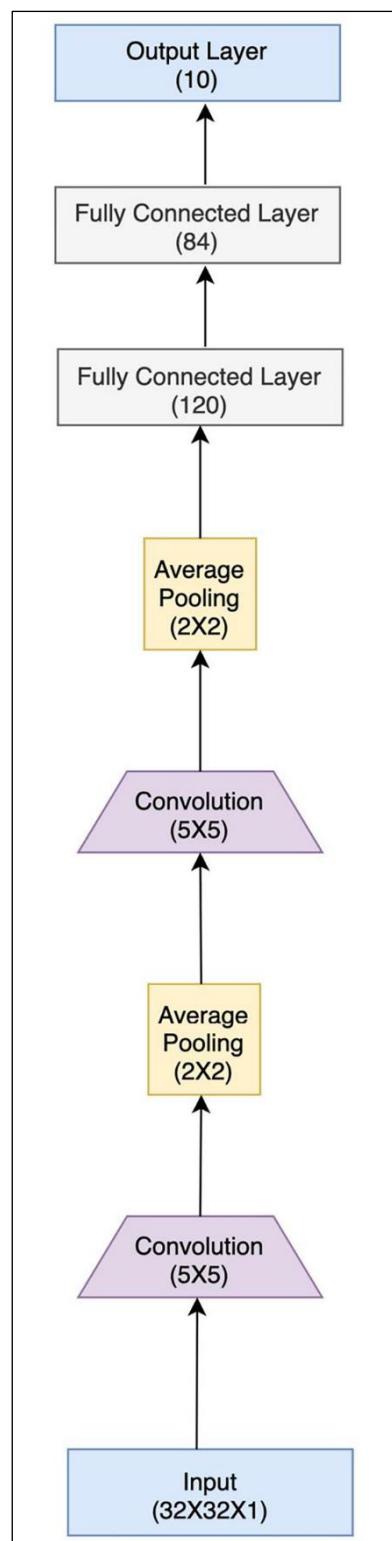
Chapter 2: Deep CNN Architectures

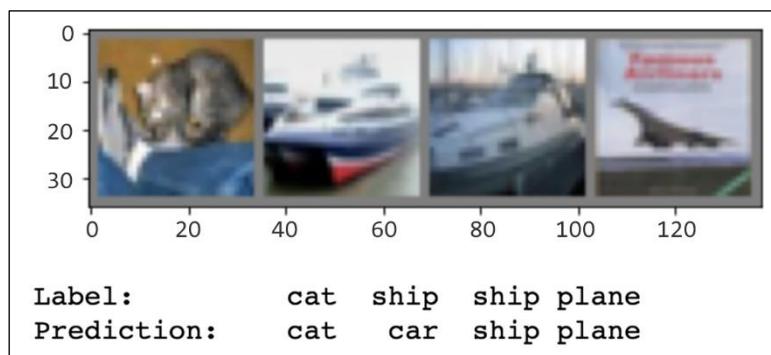
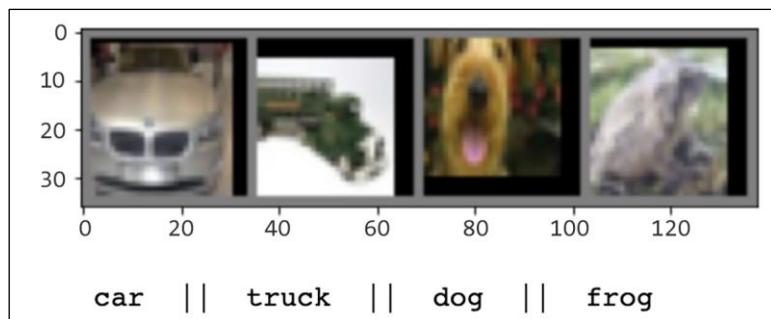


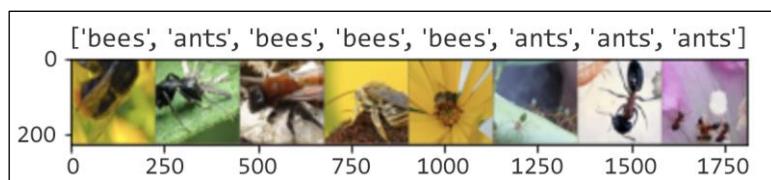
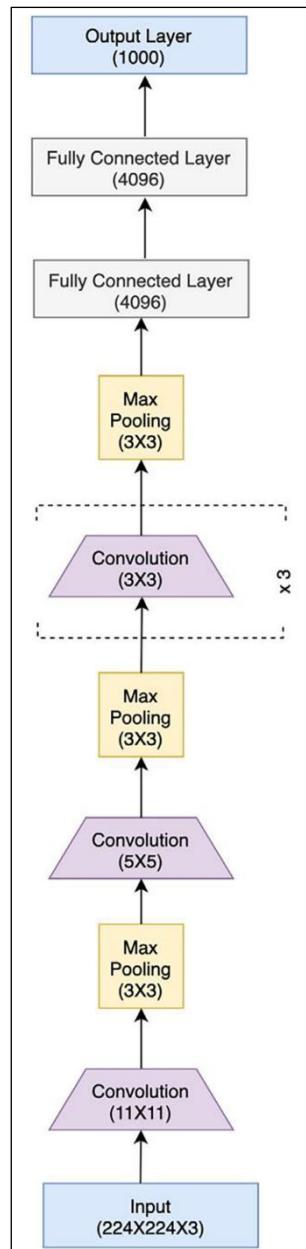












pred: bees || target: bees



pred: ants || target: ants

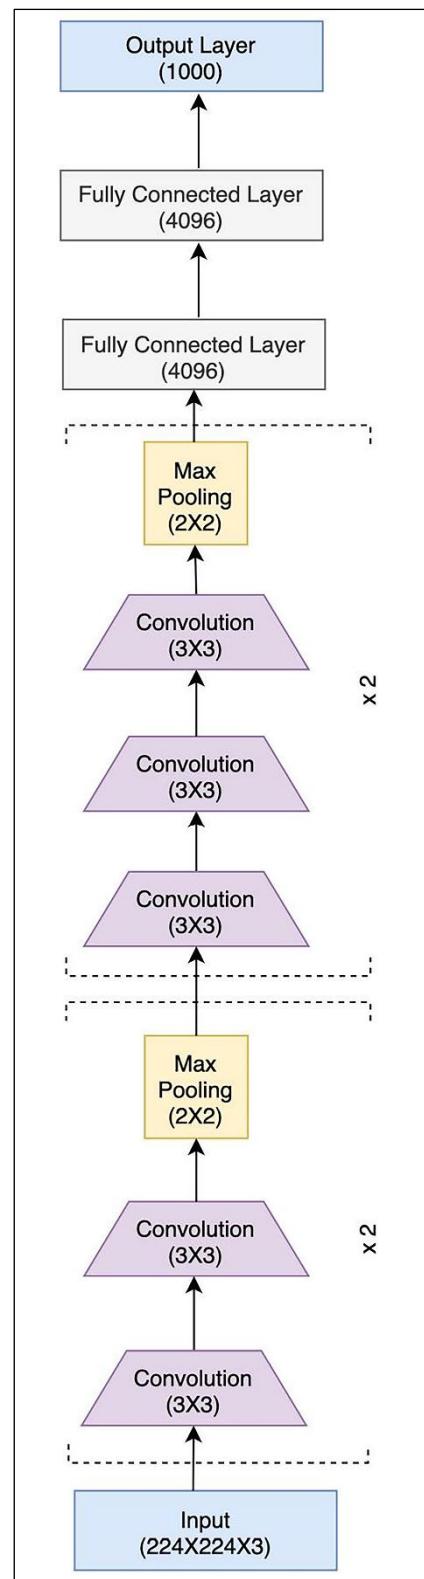


pred: ants || target: ants



pred: bees || target: bees





pred: bee



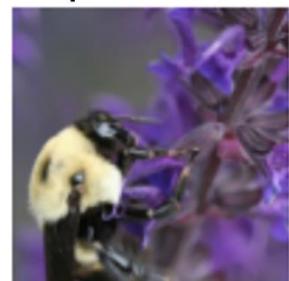
pred: ant, emmet, pismire

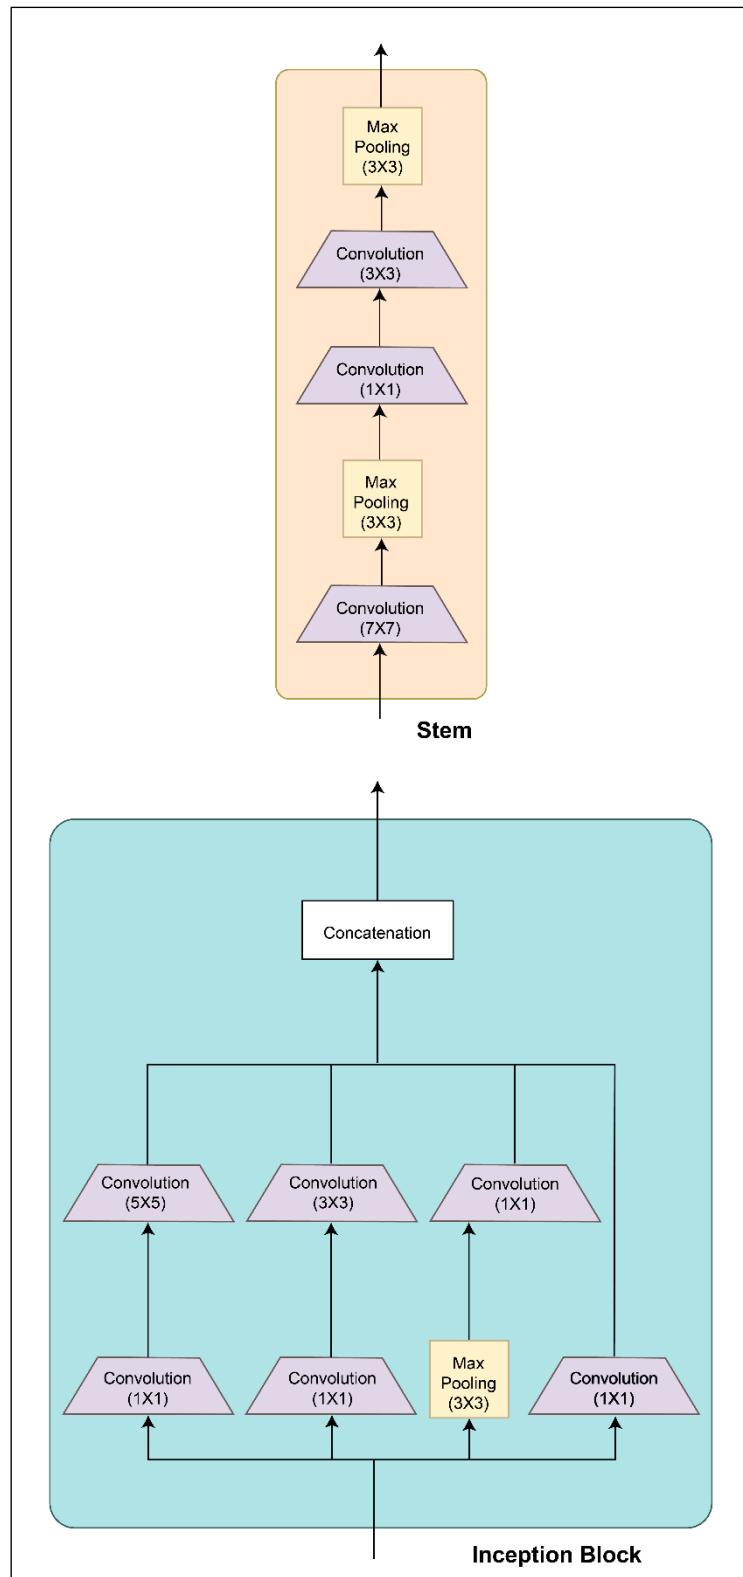


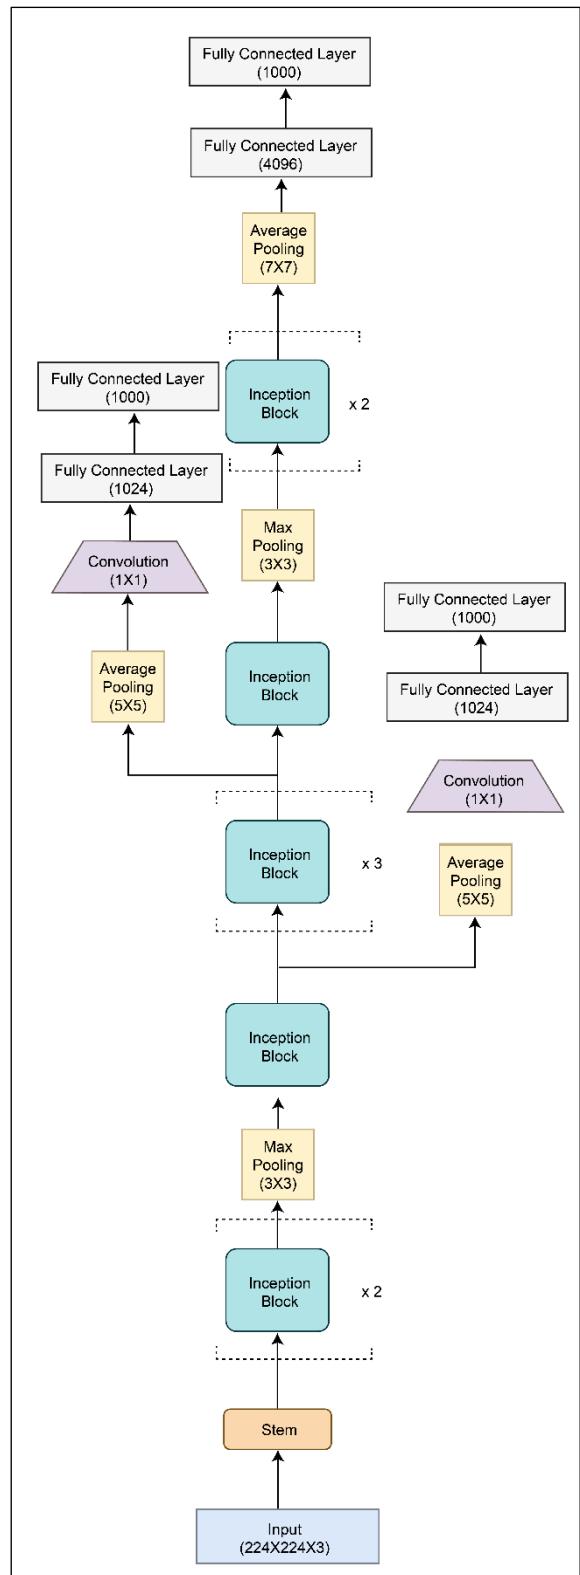
pred: ant, emmet, pismire

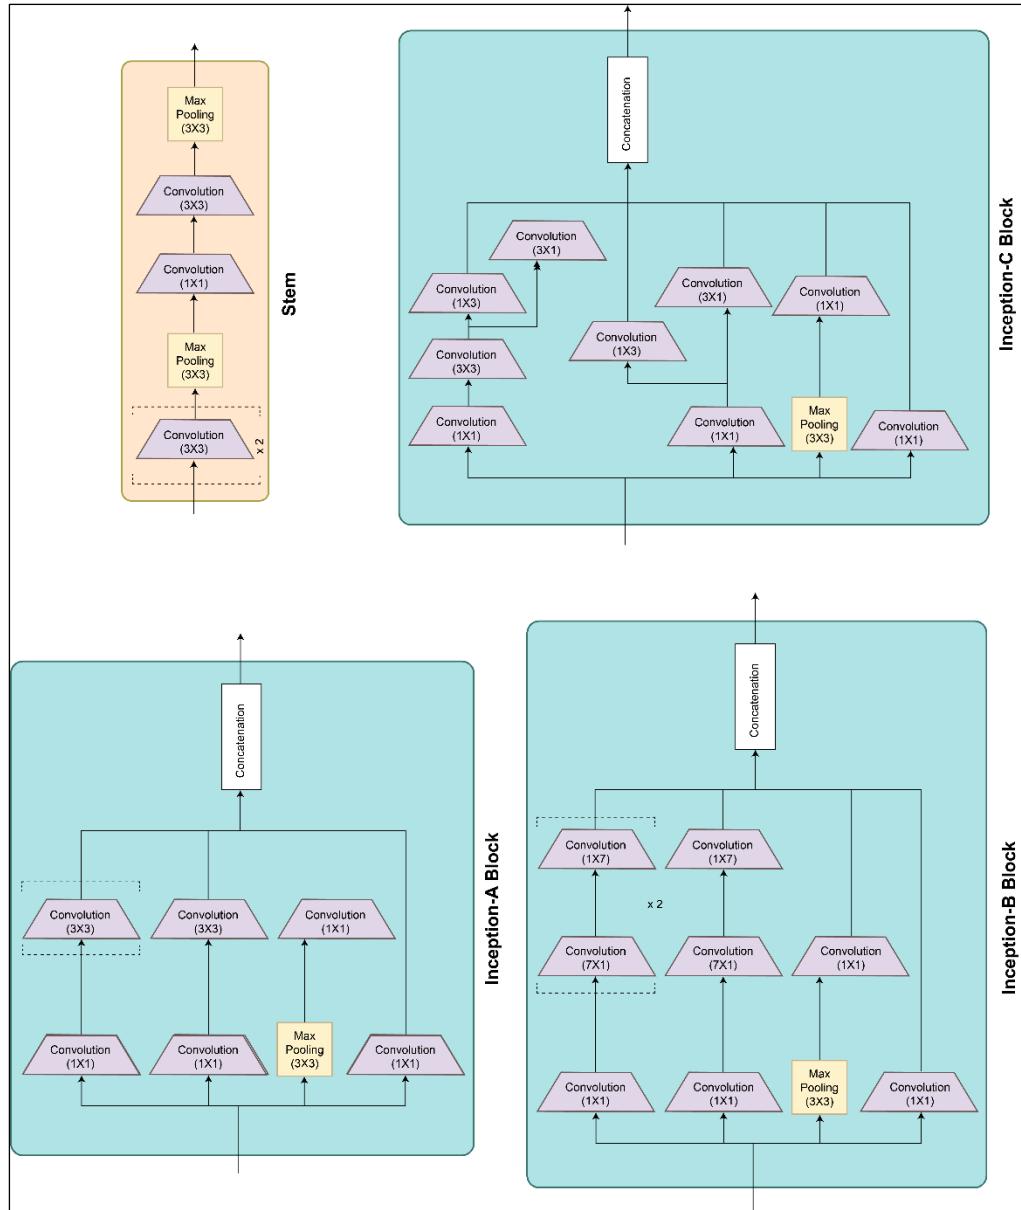


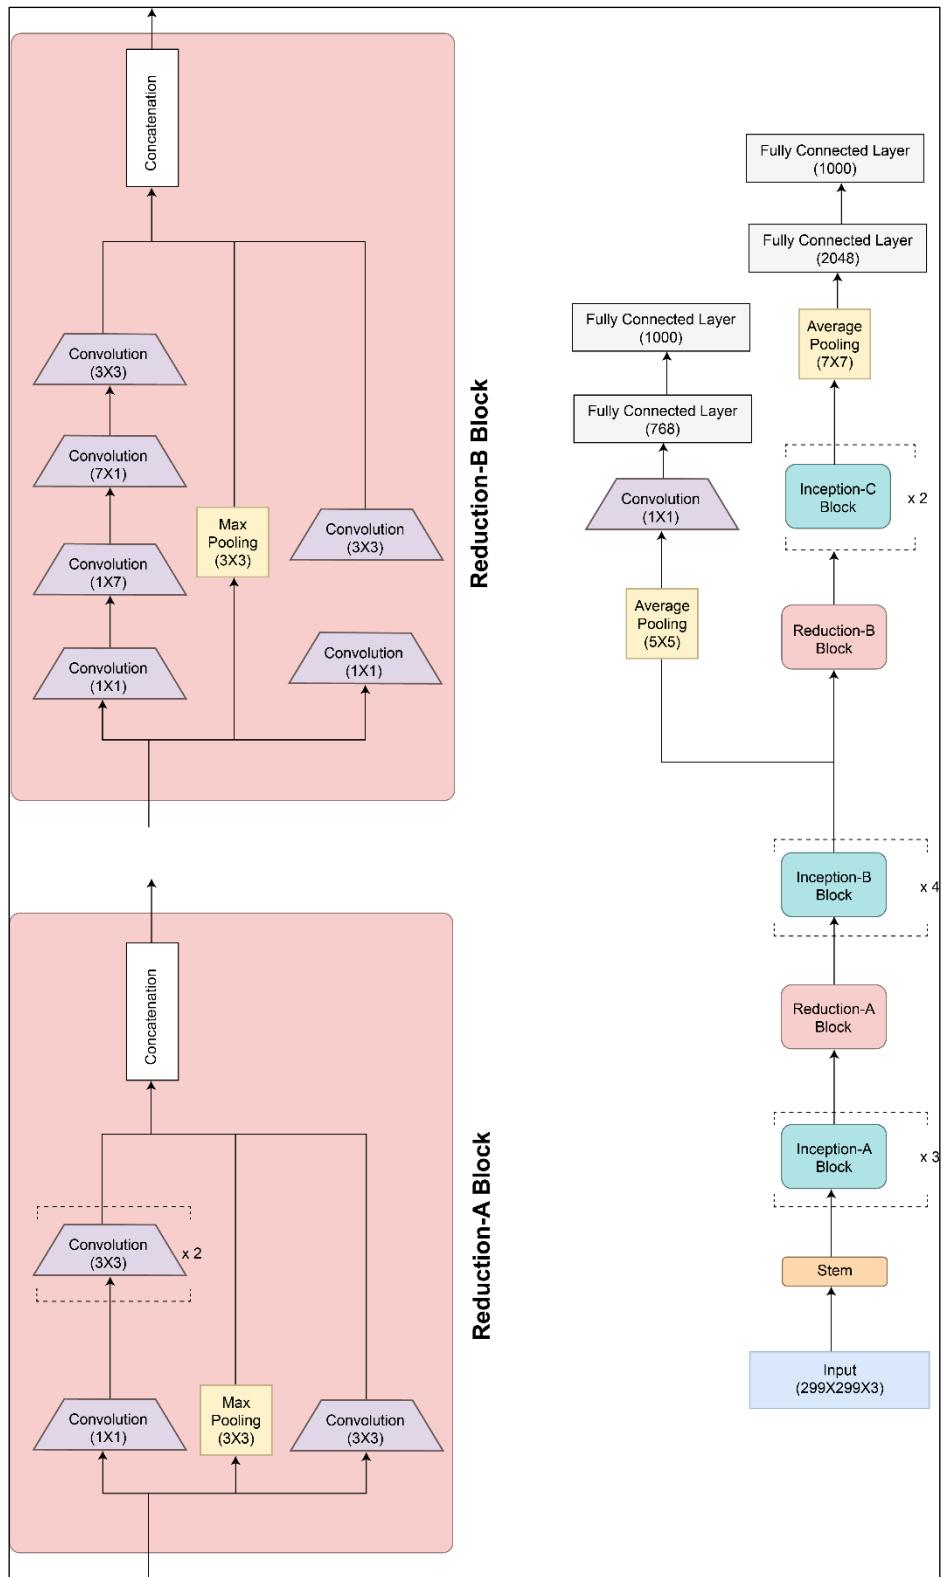
pred: bee







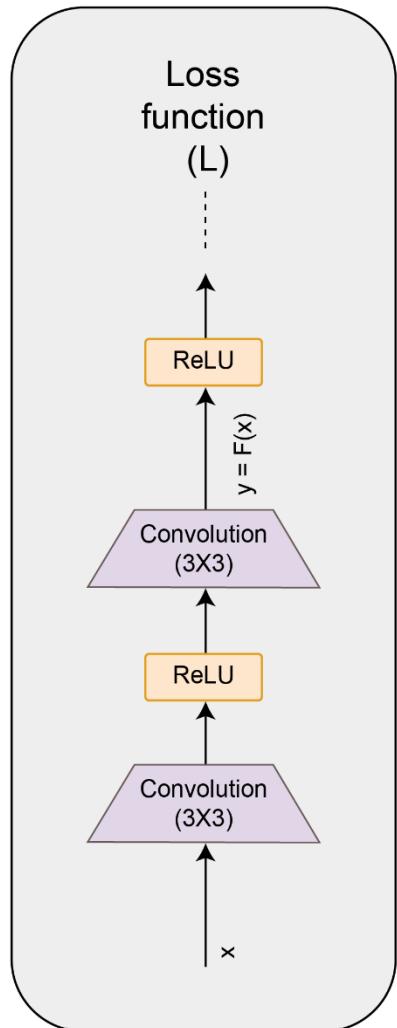




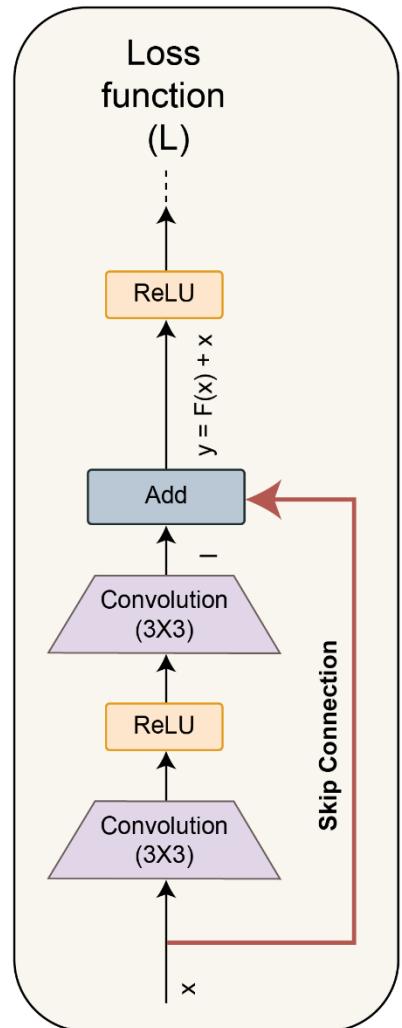
$$\begin{aligned}\frac{\delta L}{\delta x} &= \frac{\delta L}{\delta y} * \frac{\delta y}{\delta x} \\ &= \frac{\delta L}{\delta y} * F'(x)\end{aligned}$$

$$\begin{aligned}\frac{\delta L}{\delta x} &= \frac{\delta L}{\delta y} * \frac{\delta y}{\delta x} \\ &= \frac{\delta L}{\delta y} * F'(x) + \frac{\delta L}{\delta y}\end{aligned}$$

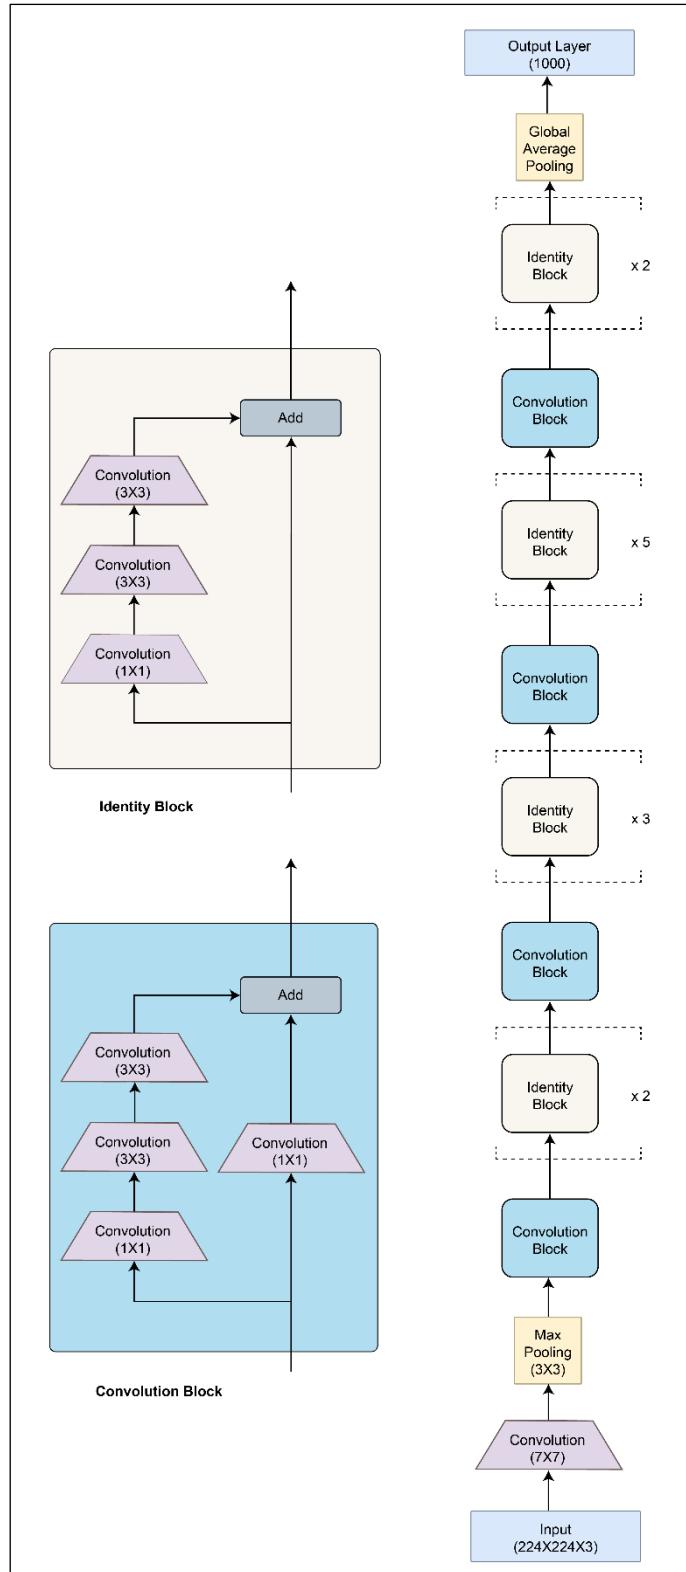
helps stabilize gradients

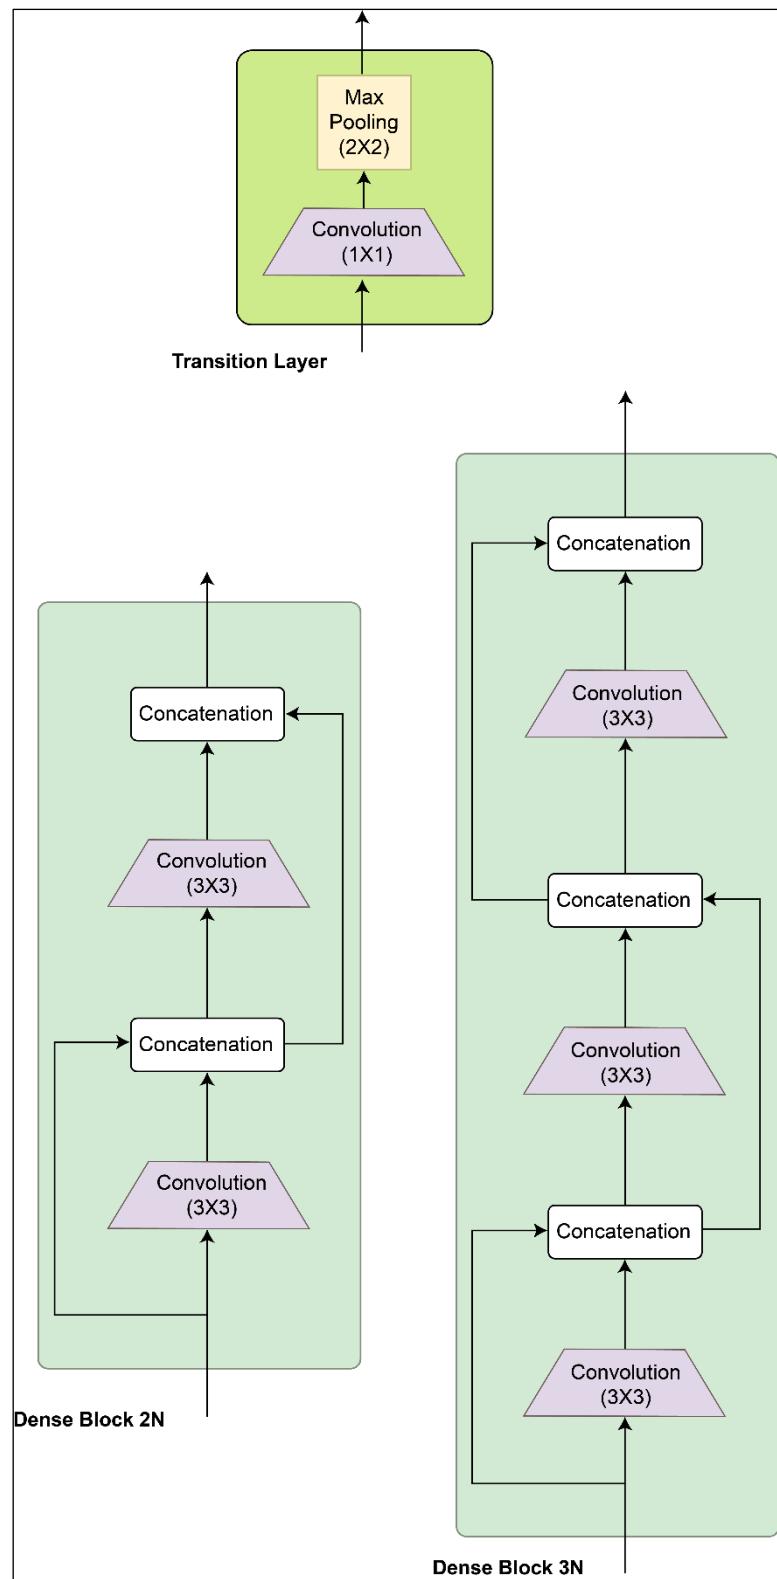


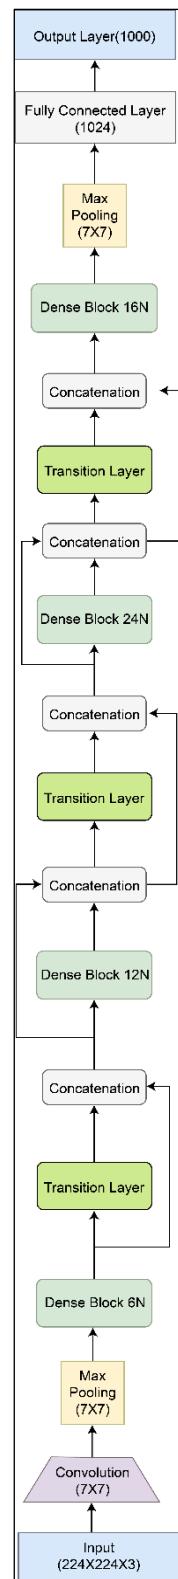
Regular Convolutional Block
(No skip connection)

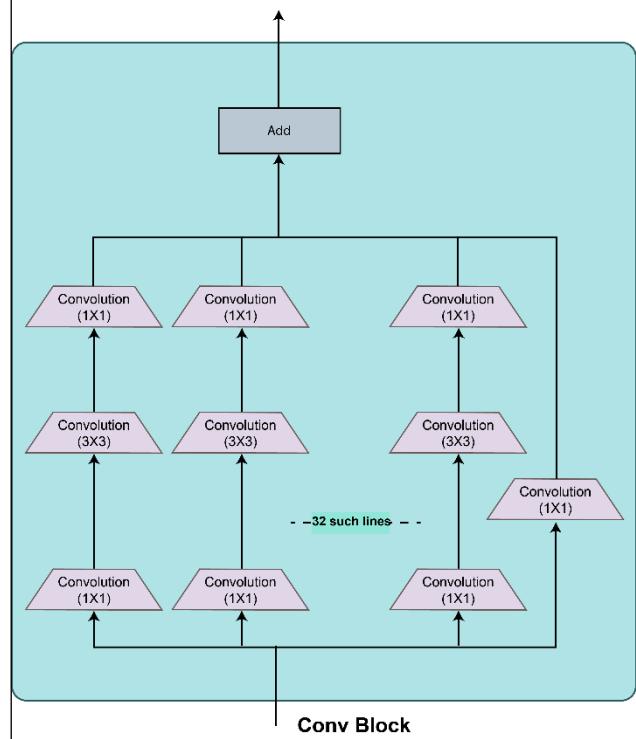
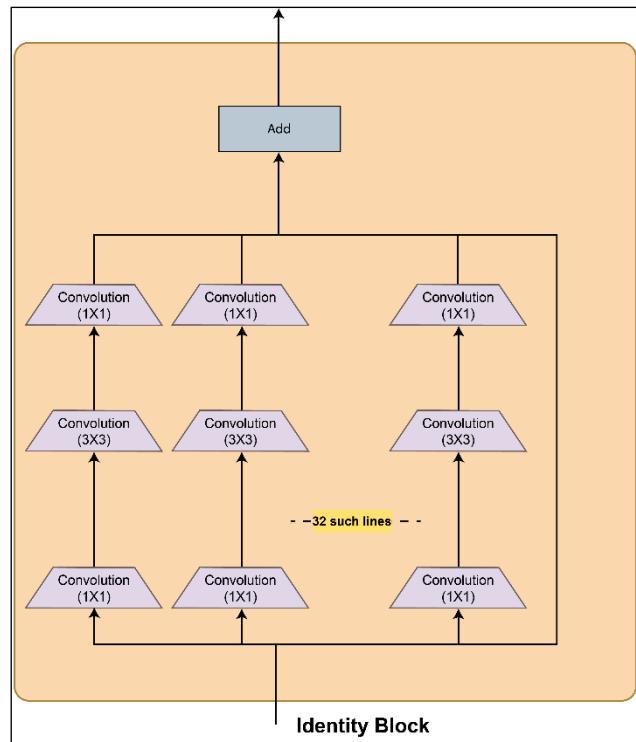


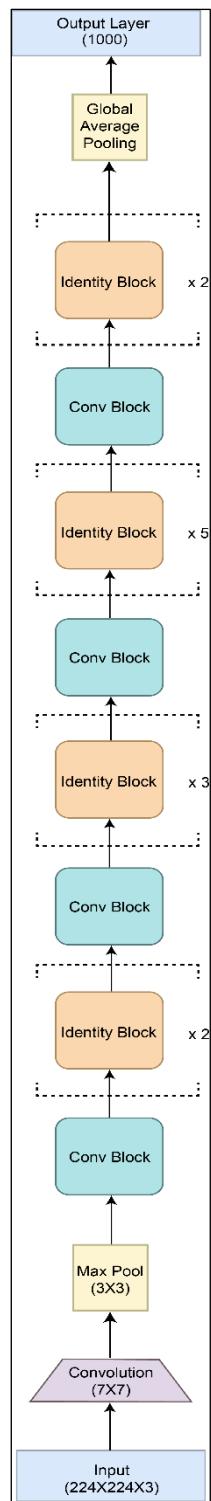
Residual Convolutional Block
(skip connection)



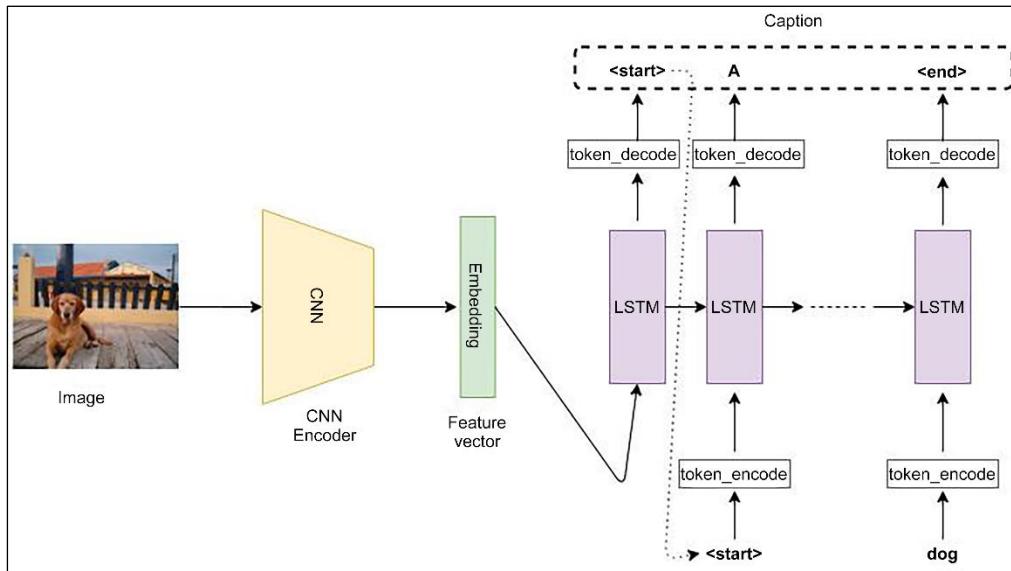




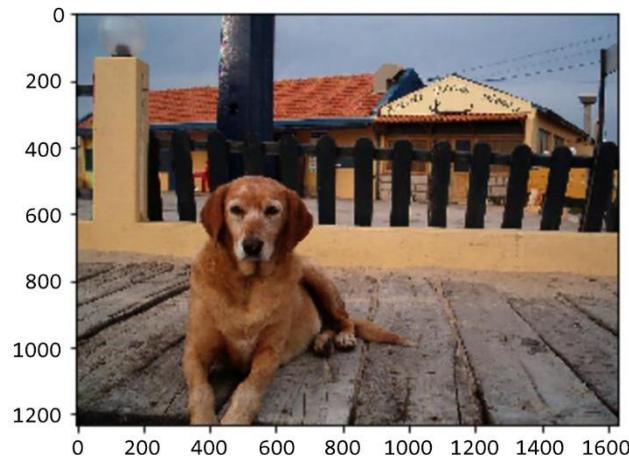




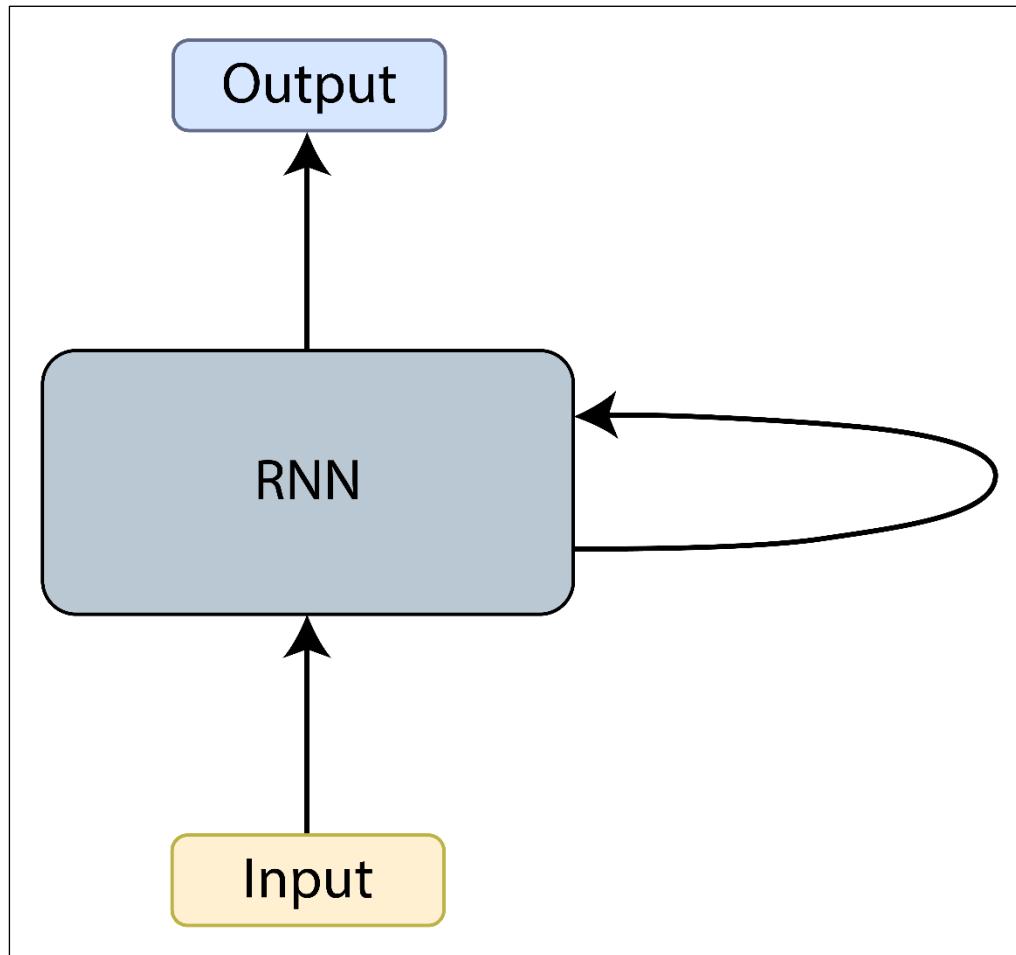
Chapter 3: Combining CNNs and LSTMs

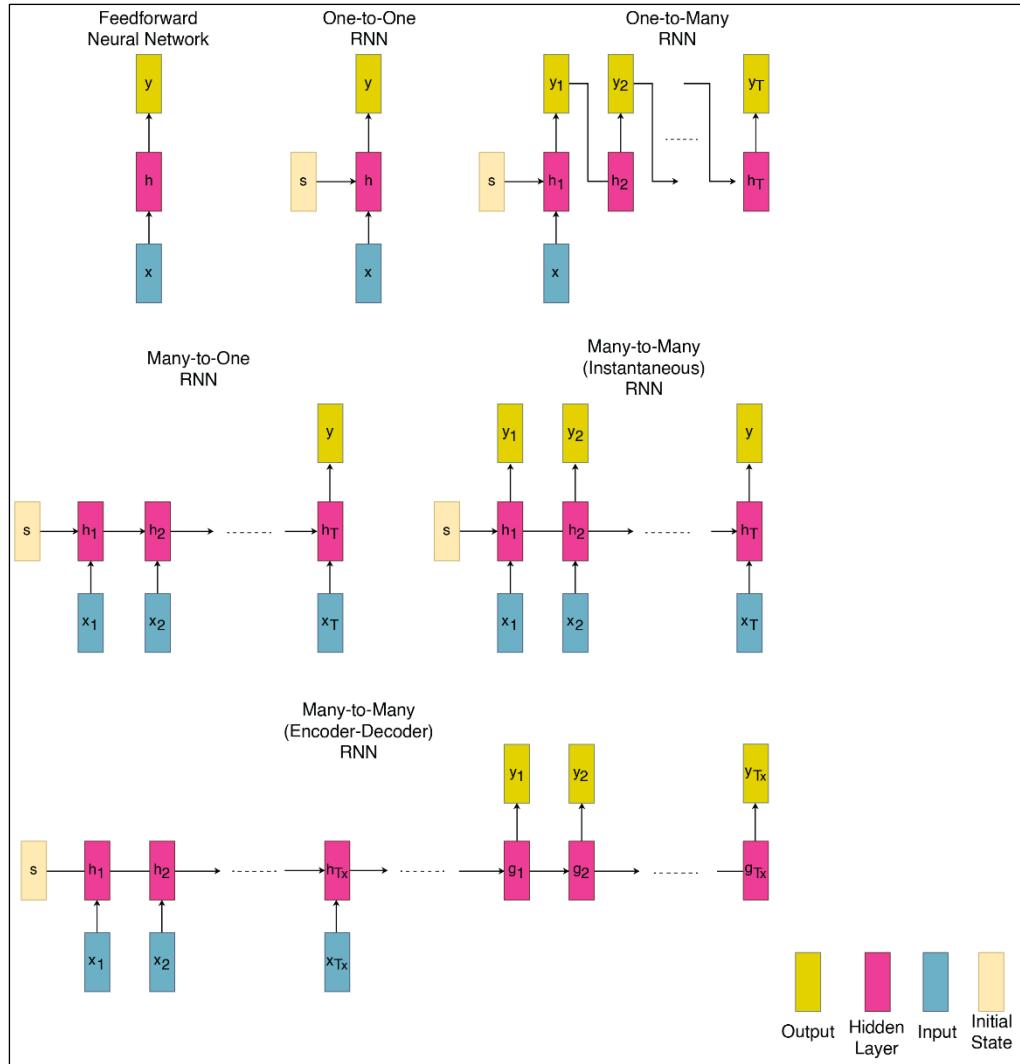


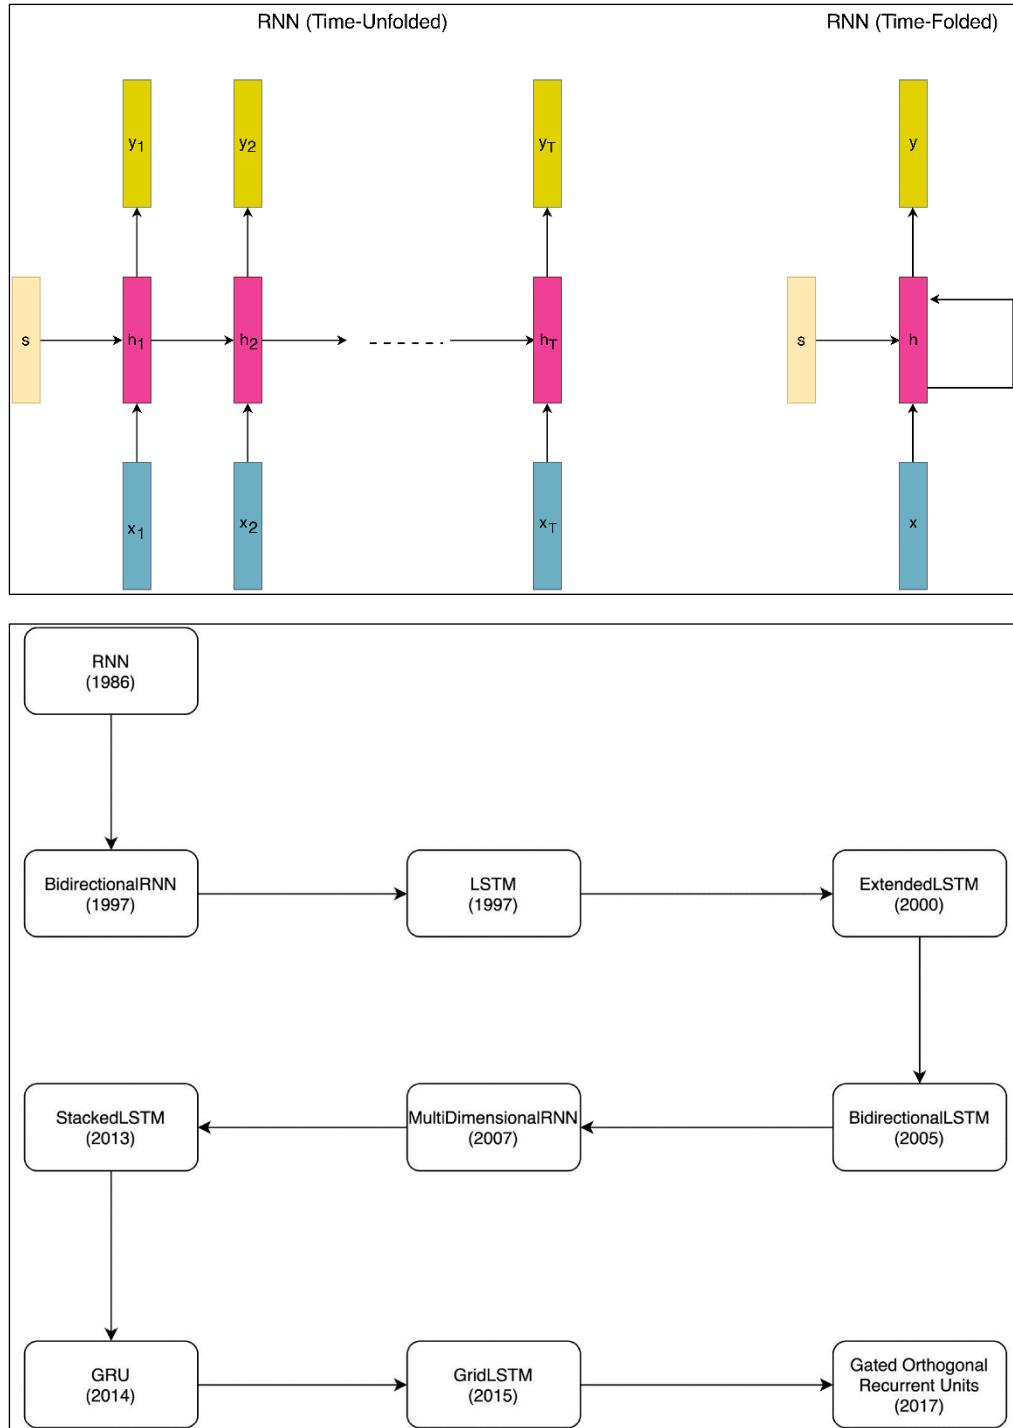
<start> a dog is standing on a sidewalk near a building . **<end>**

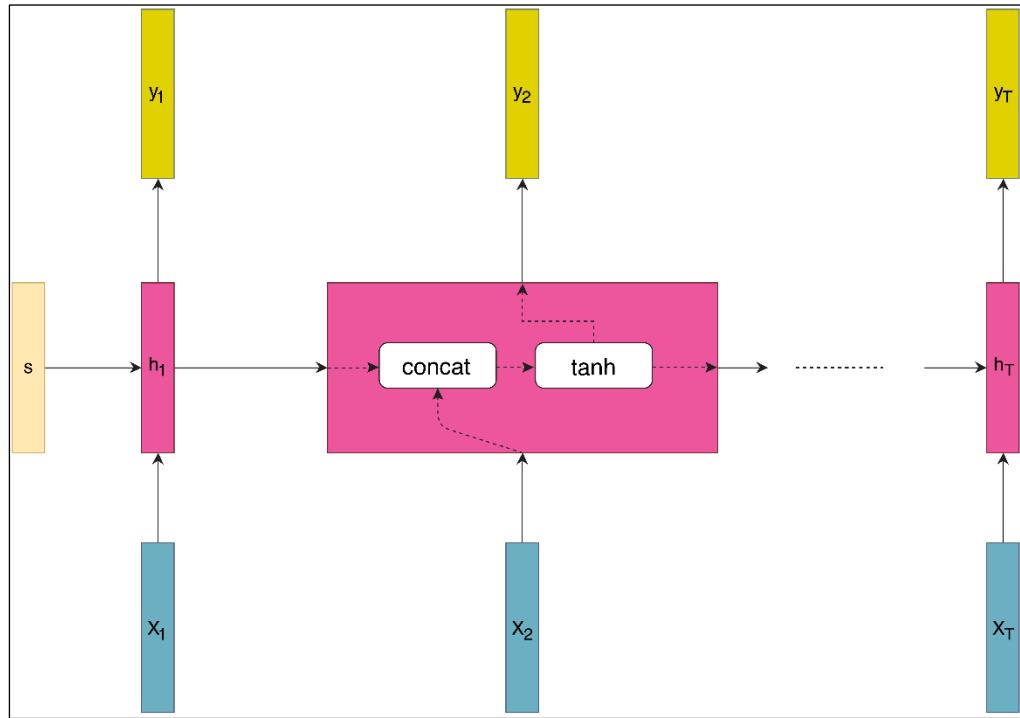
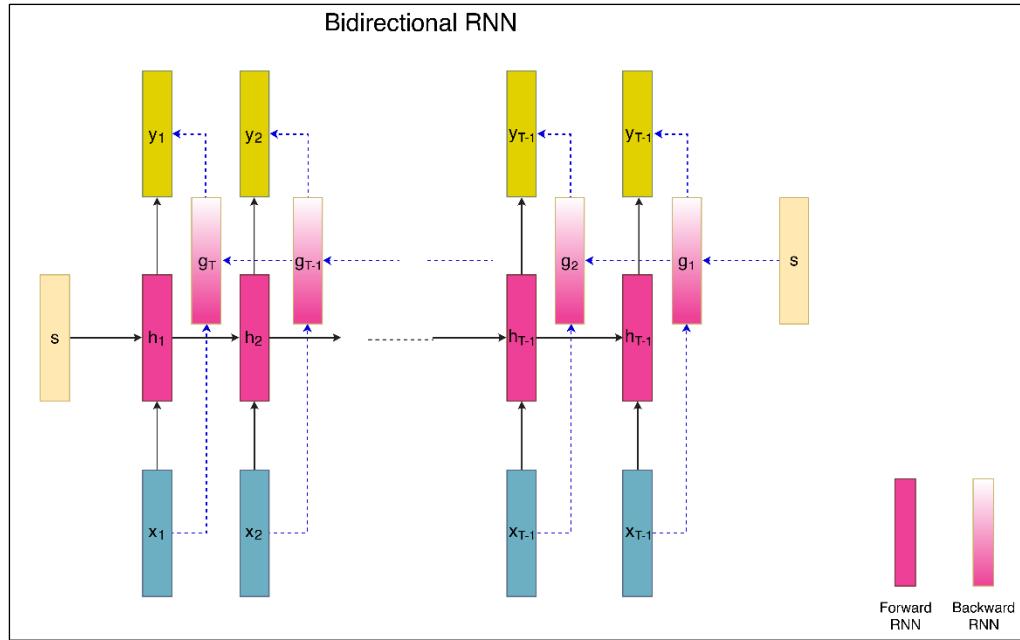


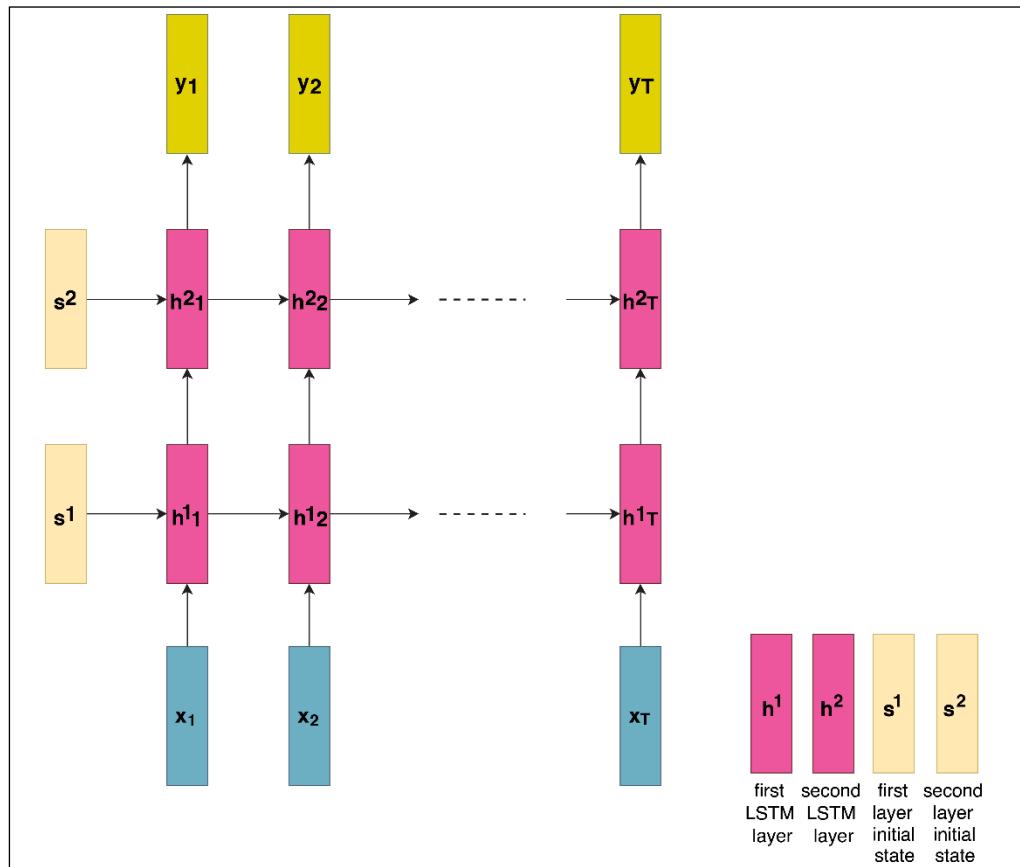
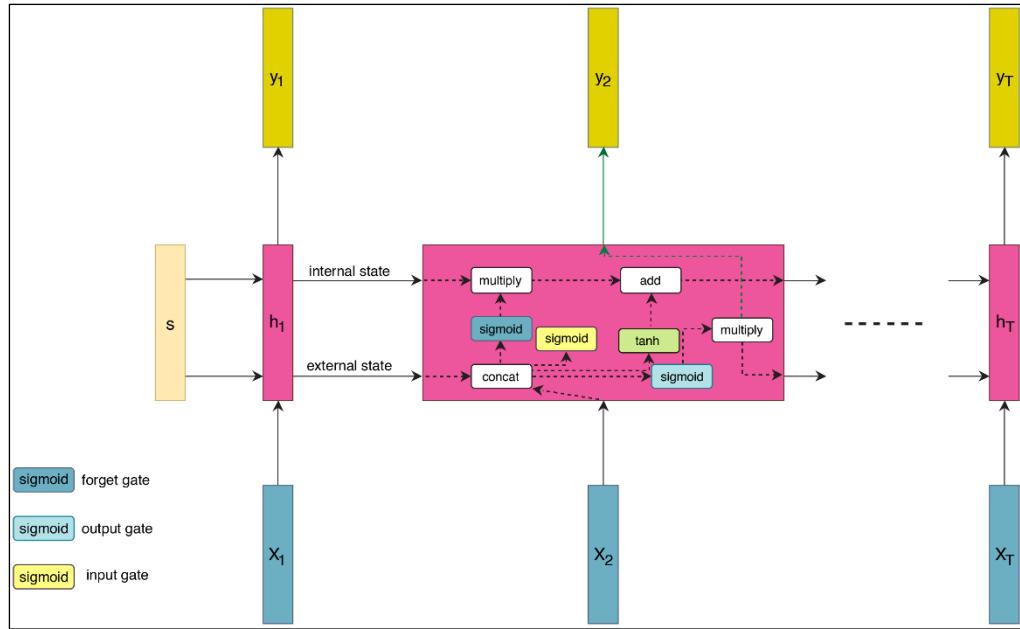
Chapter 4: Deep Recurrent Model Architectures

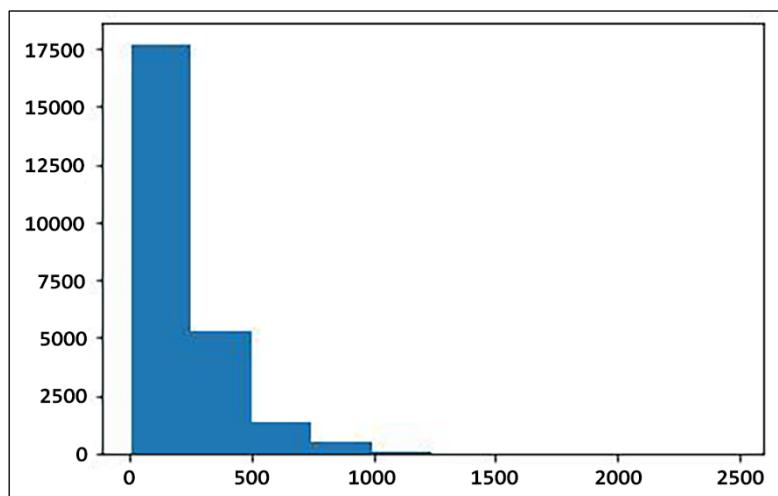
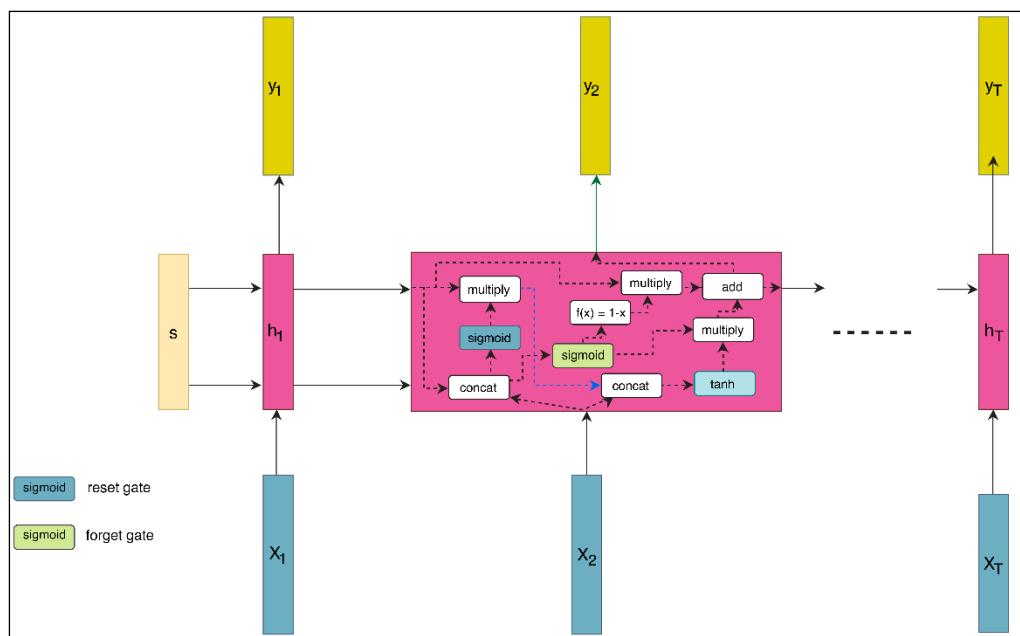


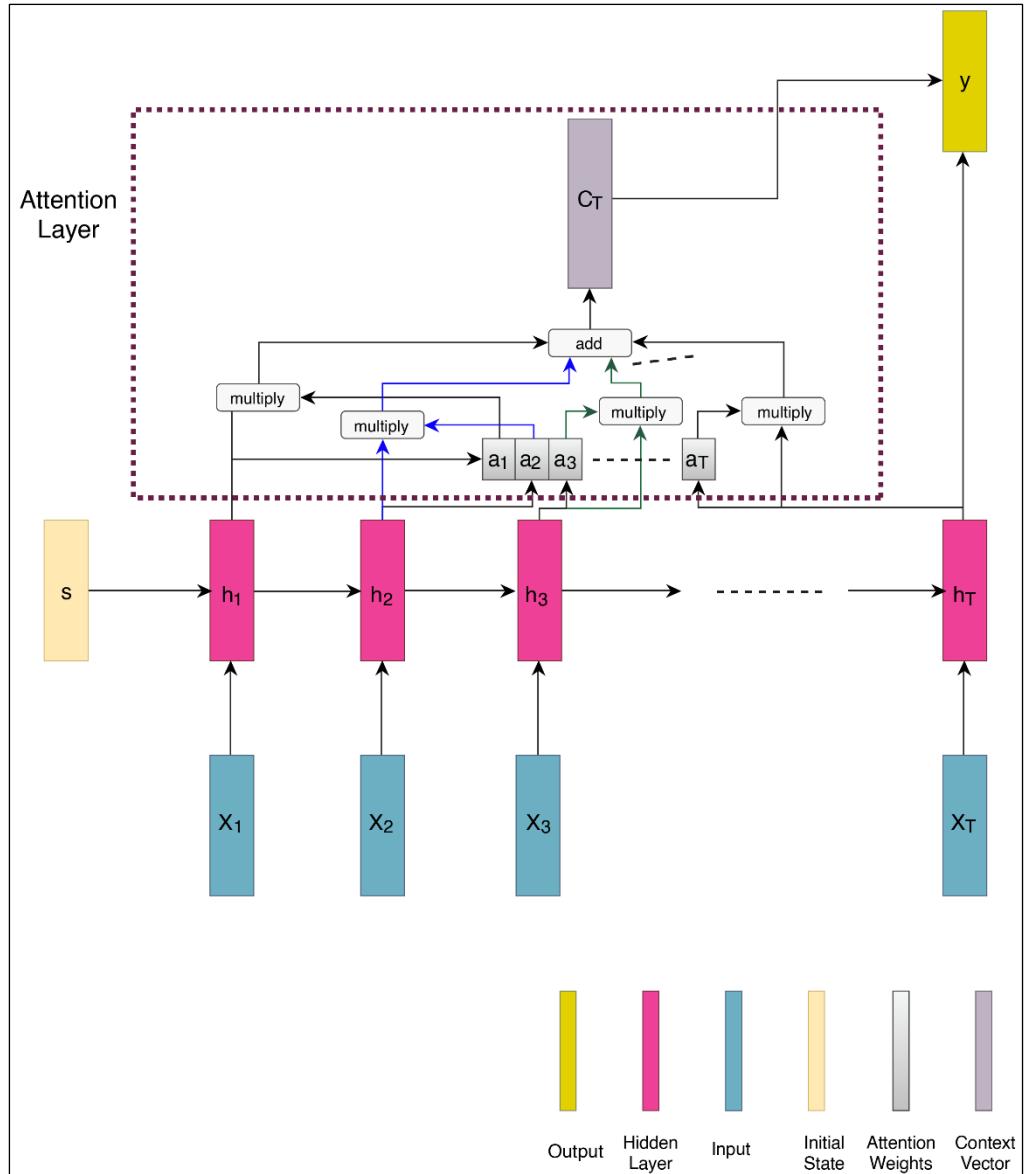




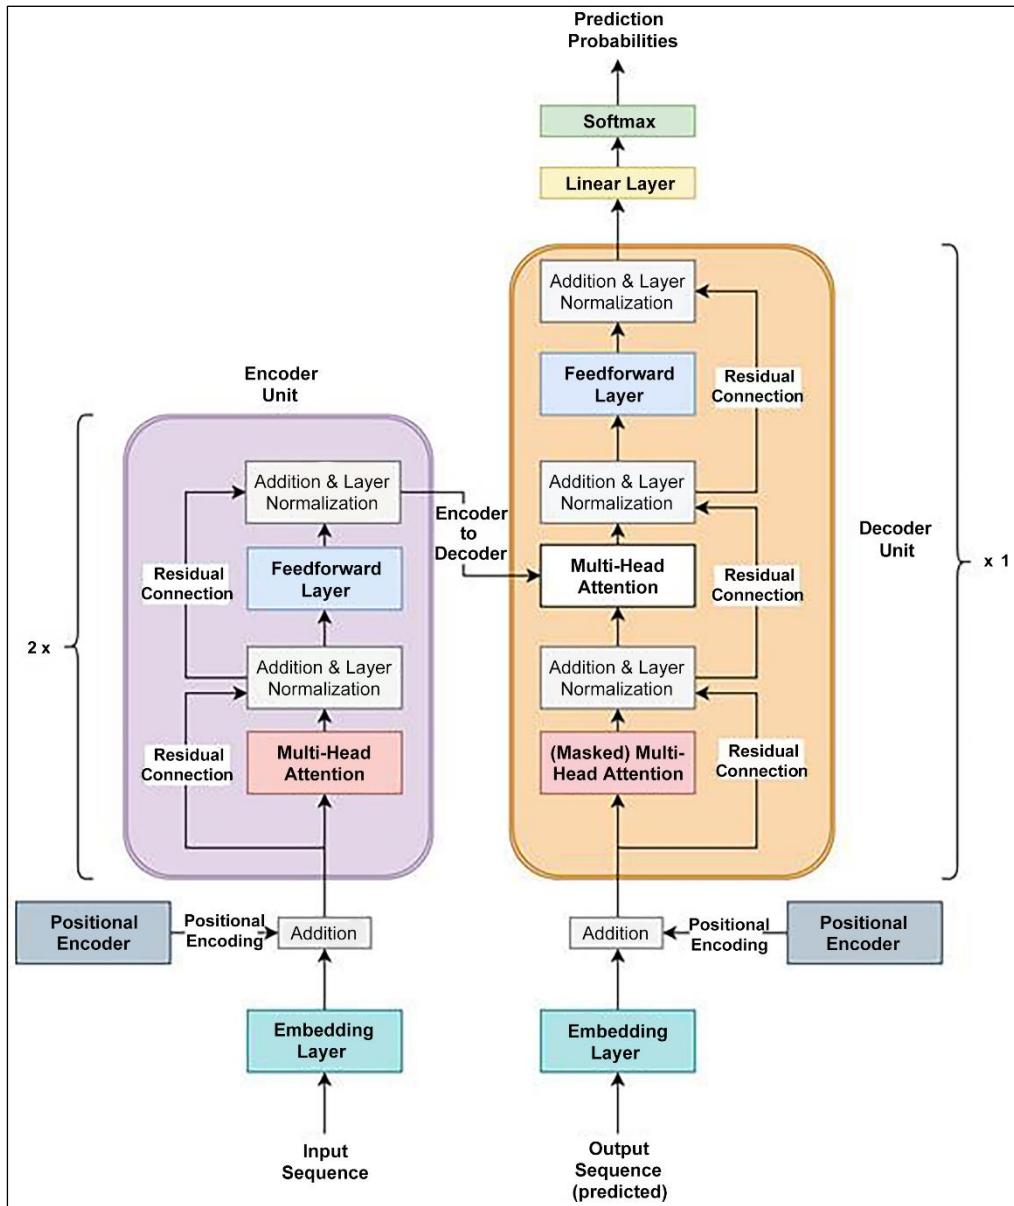


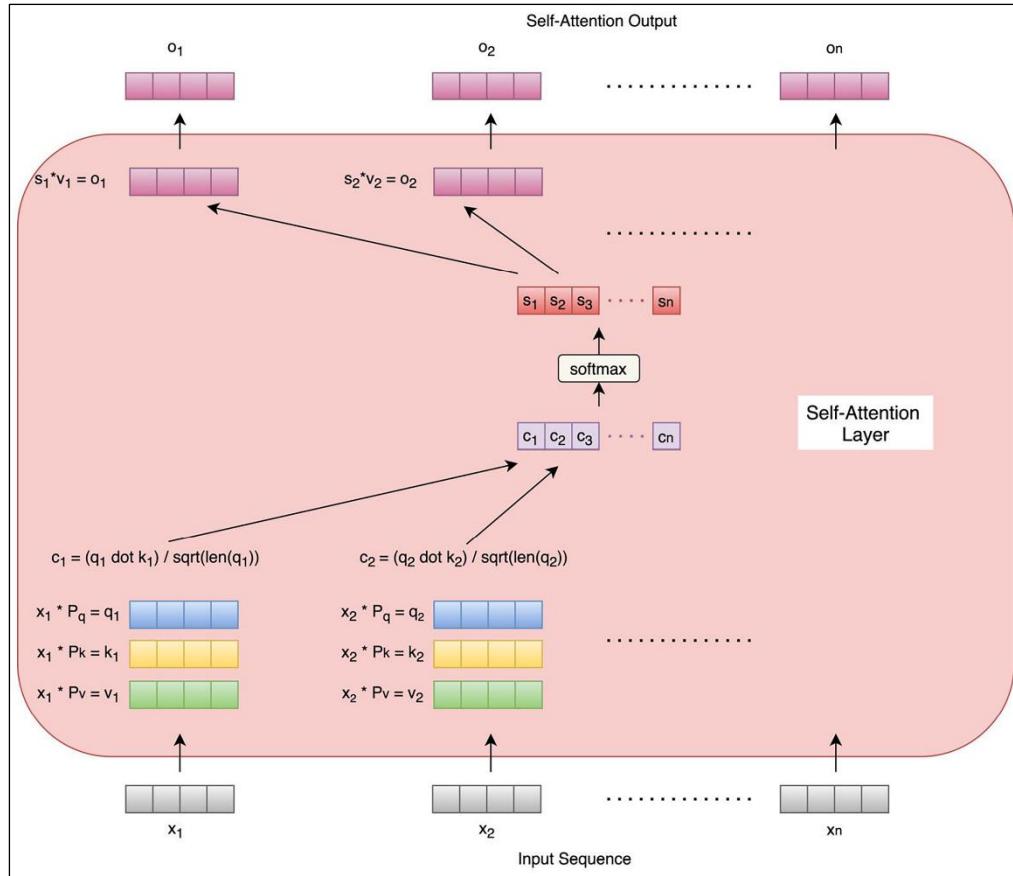


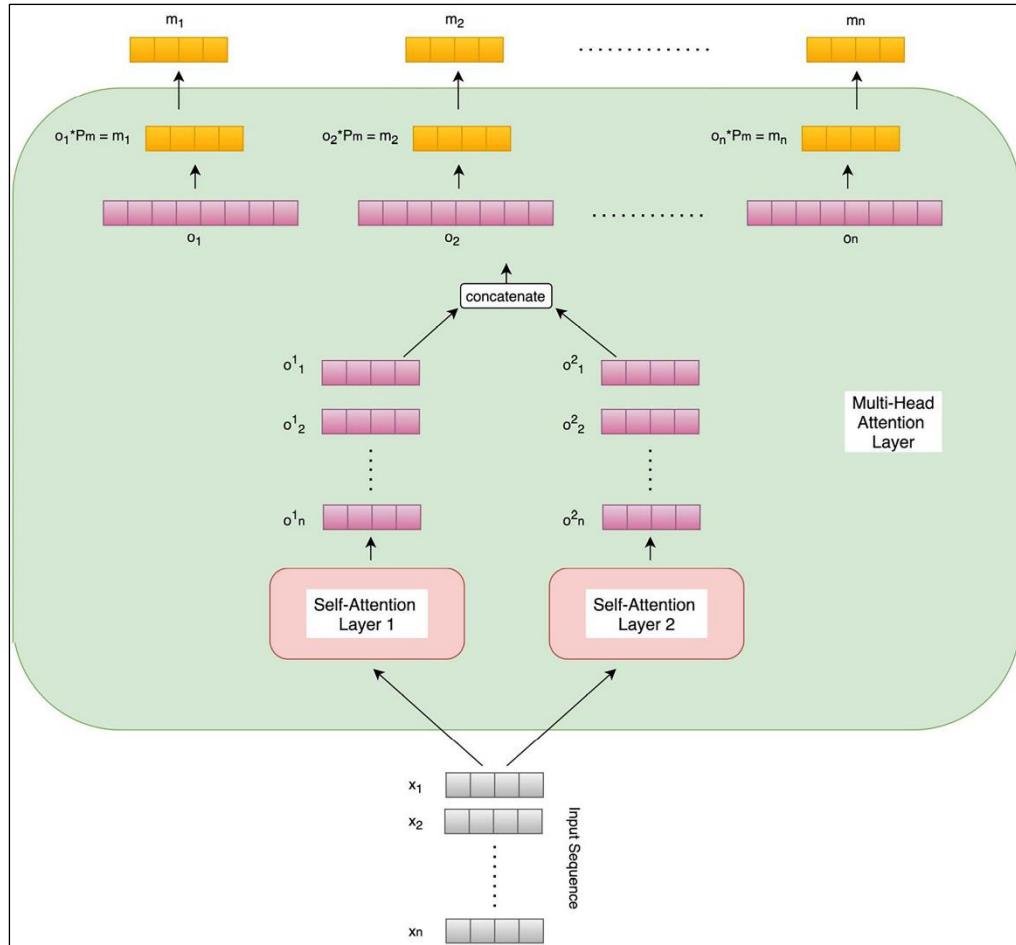


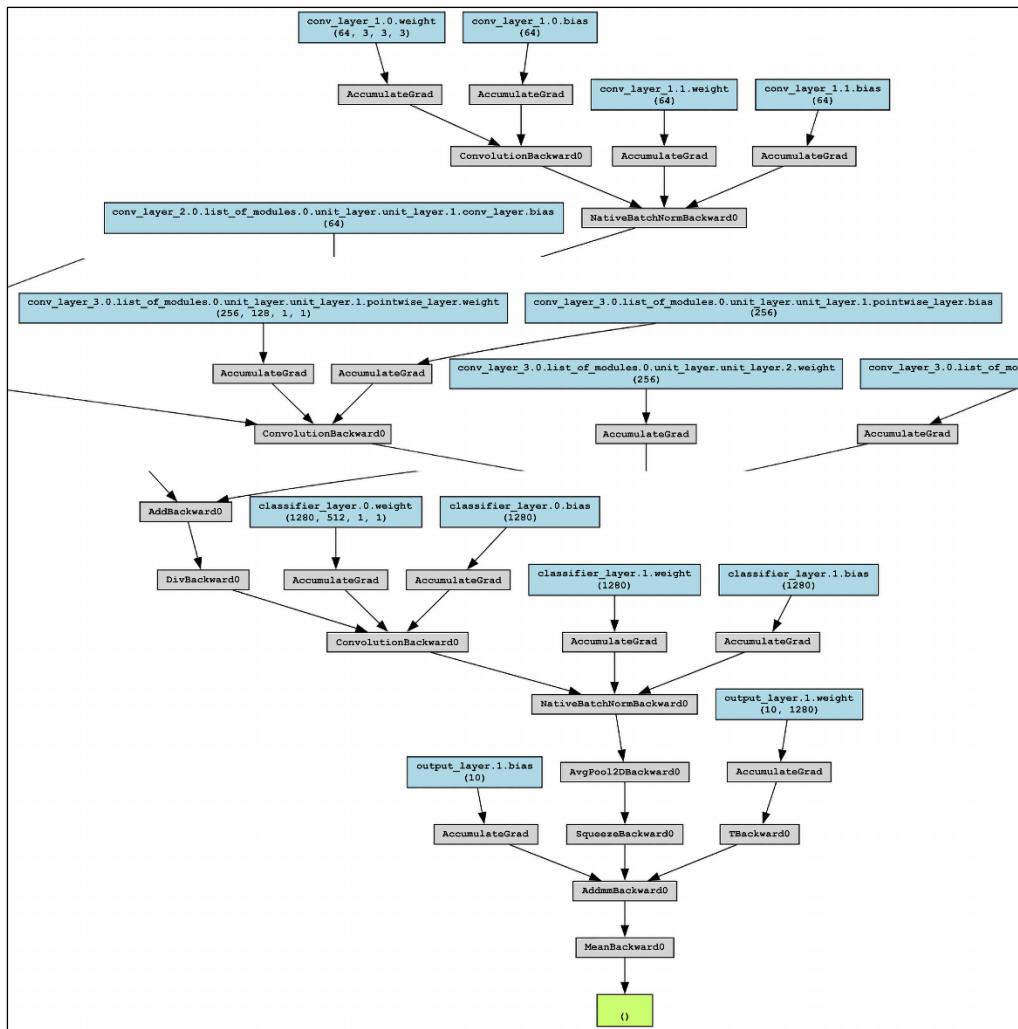


Chapter 5: Advanced Hybrid Models

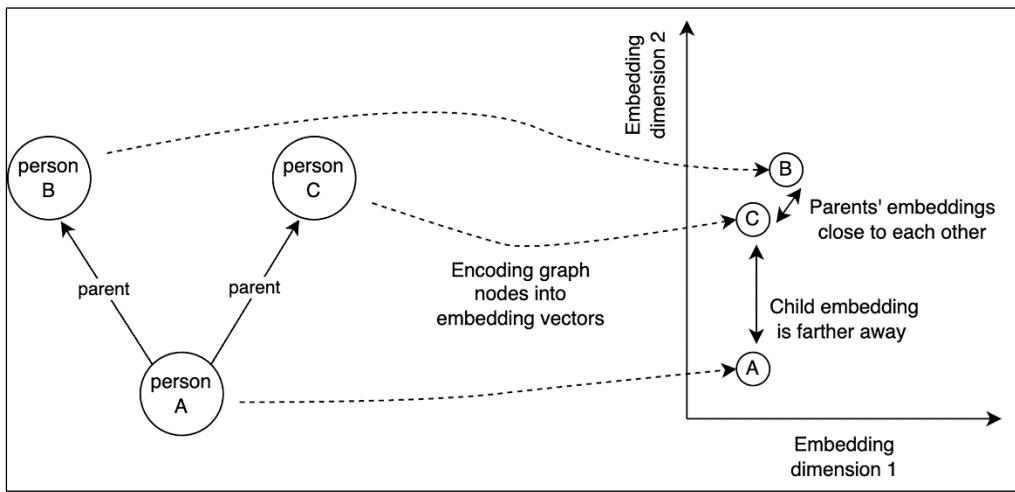
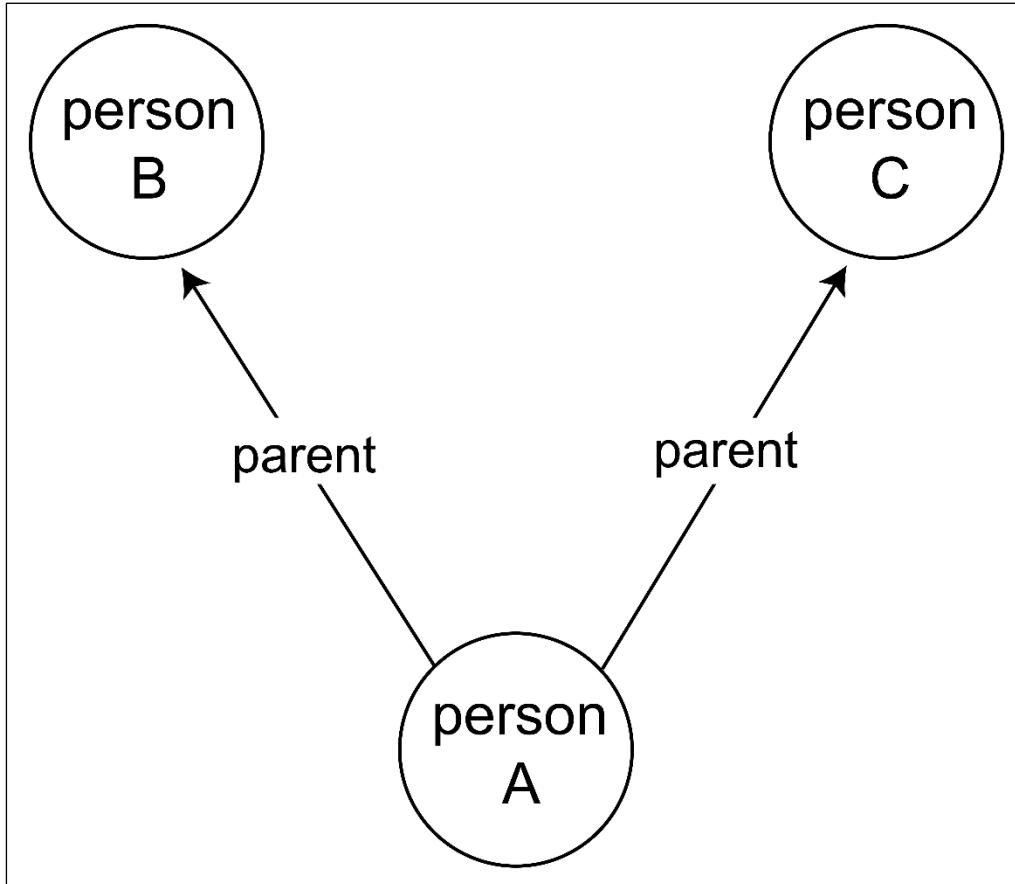


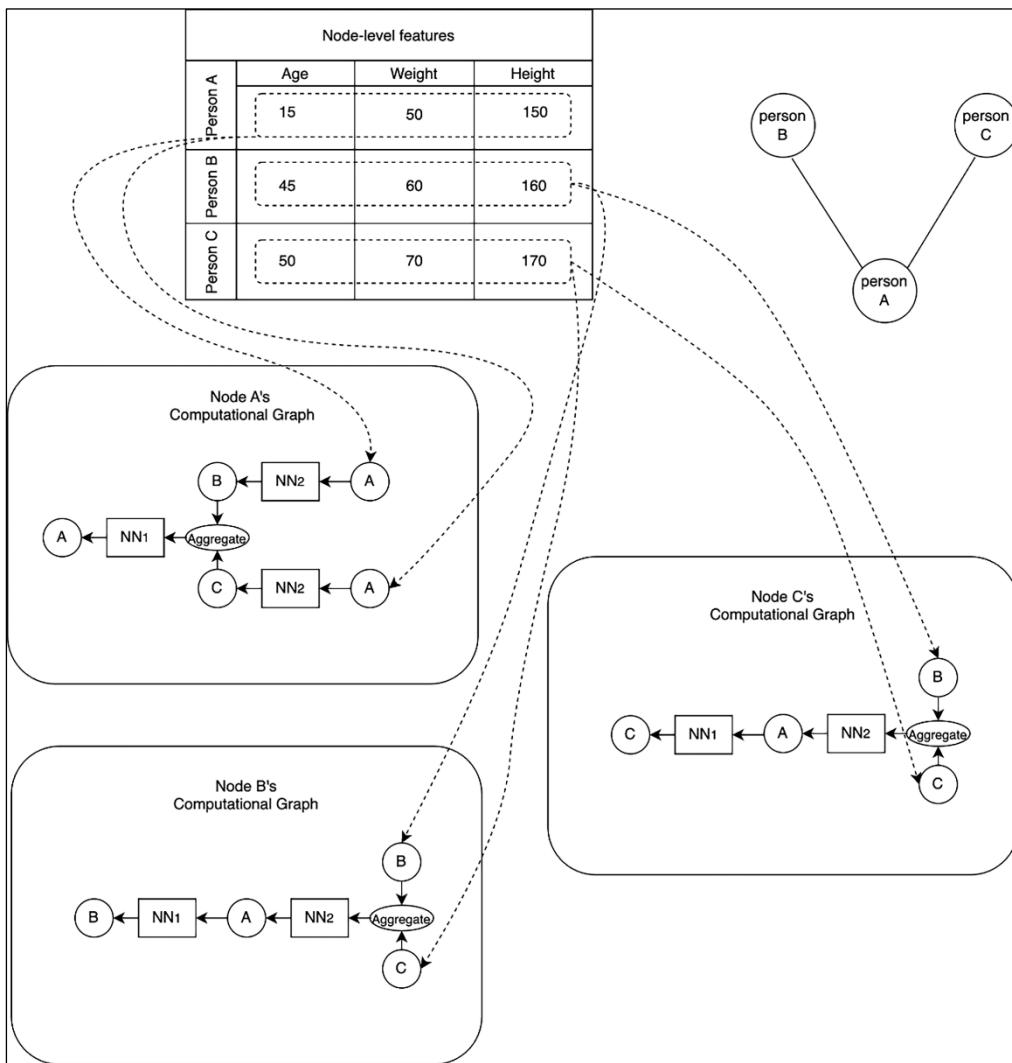
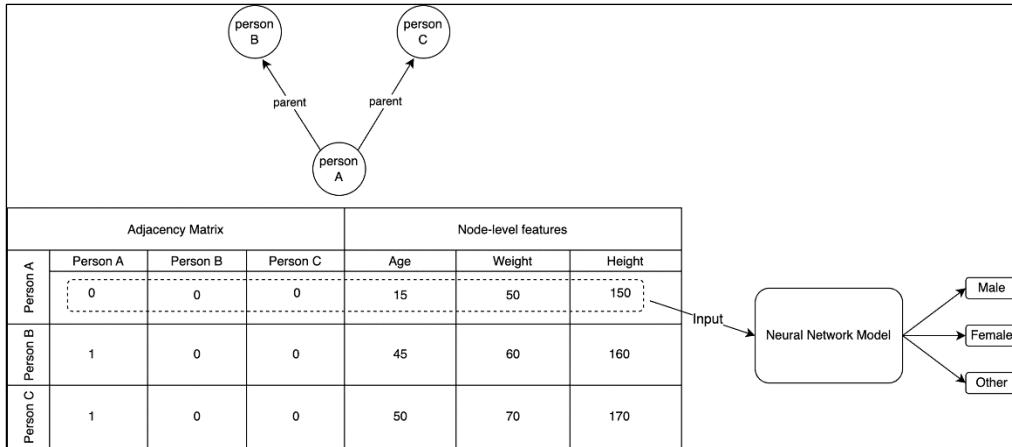


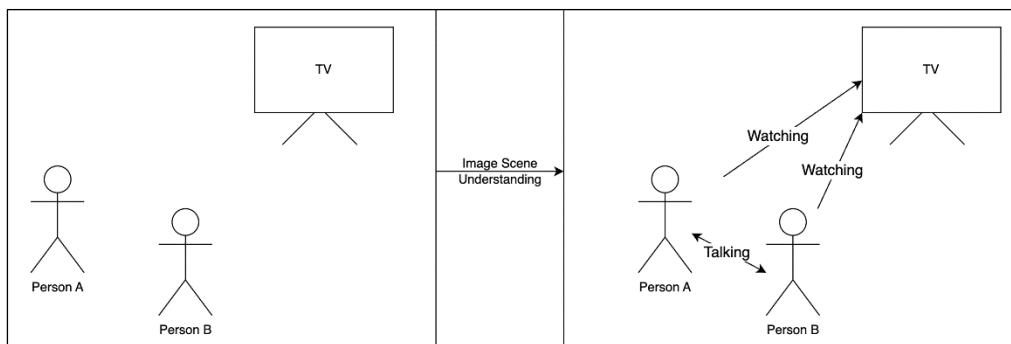
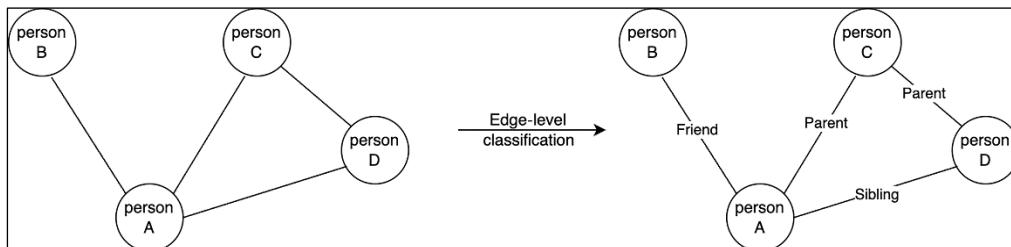
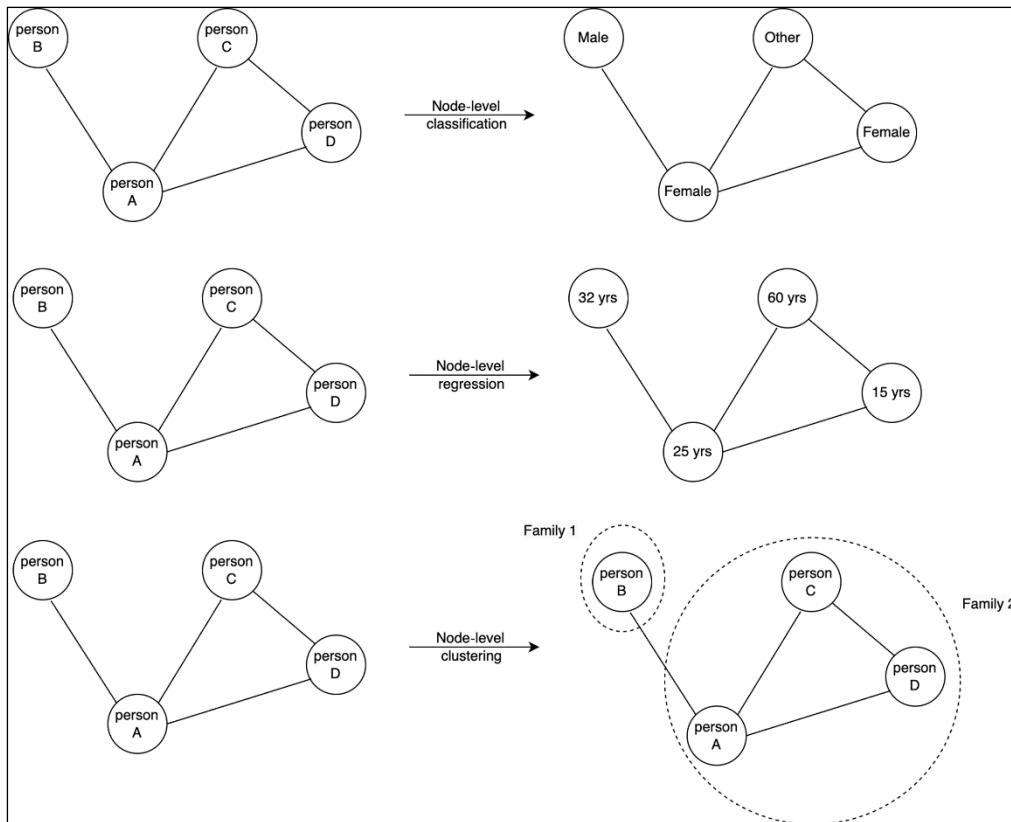


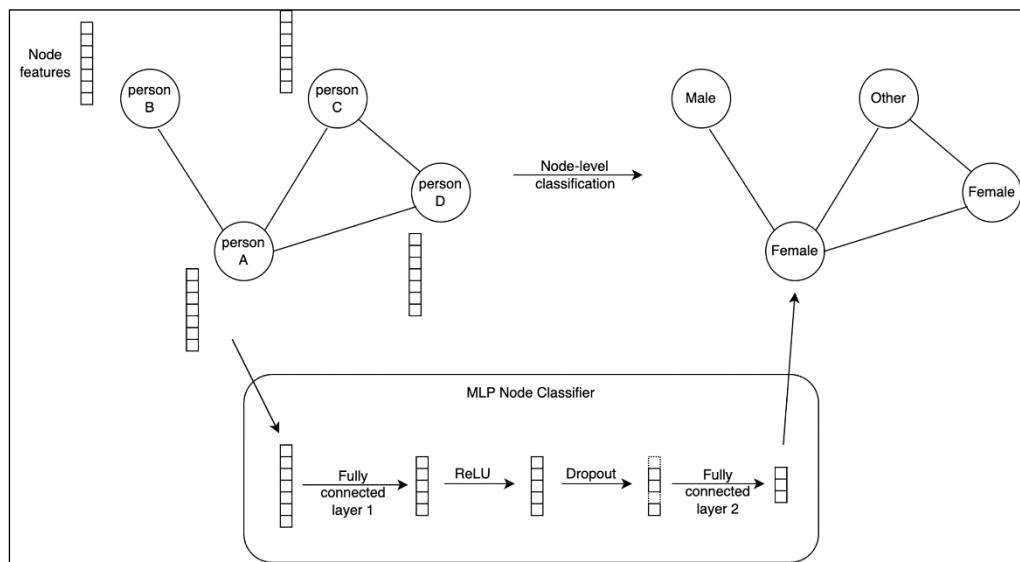
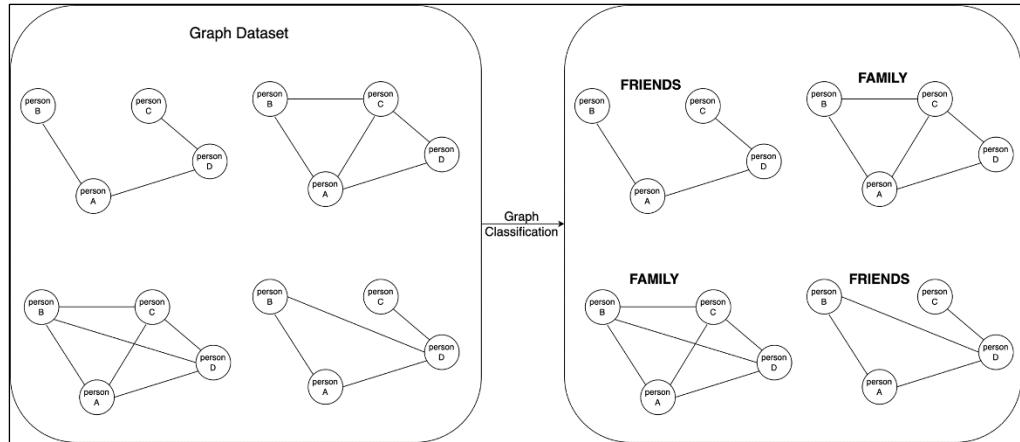


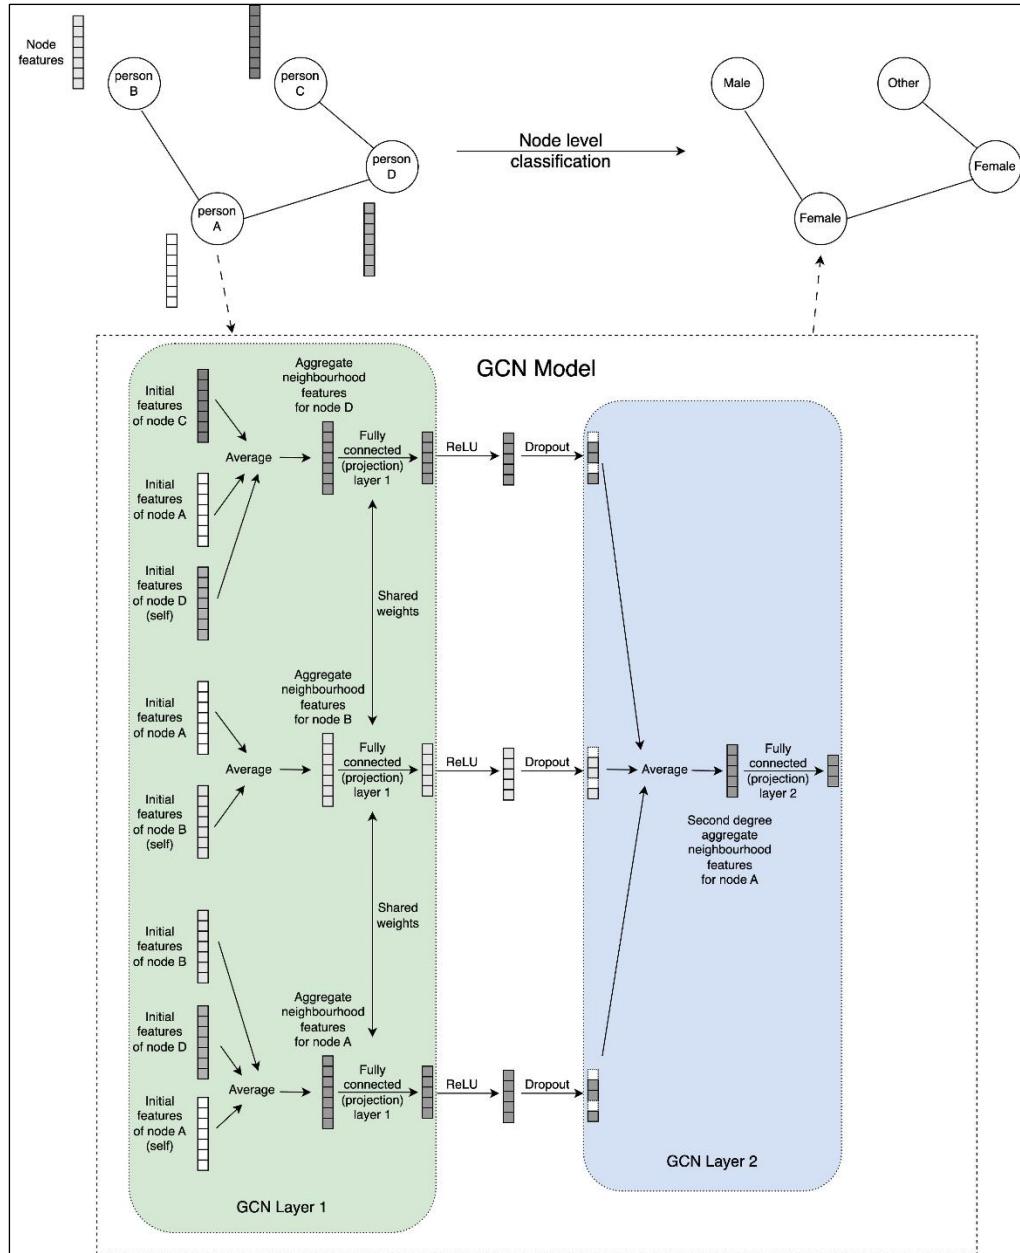
Chapter 6: Graph Neural Networks

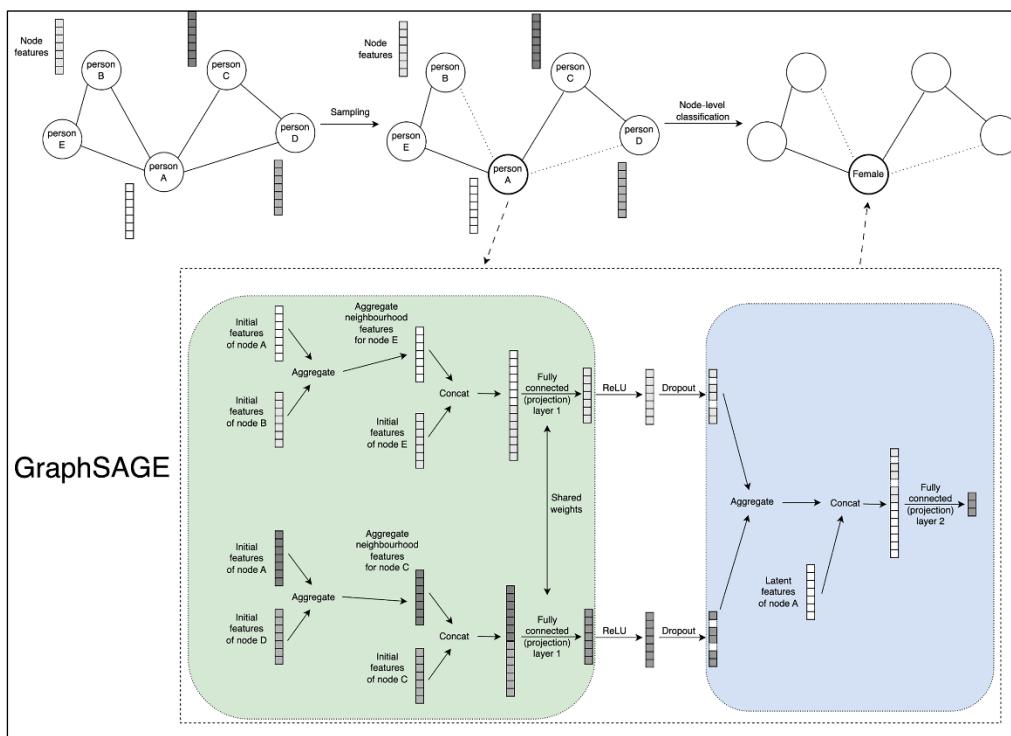
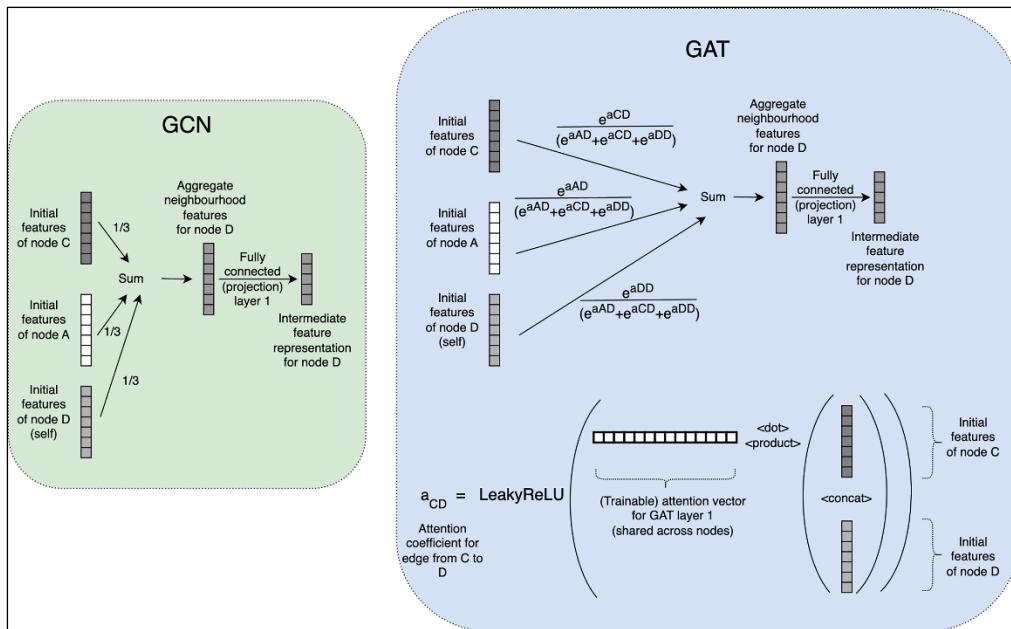


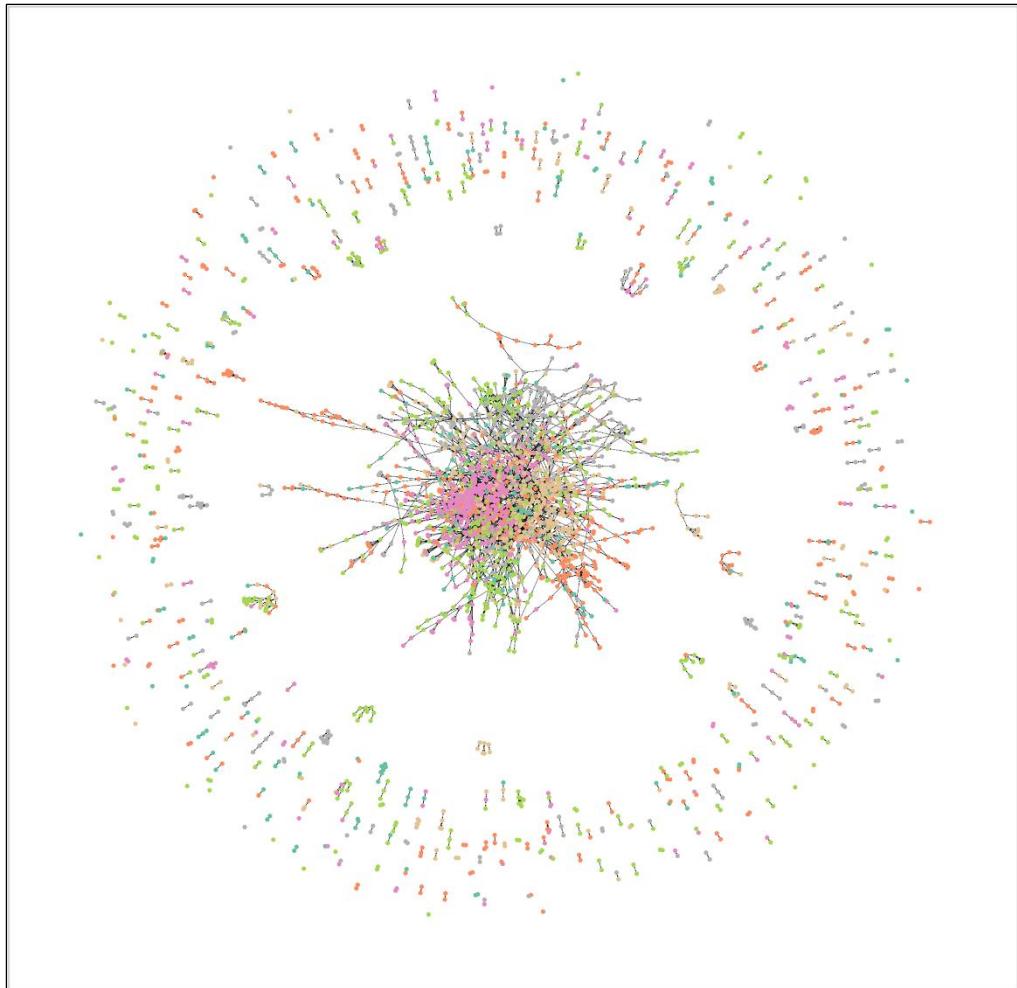


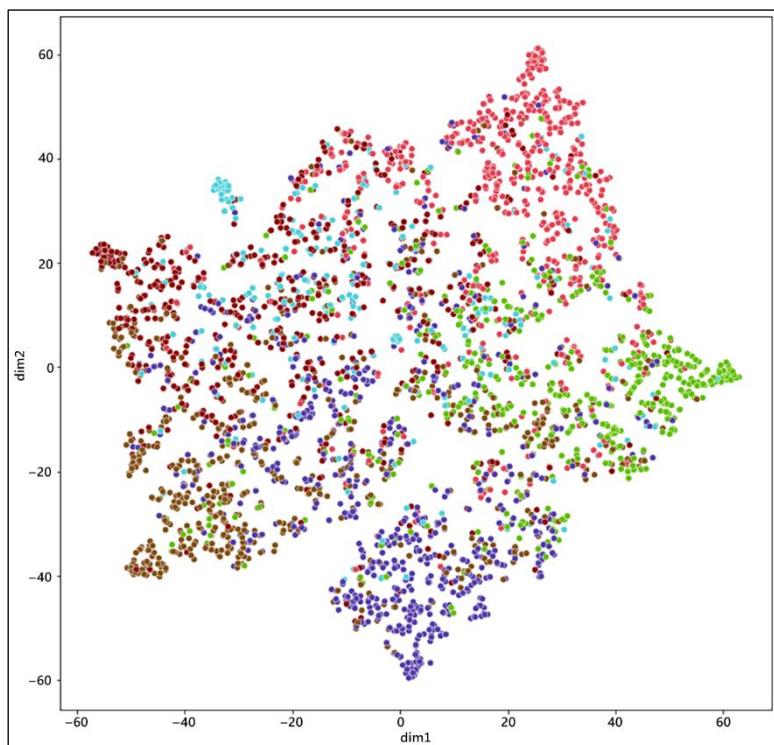
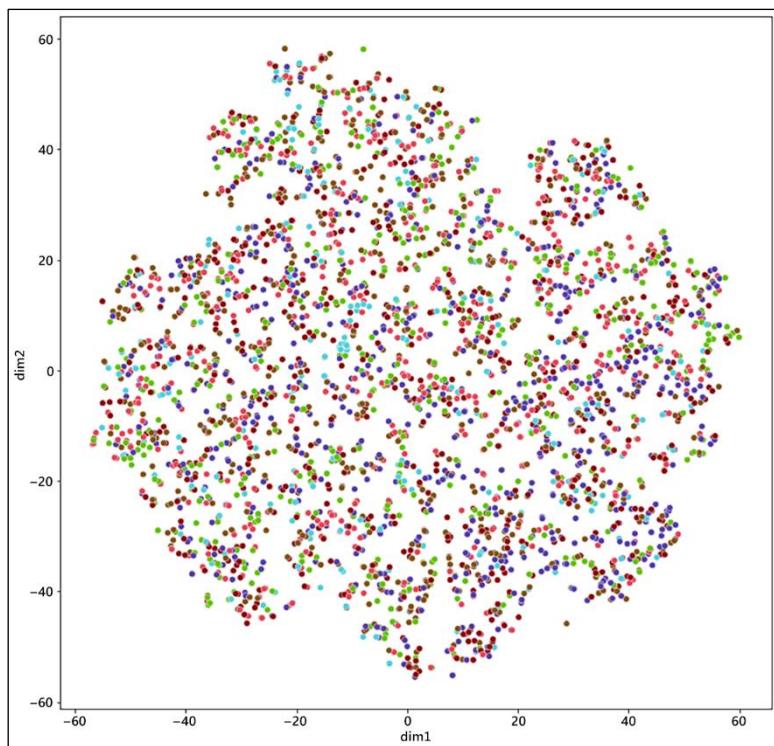


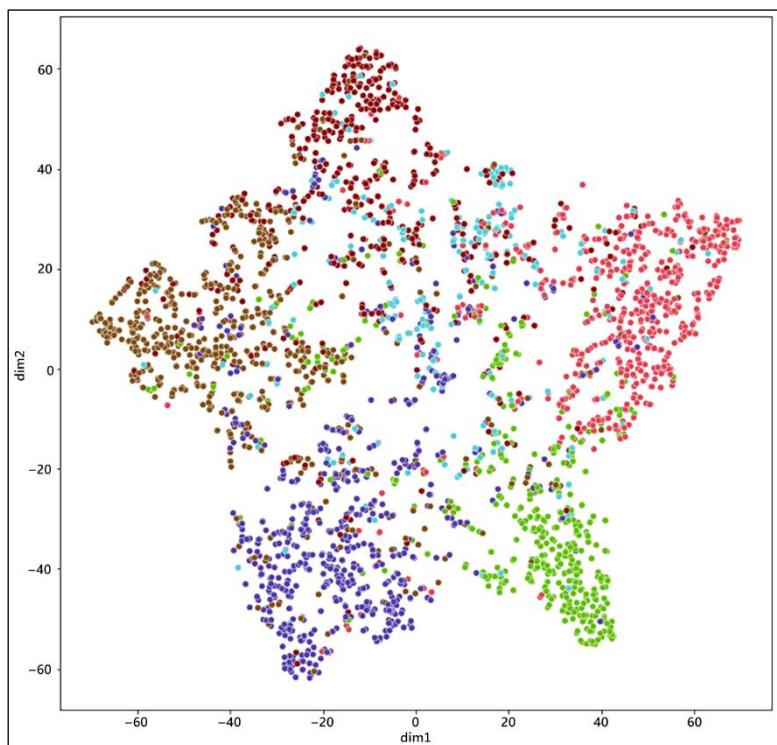
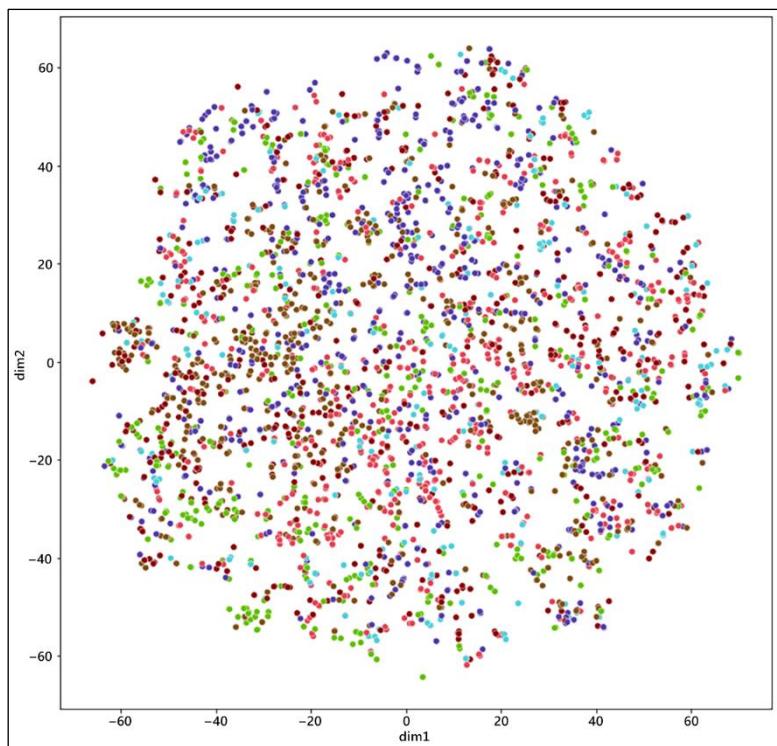


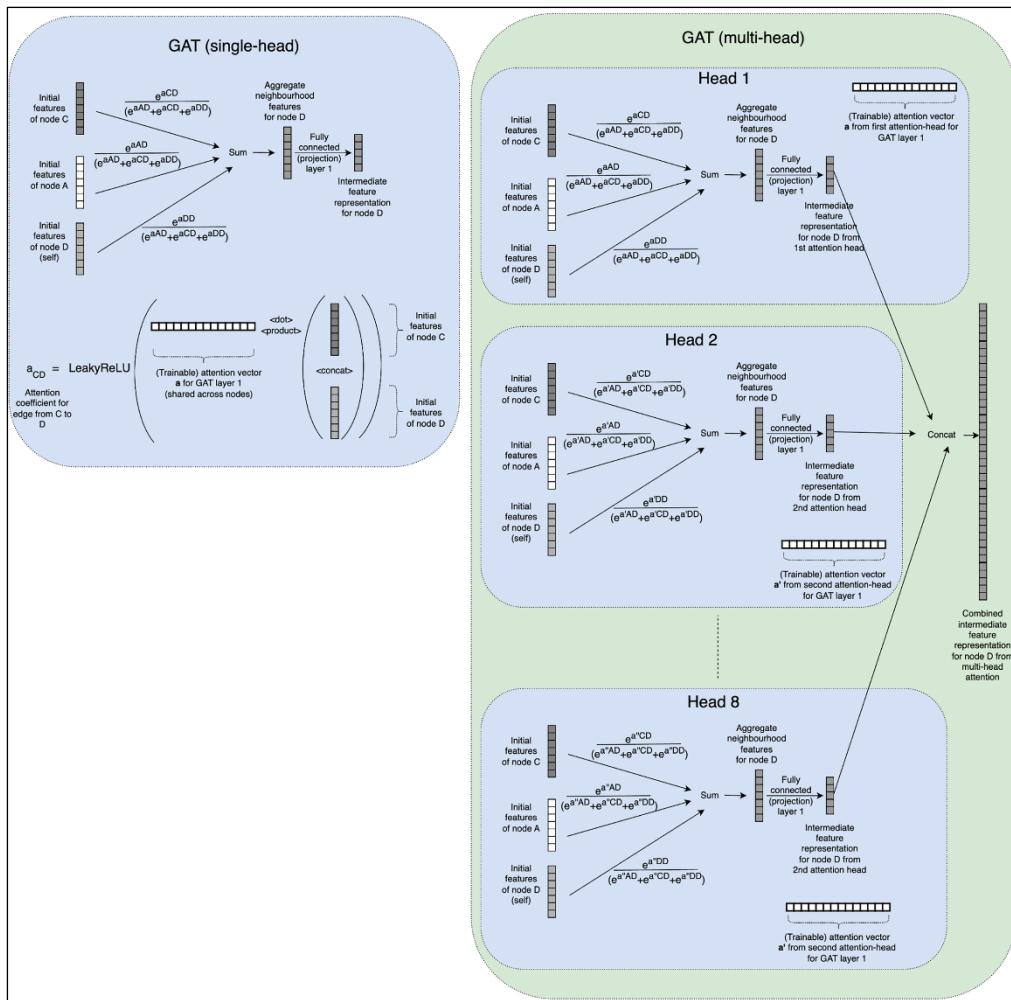


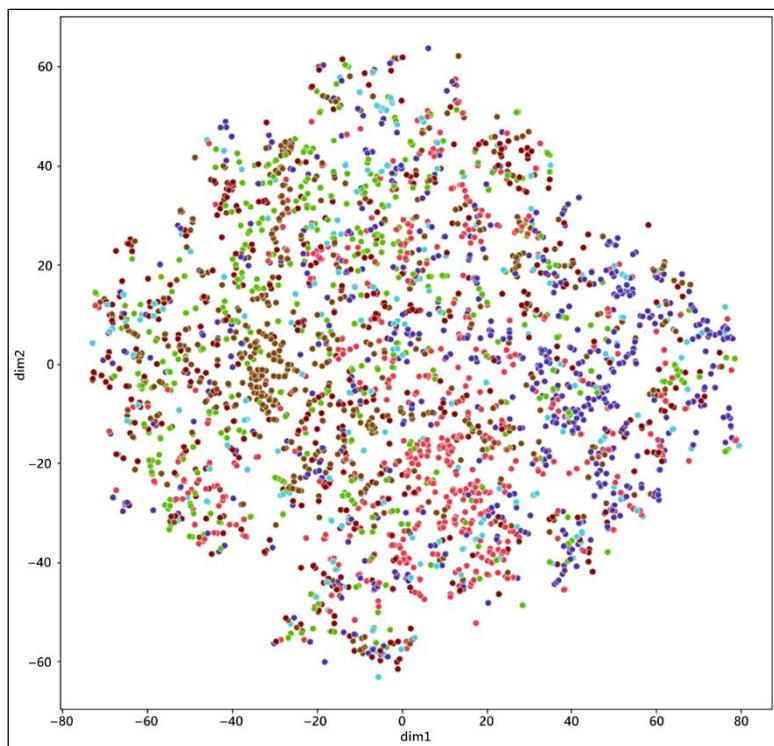


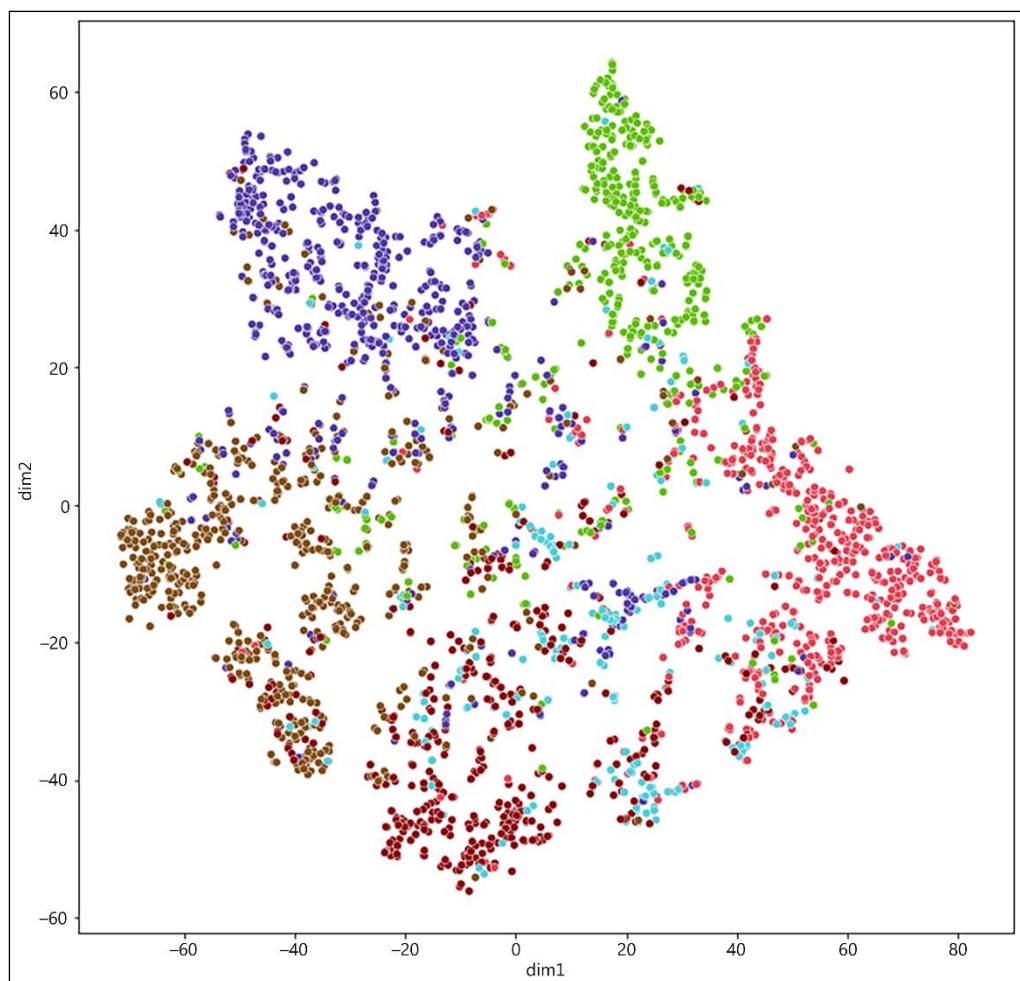




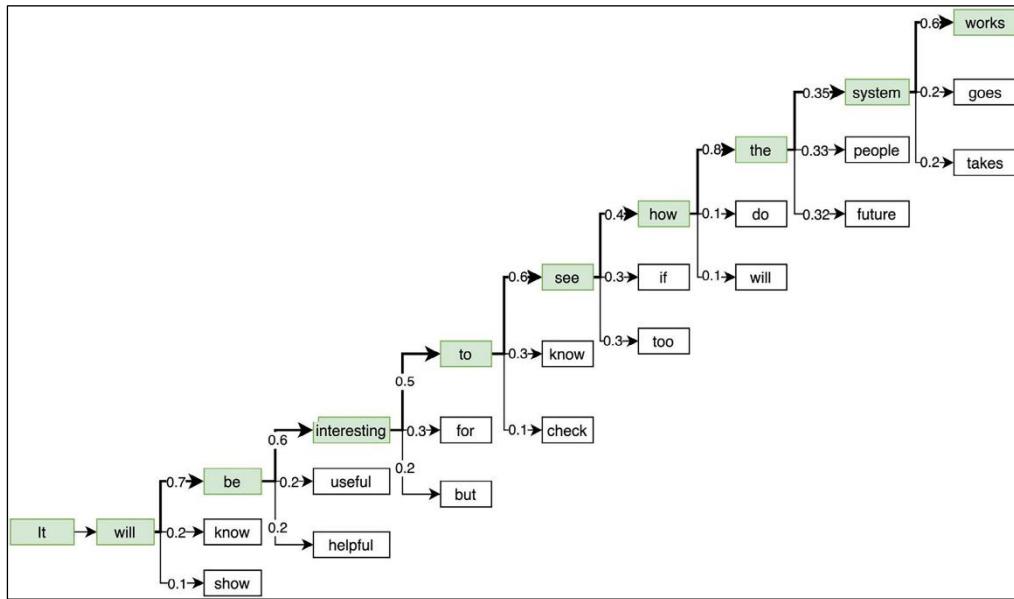


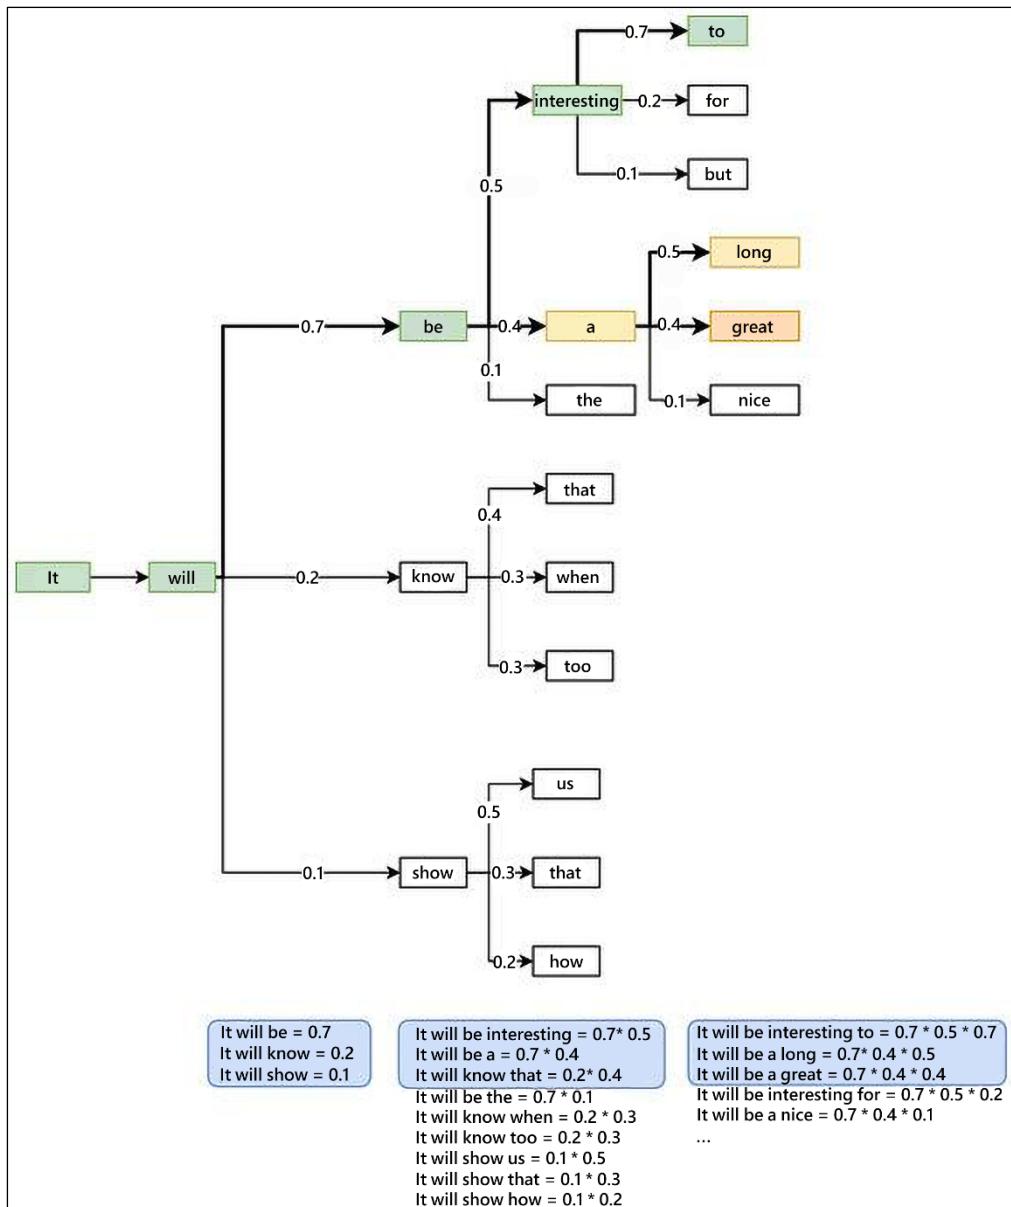


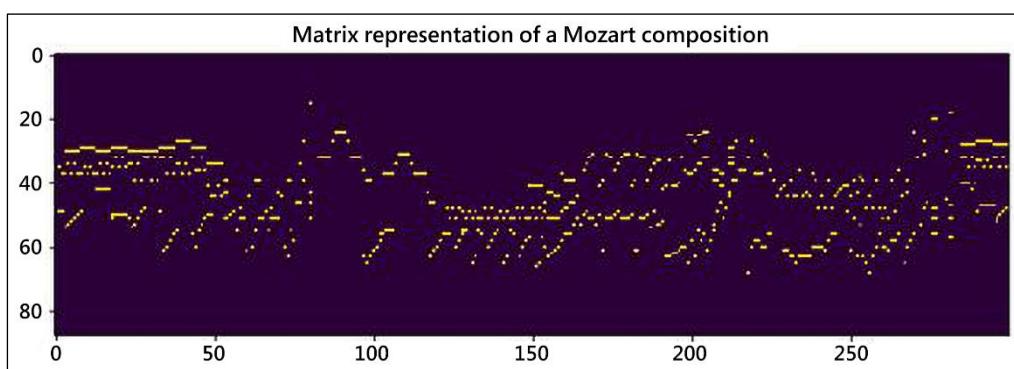


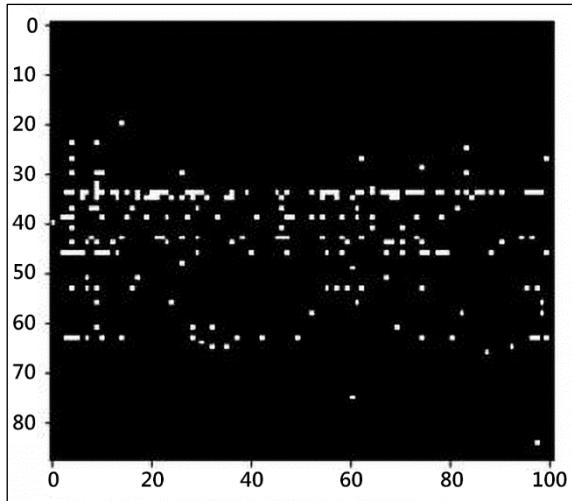


Chapter 7: Music and Text Generation with PyTorch

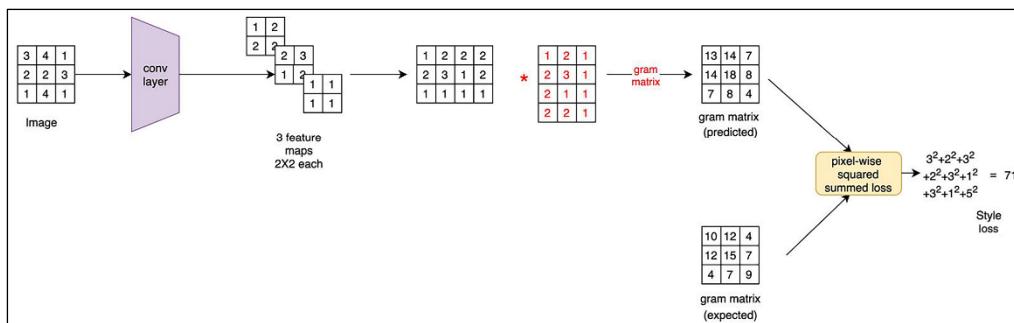
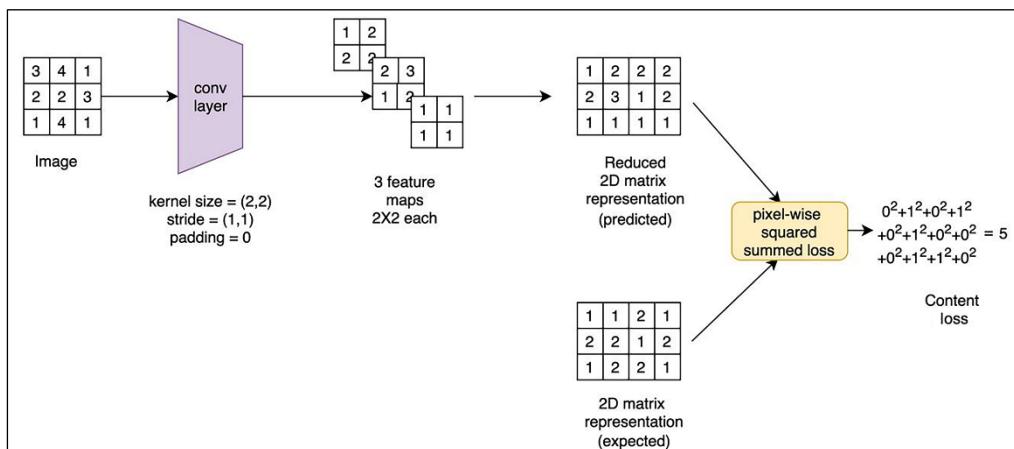
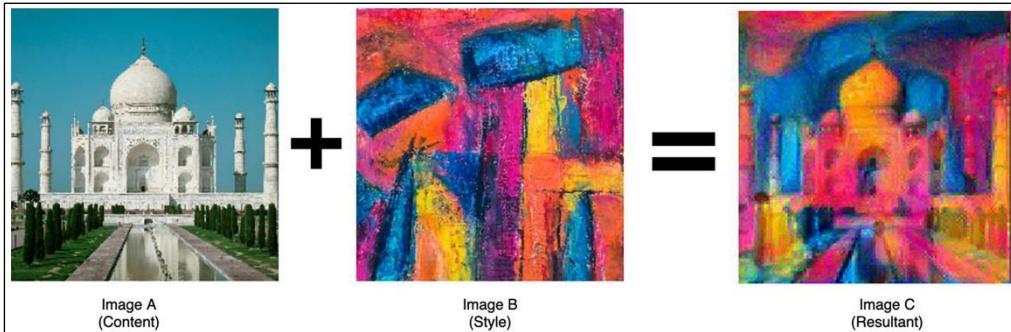


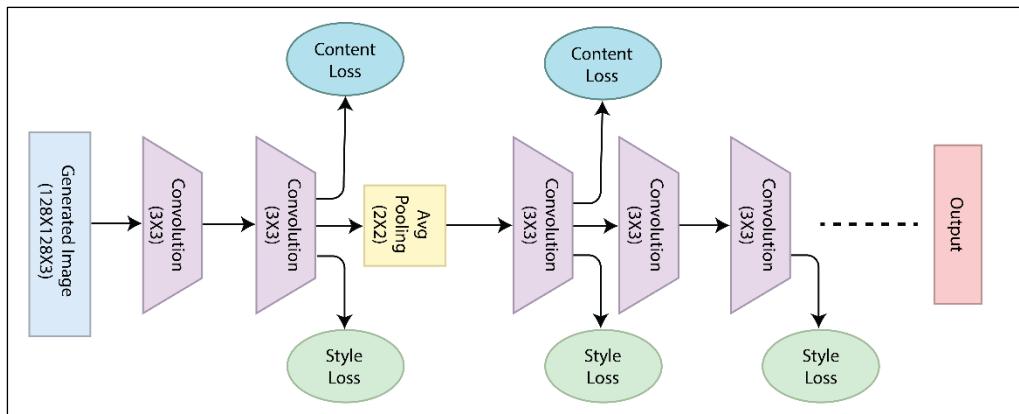


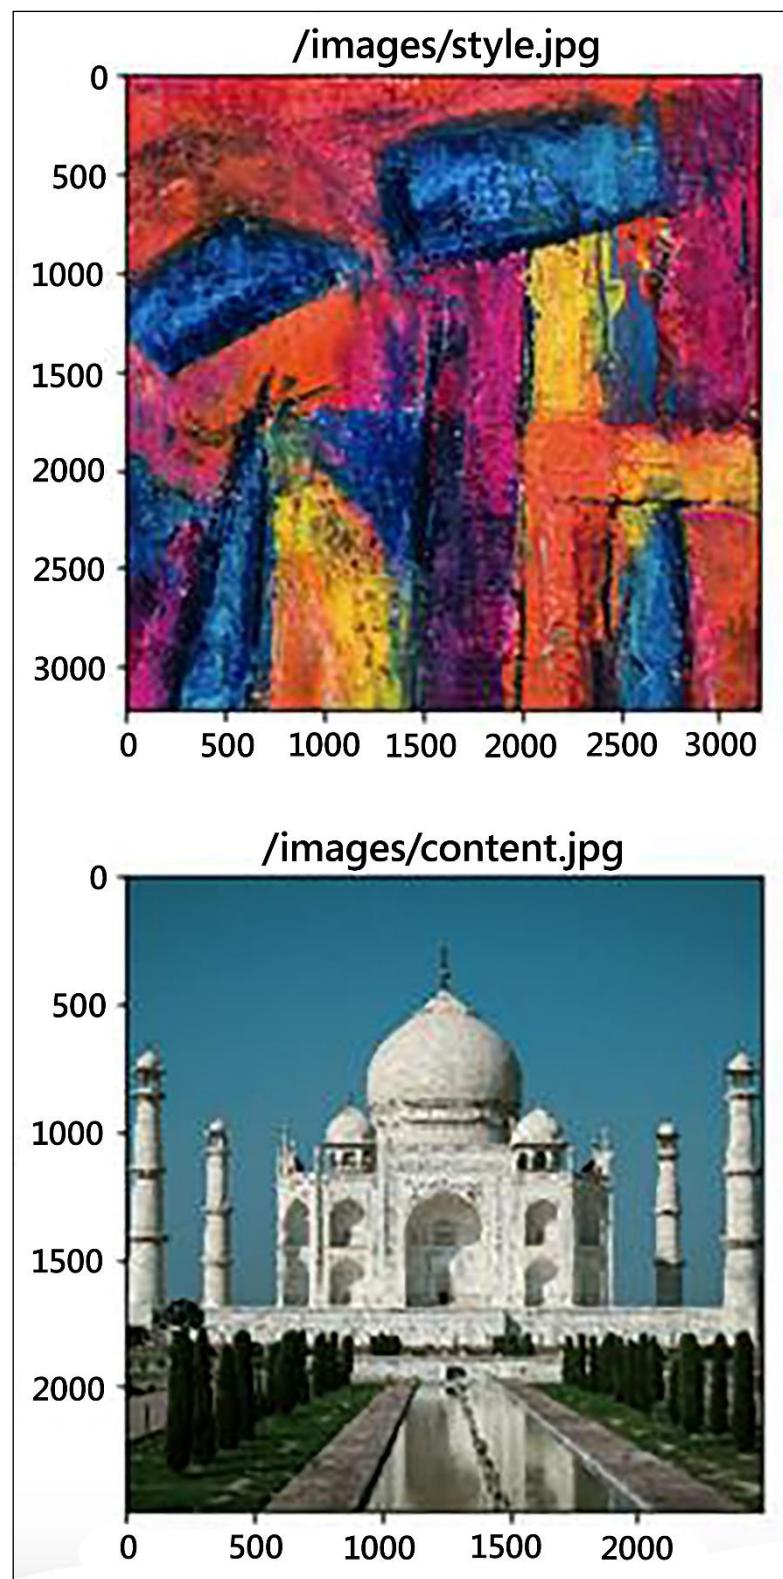


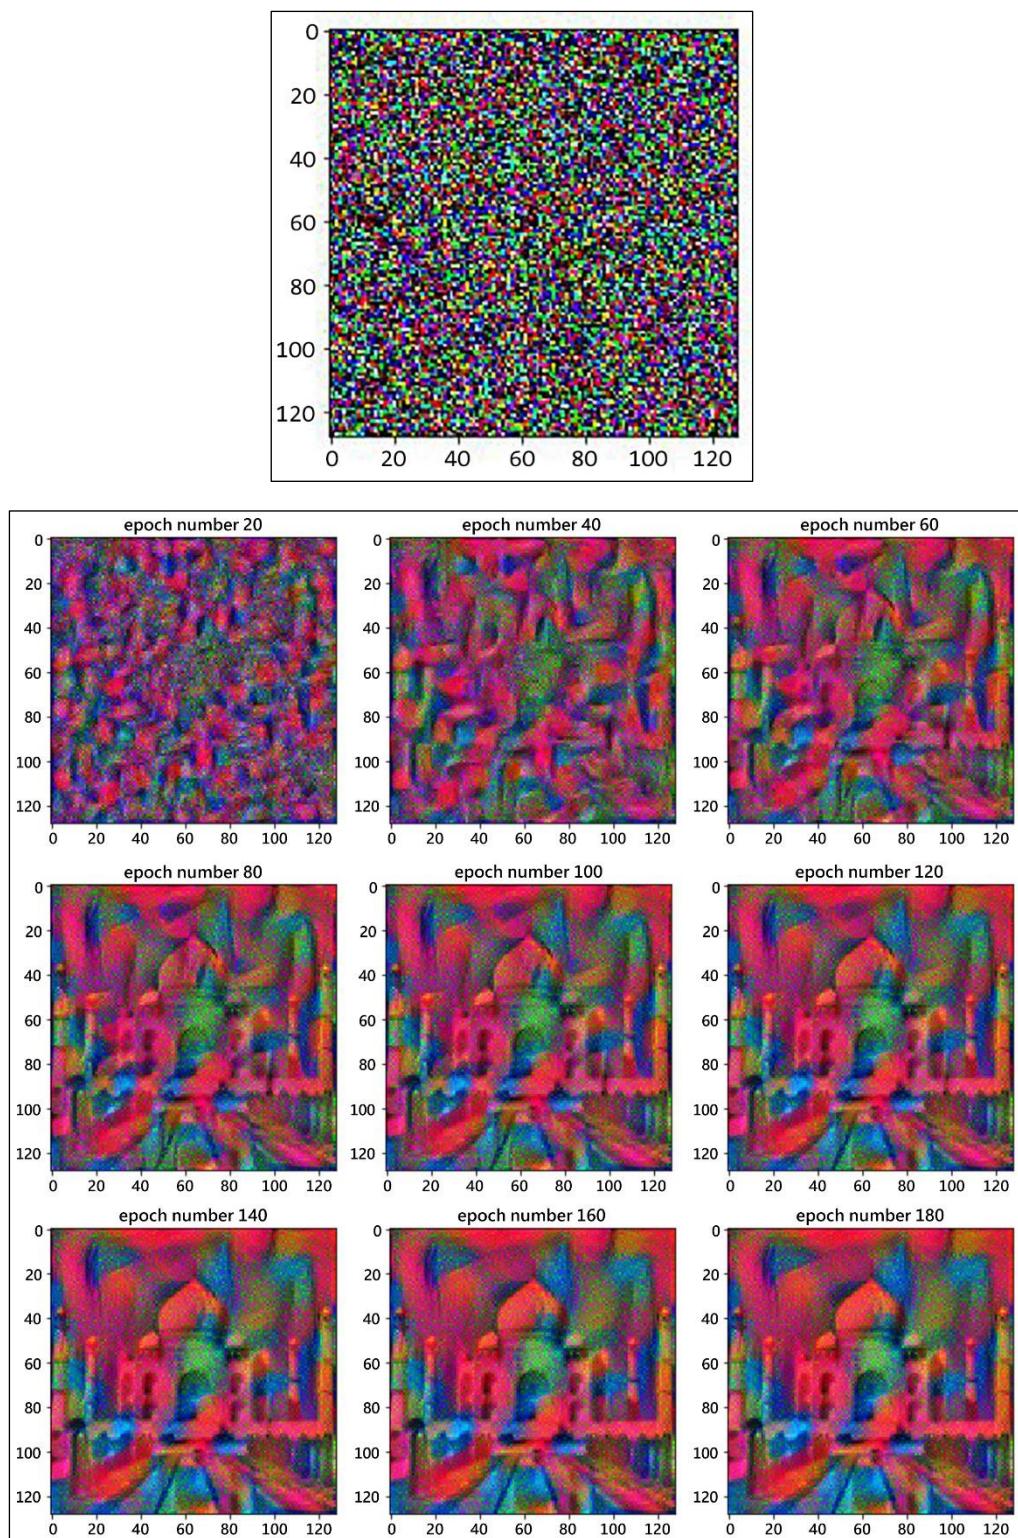


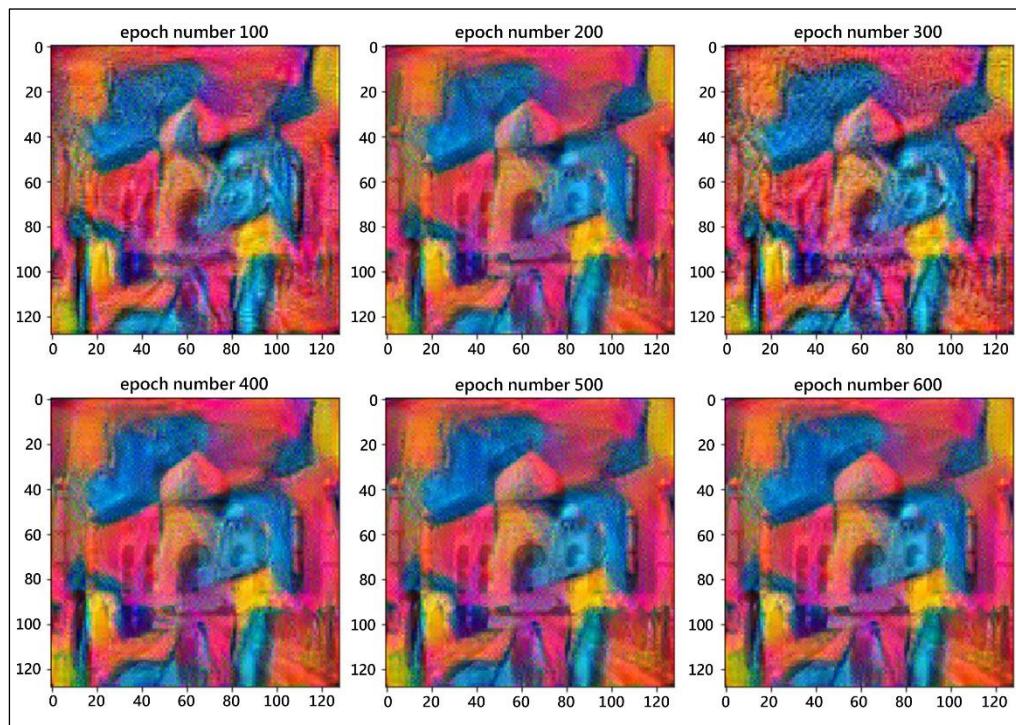
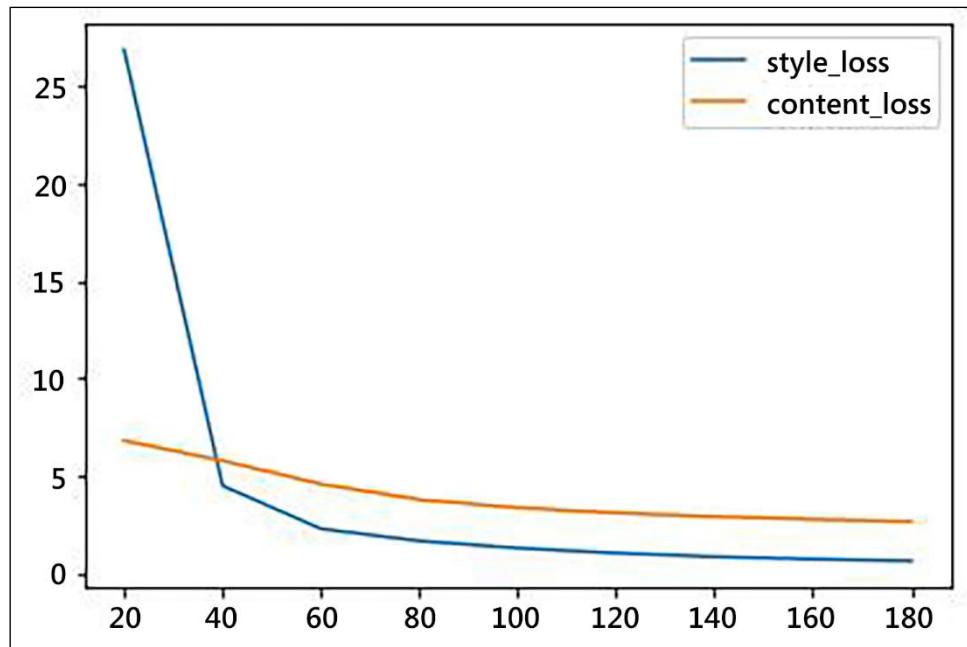
Chapter 8: Neural Style Transfer

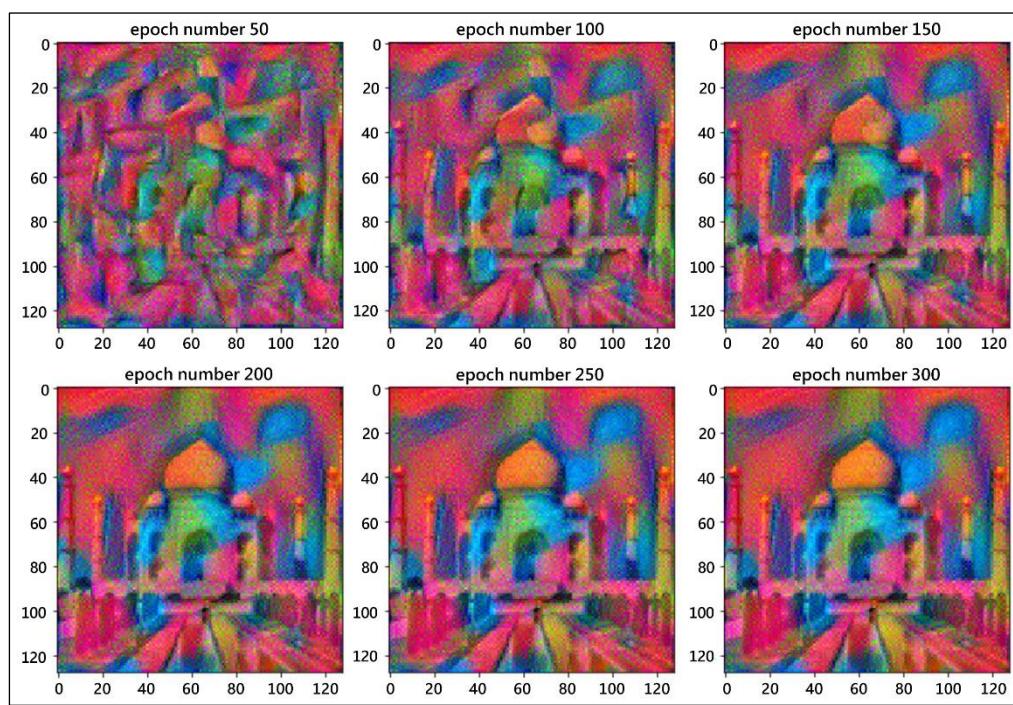
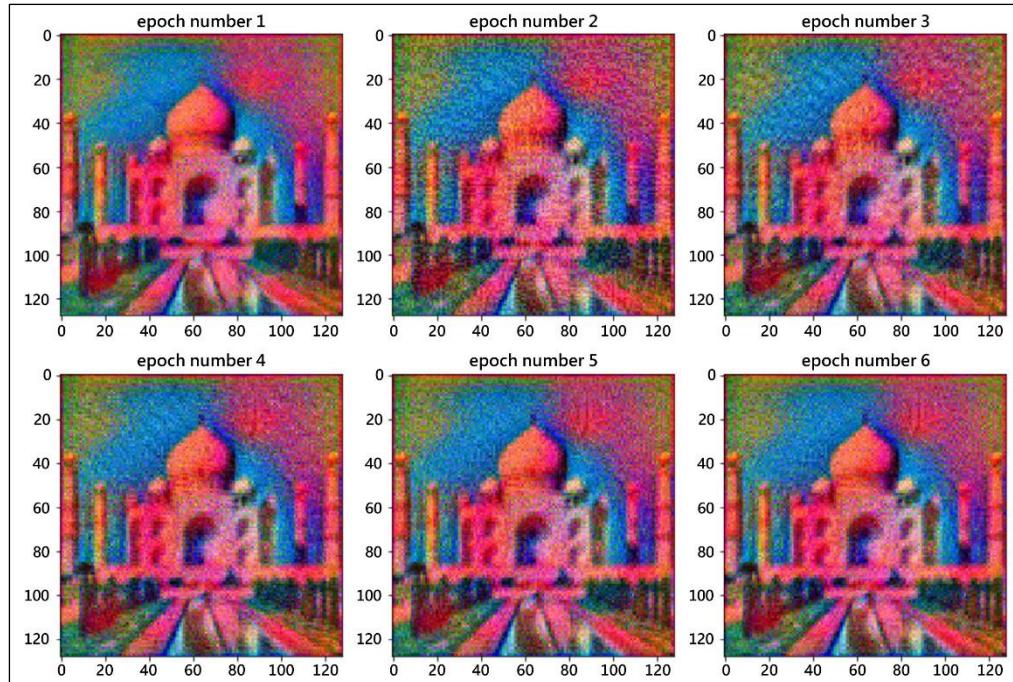


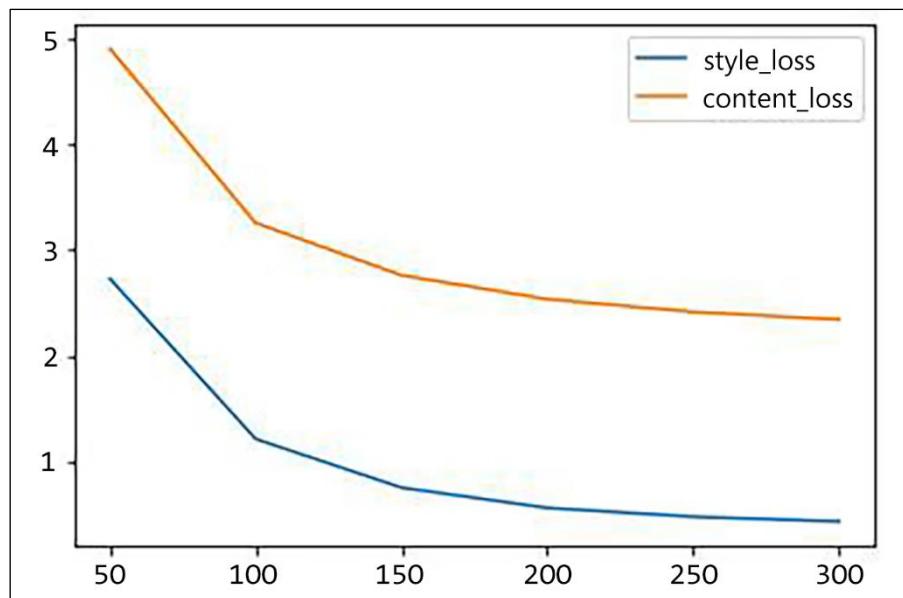




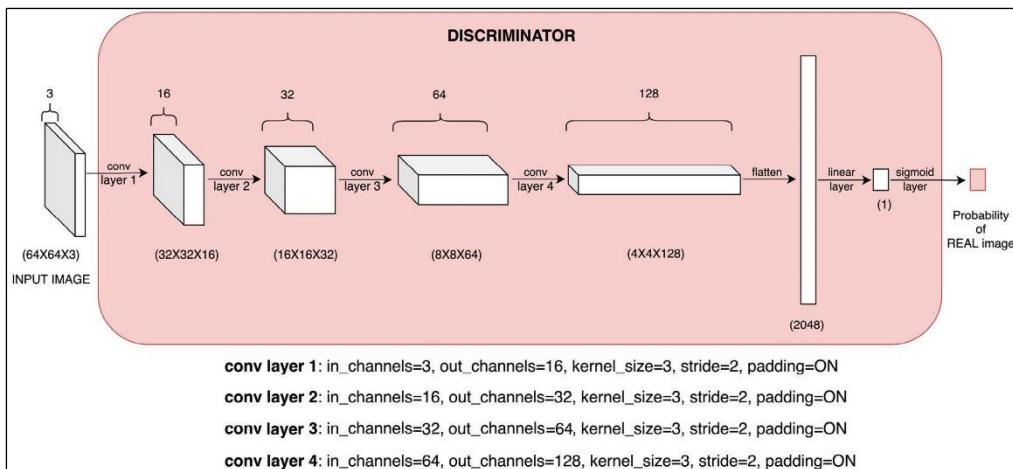
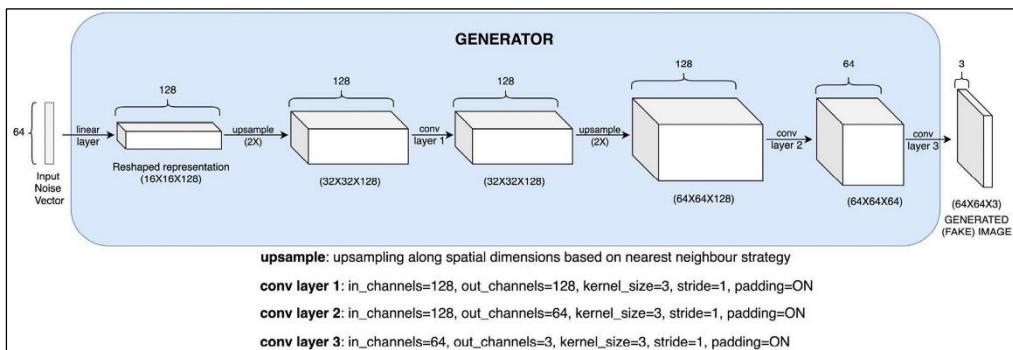
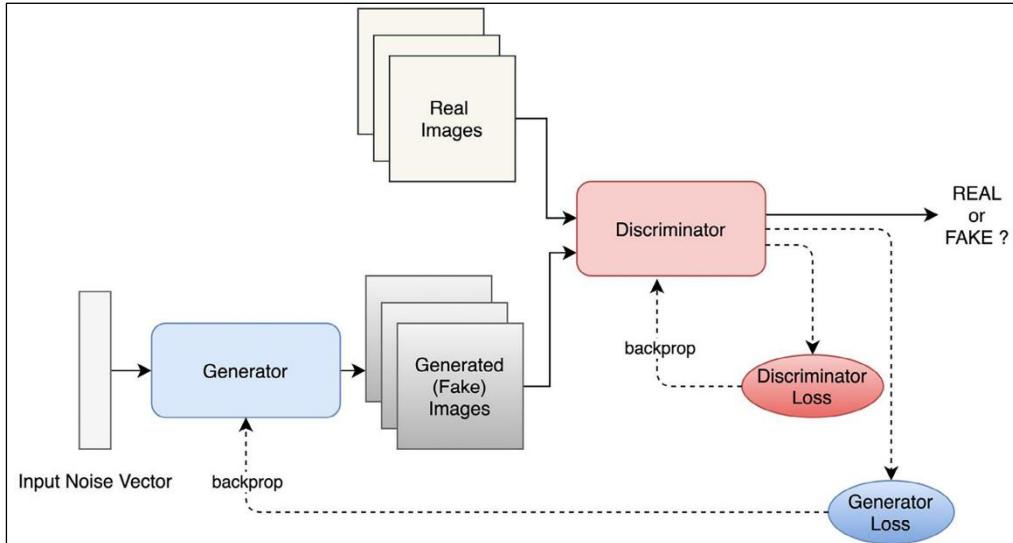






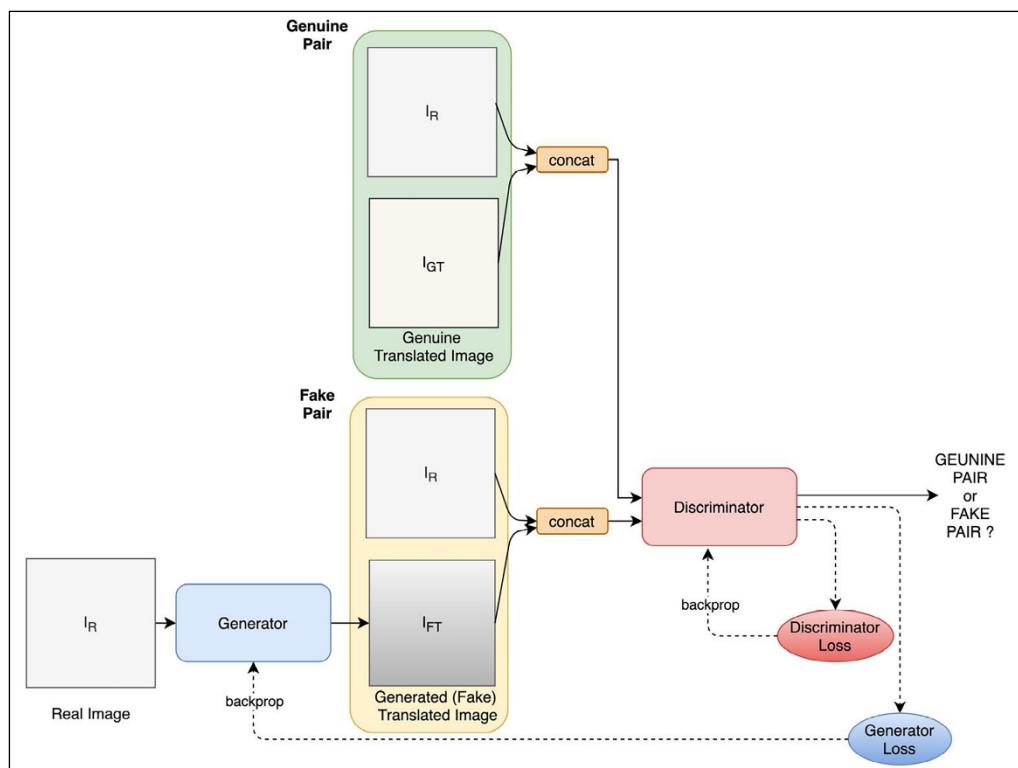


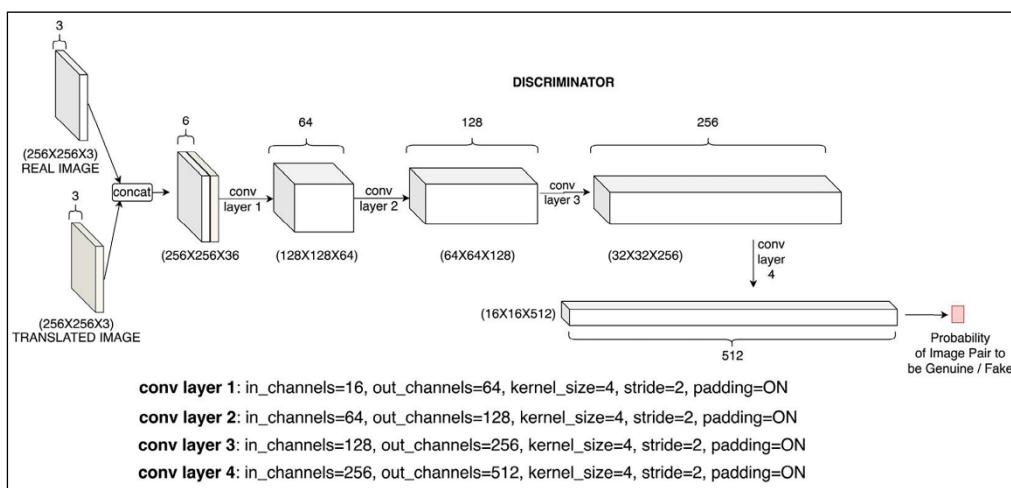
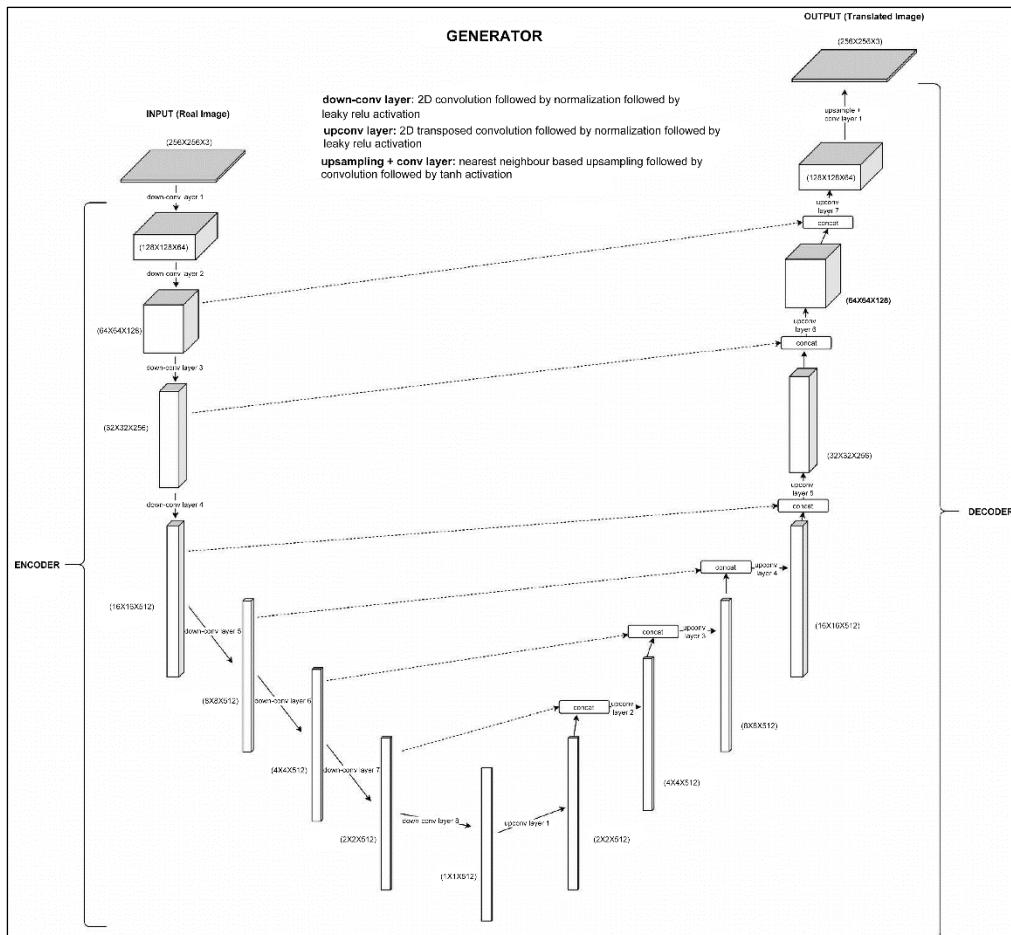
Chapter 9: Deep Convolutional GANs



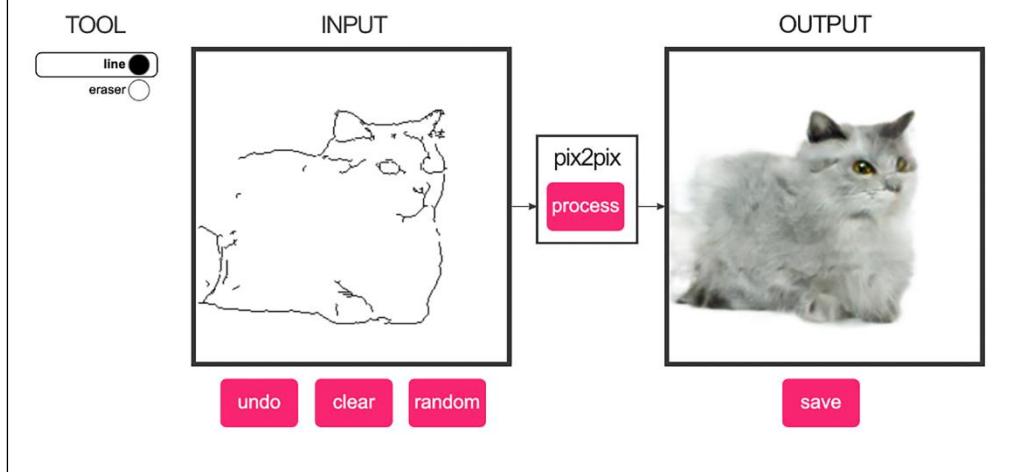
3 4 2 1 9
8 9 1 2 5
6 7 0 1 6
3 7 7 9 4
2 9 3 4 3

Epoch 0	Epoch 1	Epoch 2
	4 4 2 2 4 9 1 3 0 4 0 9 2 7 1 6 7 7 4 8 1 9 0 8 7	6 3 3 4 9 0 1 9 2 3 2 1 4 3 8 5 4 6 5 7 8 2 4 0 1
Epoch 3	Epoch 4	Epoch 5
1 9 4 8 9 0 3 5 0 9 8 1 9 8 6 4 0 3 3 7 1 8 6 2 1	7 6 8 2 0 7 2 3 8 0 8 3 5 1 6 8 7 9 5 8 9 3 5 0 7	2 0 2 1 8 3 3 1 2 3 2 8 0 4 3 4 5 0 9 4 4 2 2 0 4





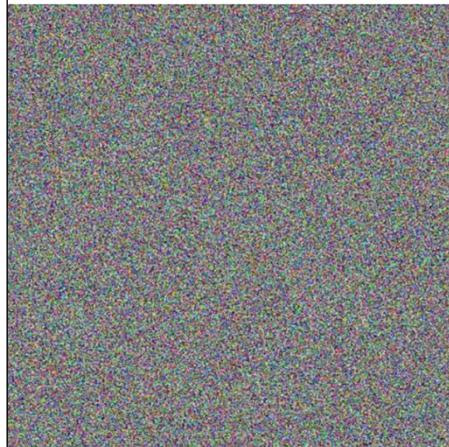
edges2cats



Chapter 10: Image Generation Using Diffusion

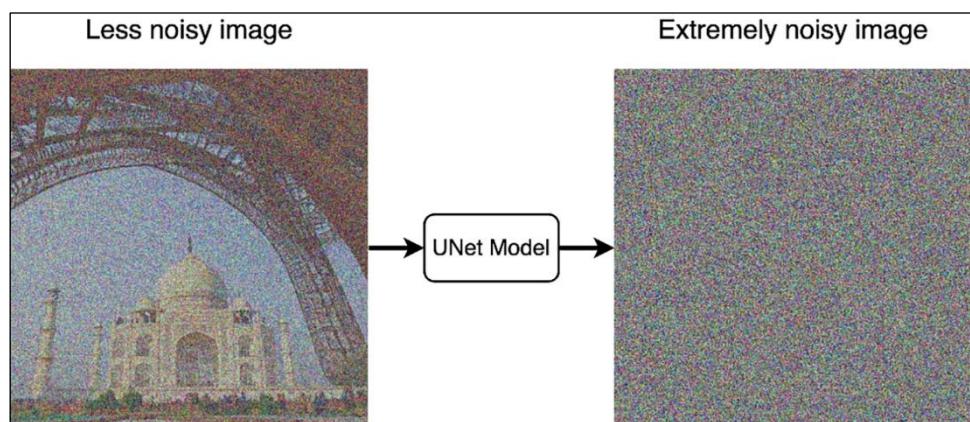
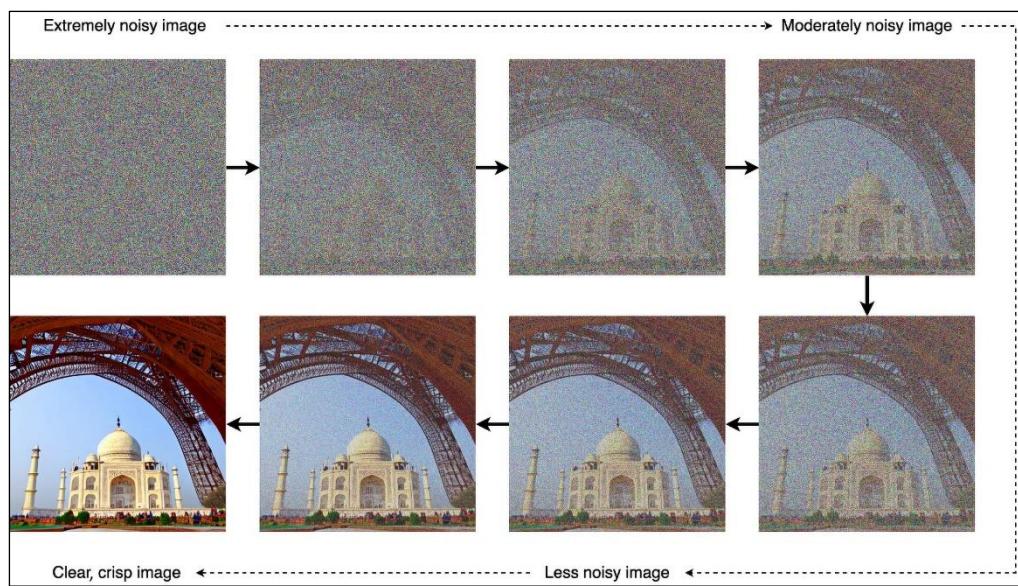


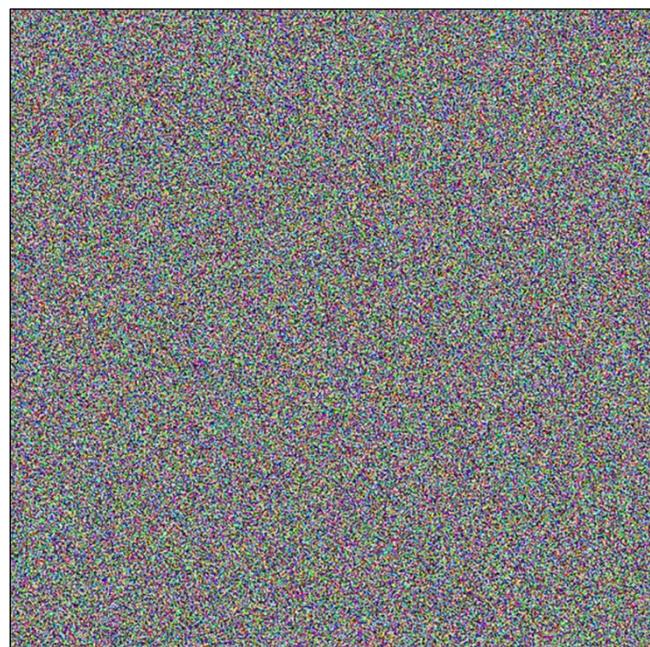
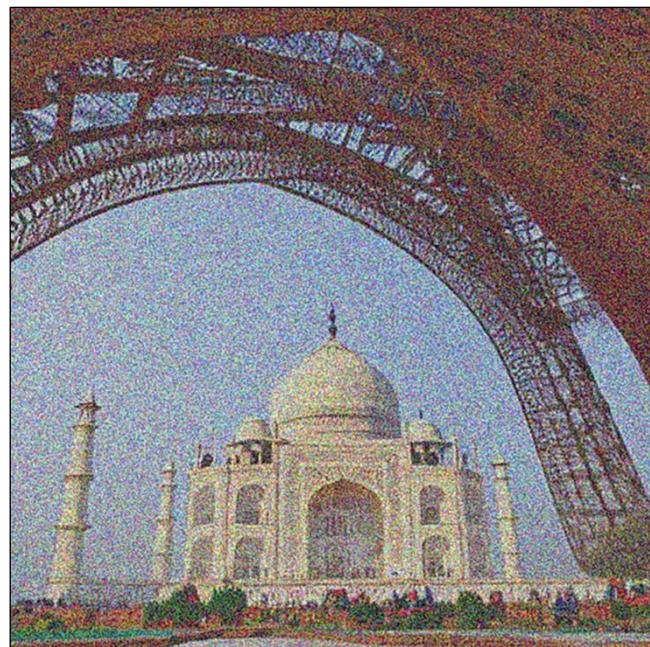
Extremely noisy image

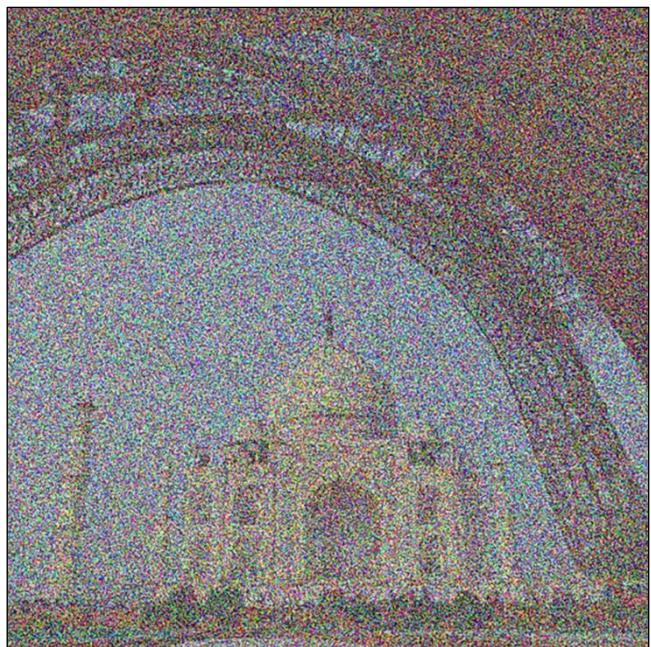


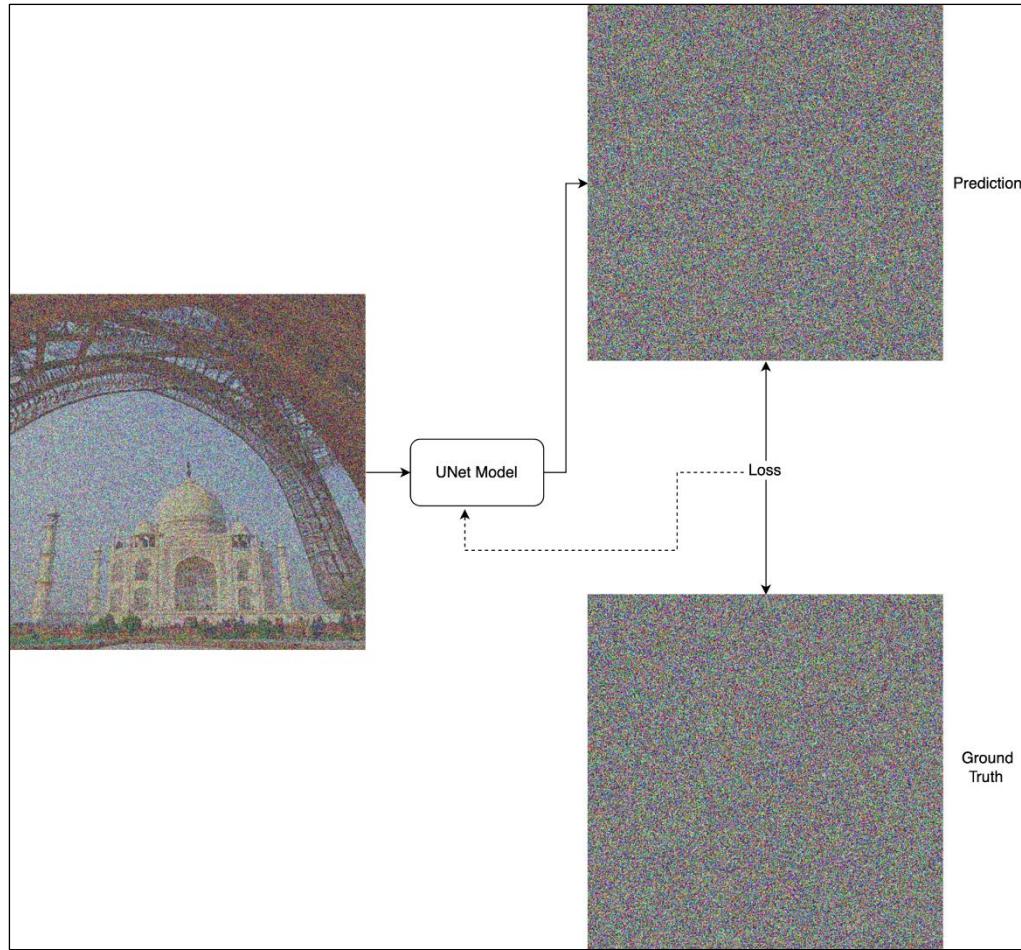
Clear, crisp image

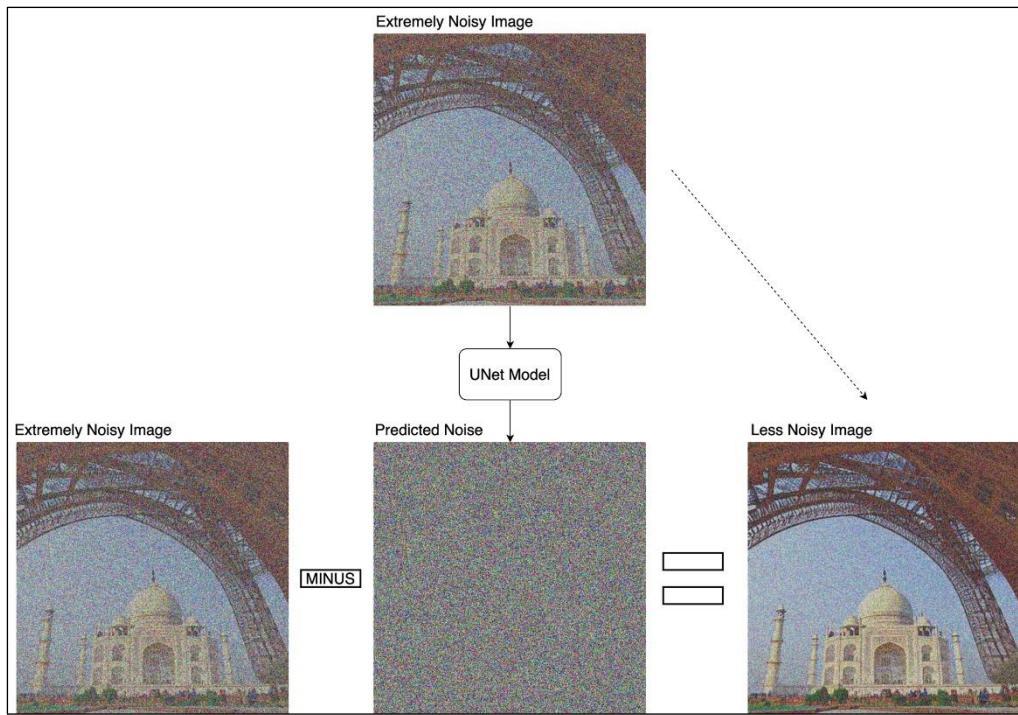


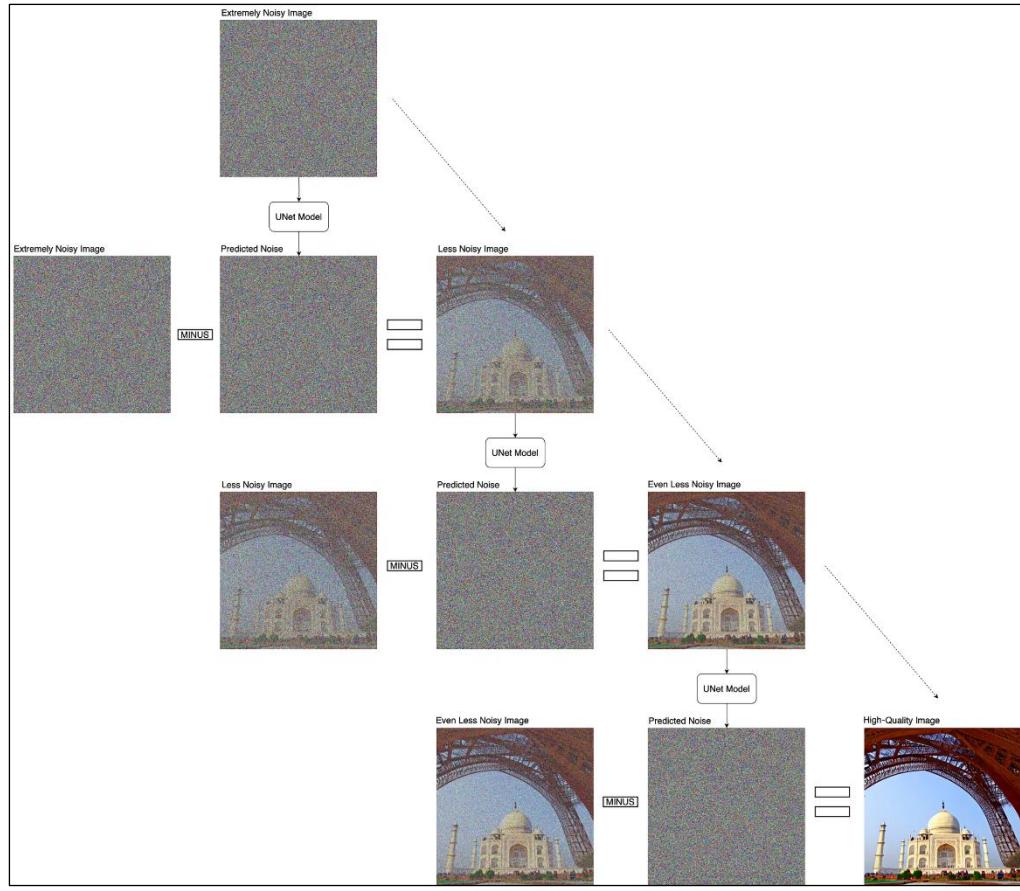












Hugging Face Search models, datasets, users...

Datasets: huggan/selfie2anime like 0

Dataset card Files and versions Community

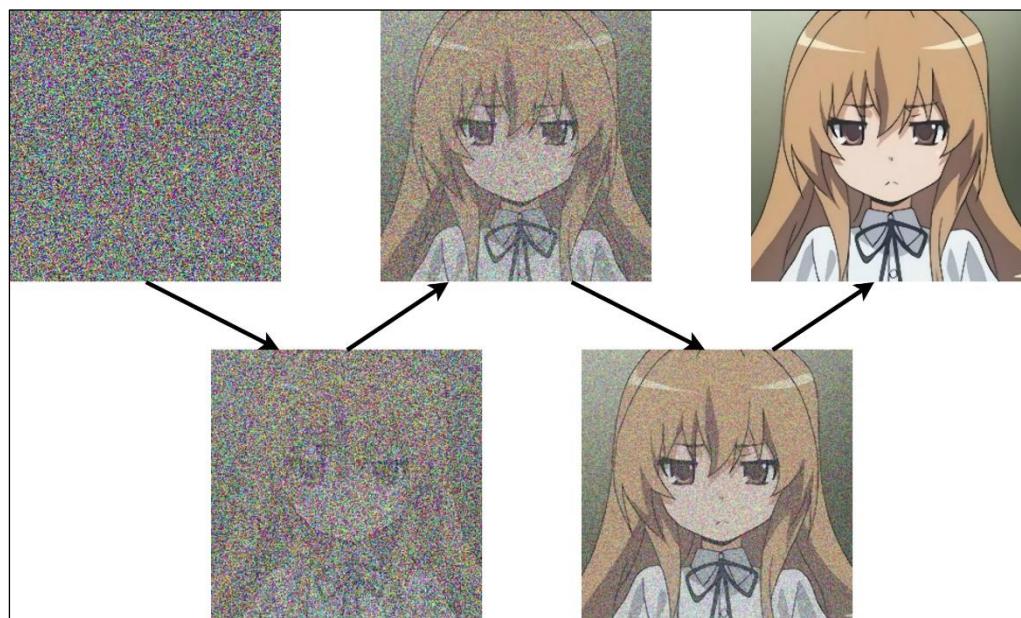
Dataset Viewer Split train (3.4k rows)

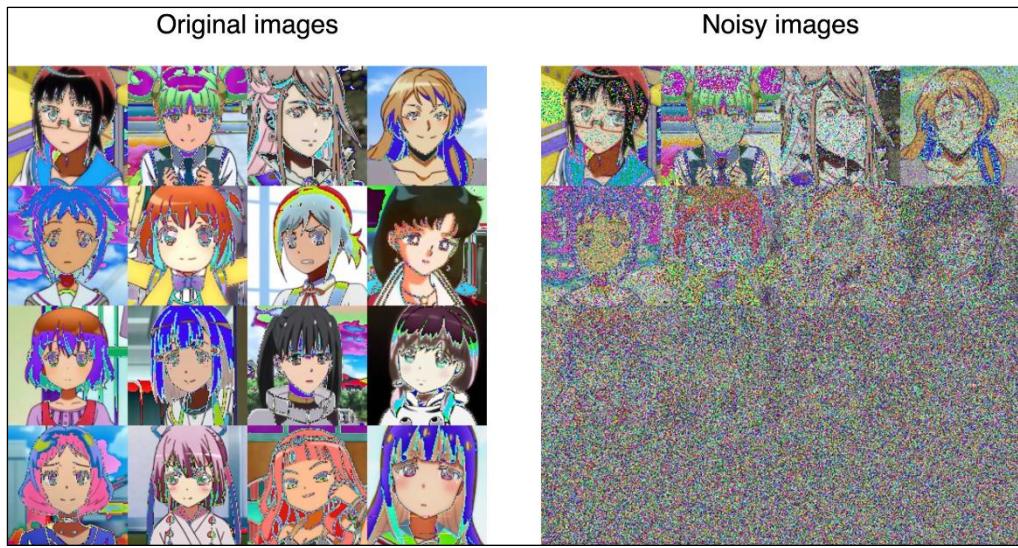
The full dataset viewer is not available (click to read why). Only showing a preview of the rows.

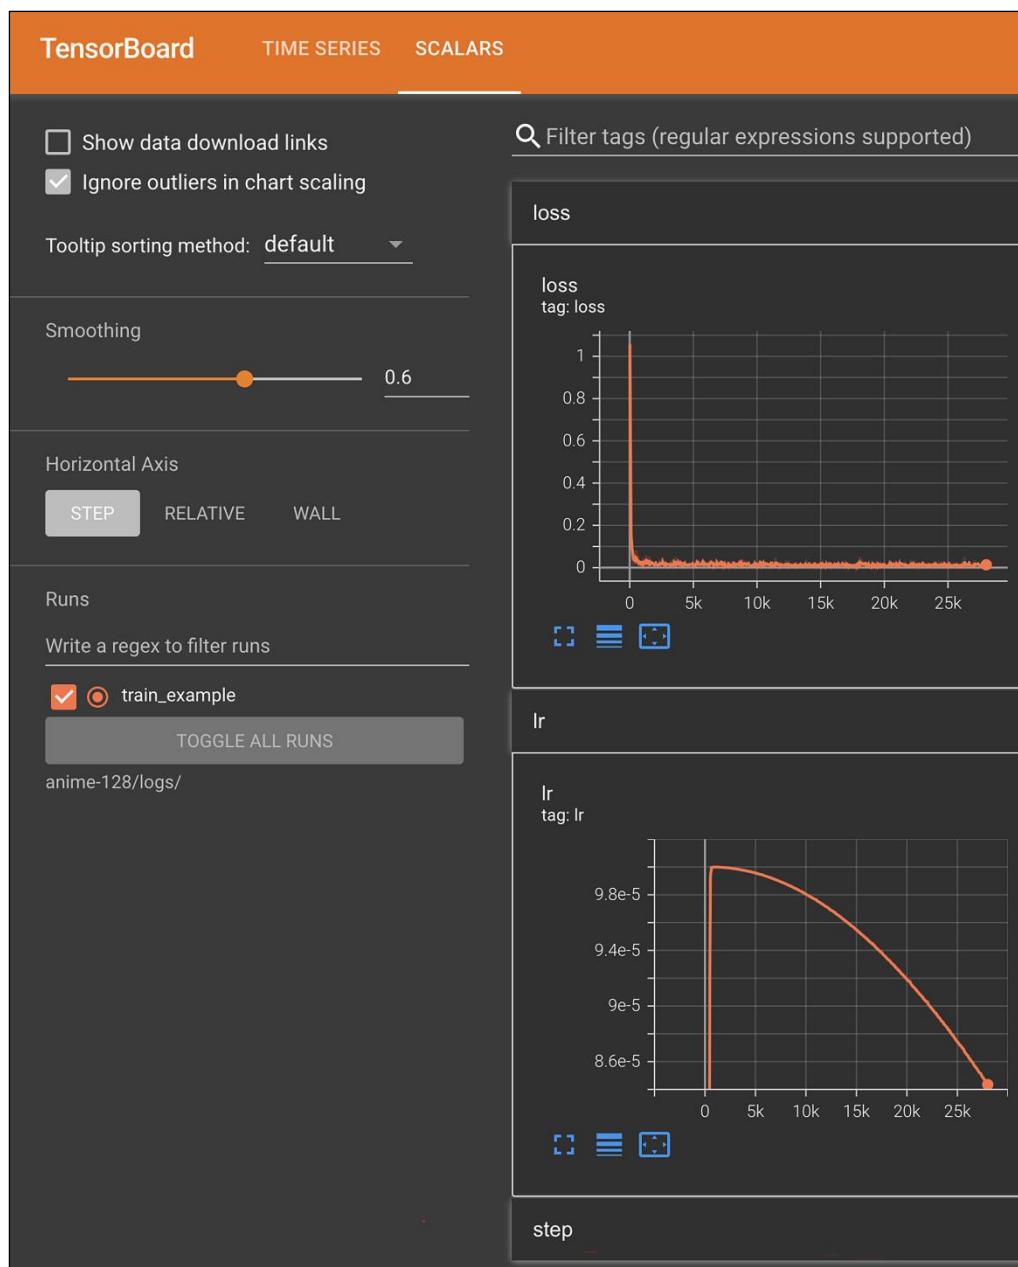
imageA	imageB
image	image

< Previous 1 2 3 ... 34 Next >

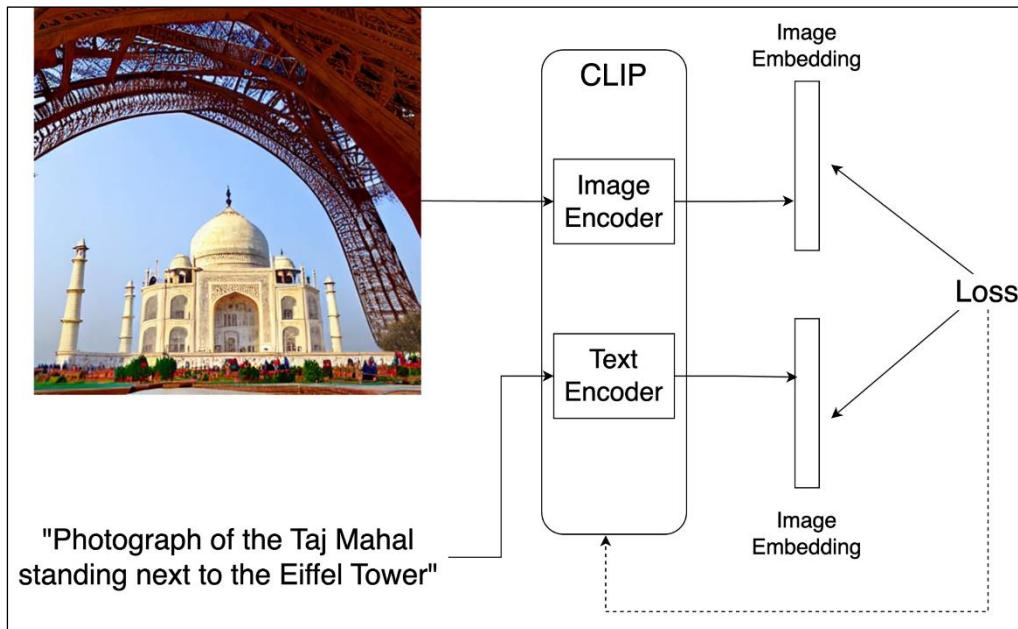
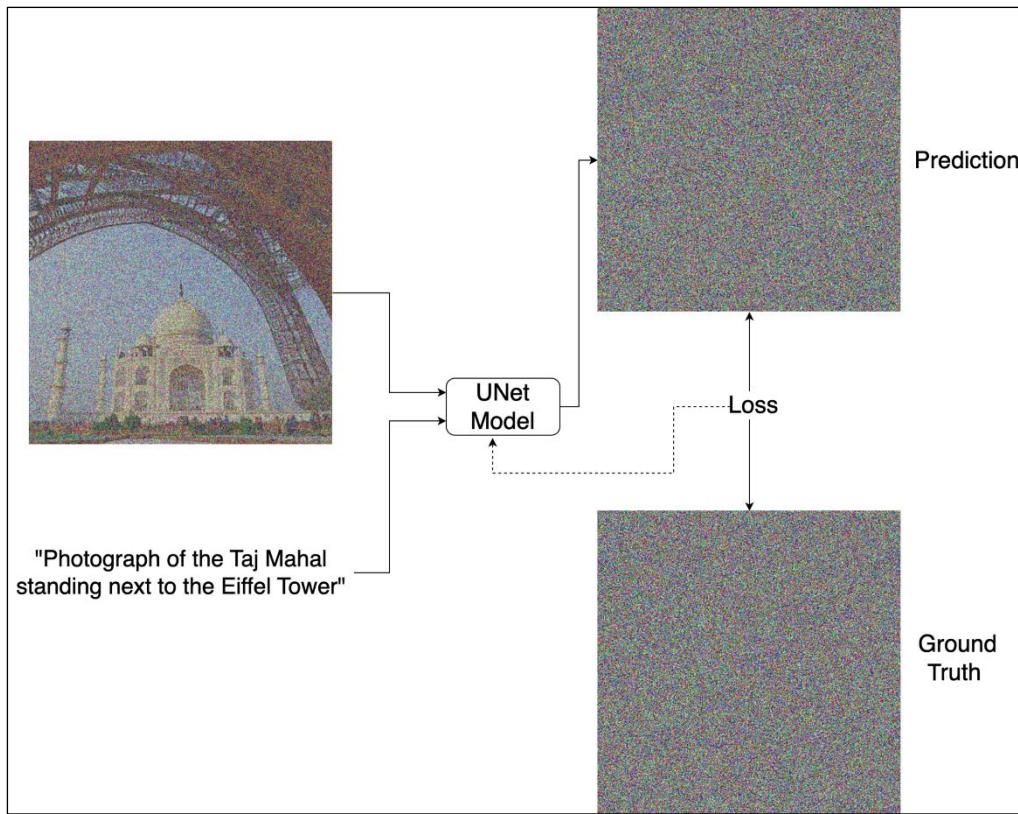


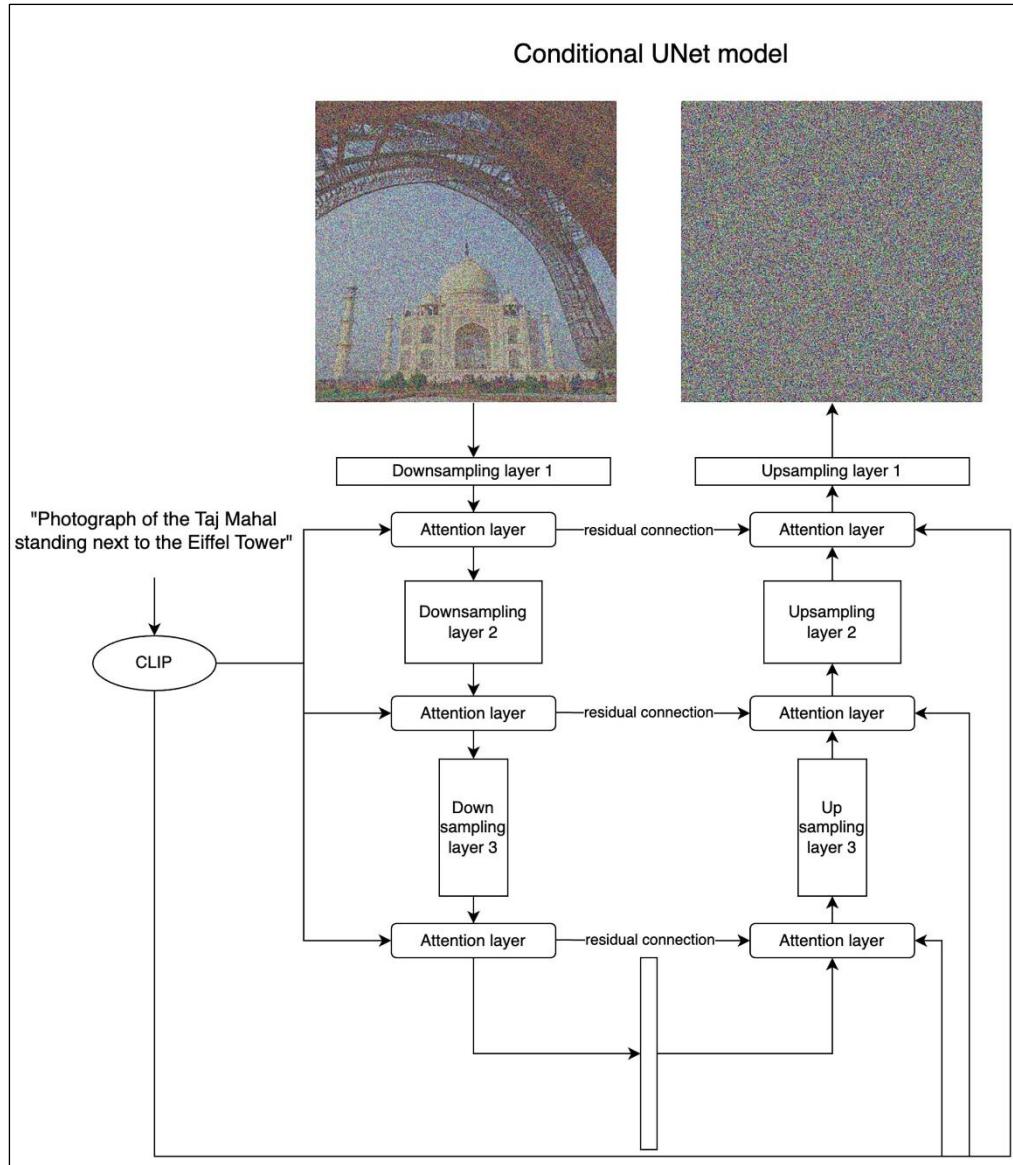


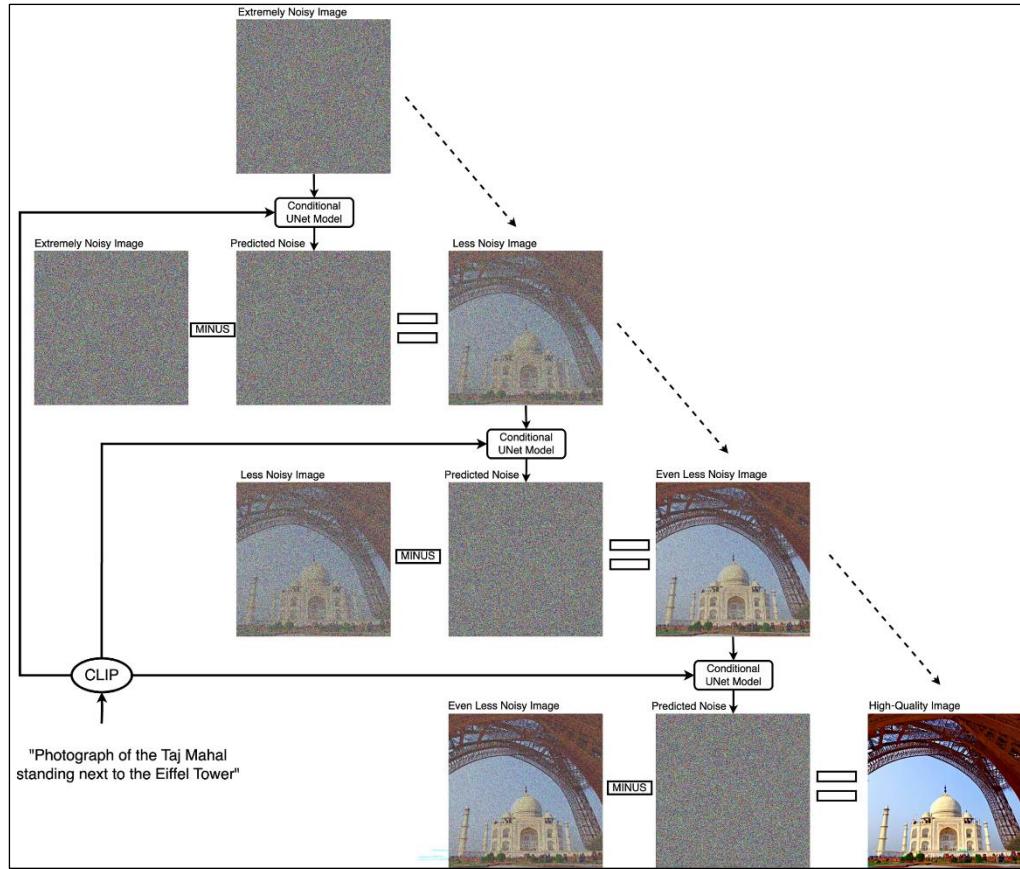


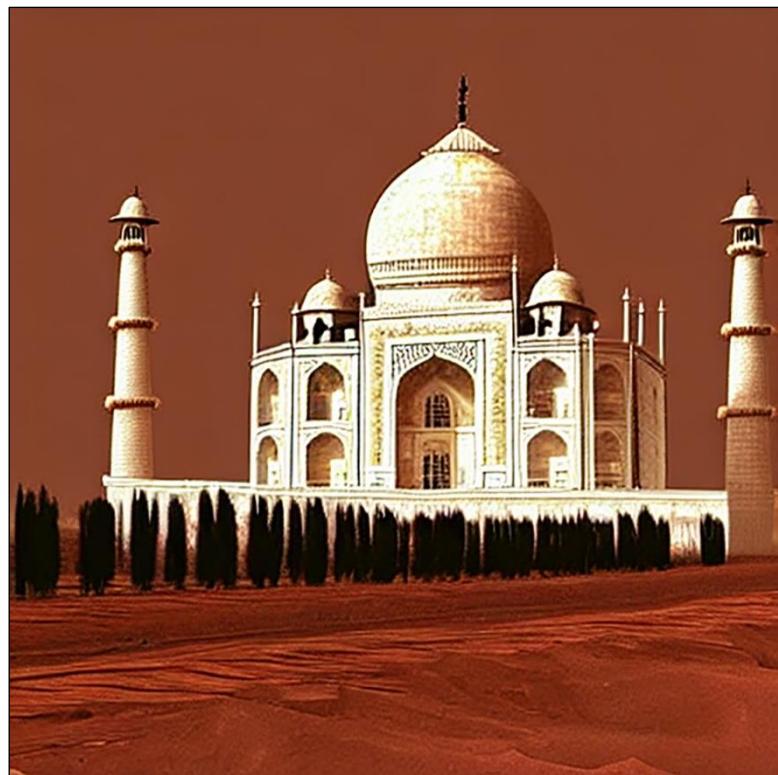




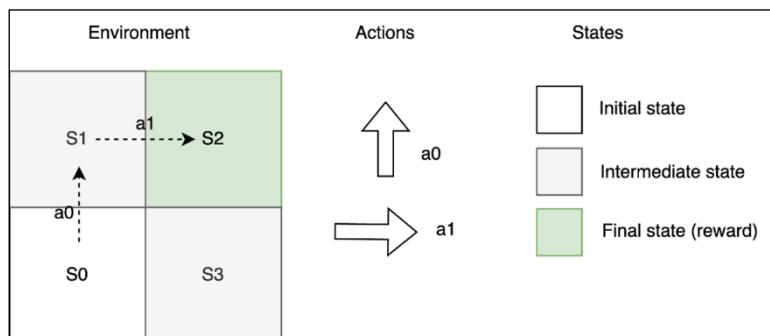
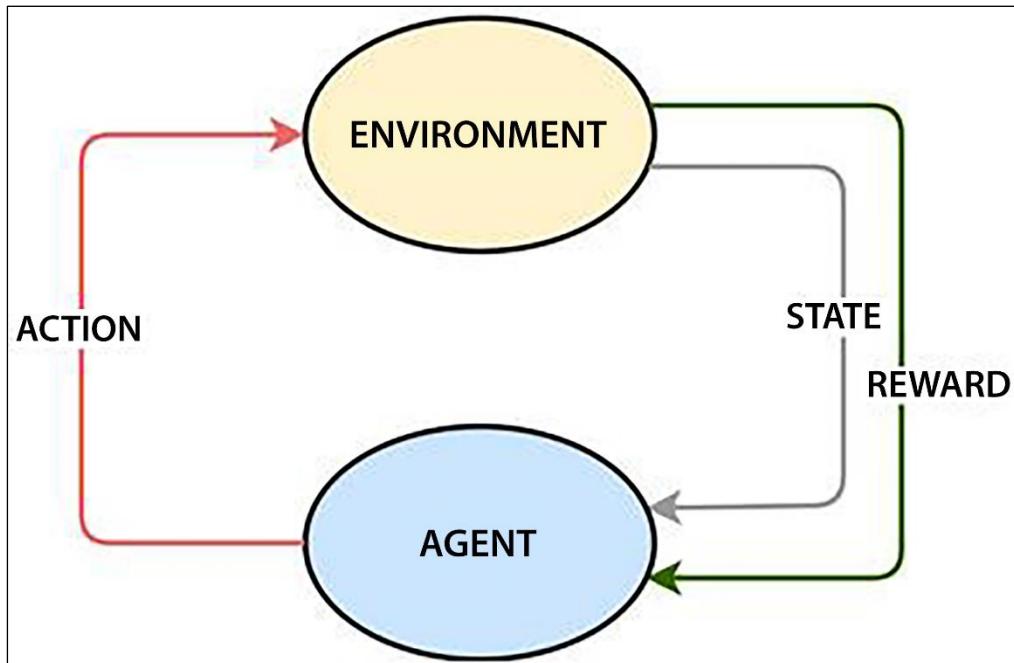


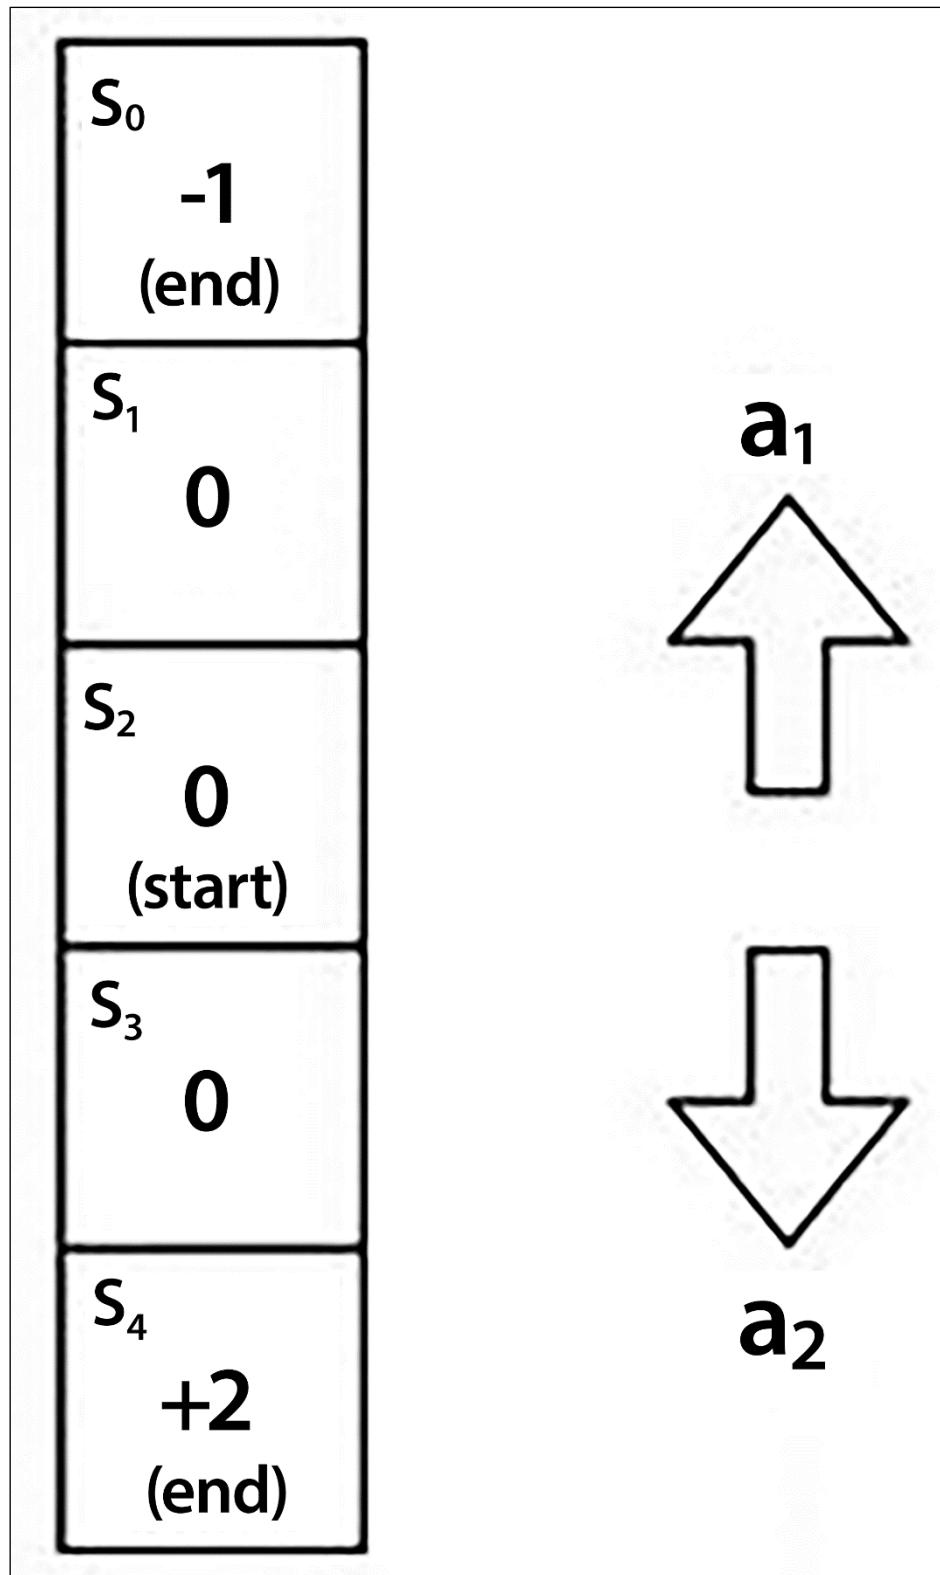


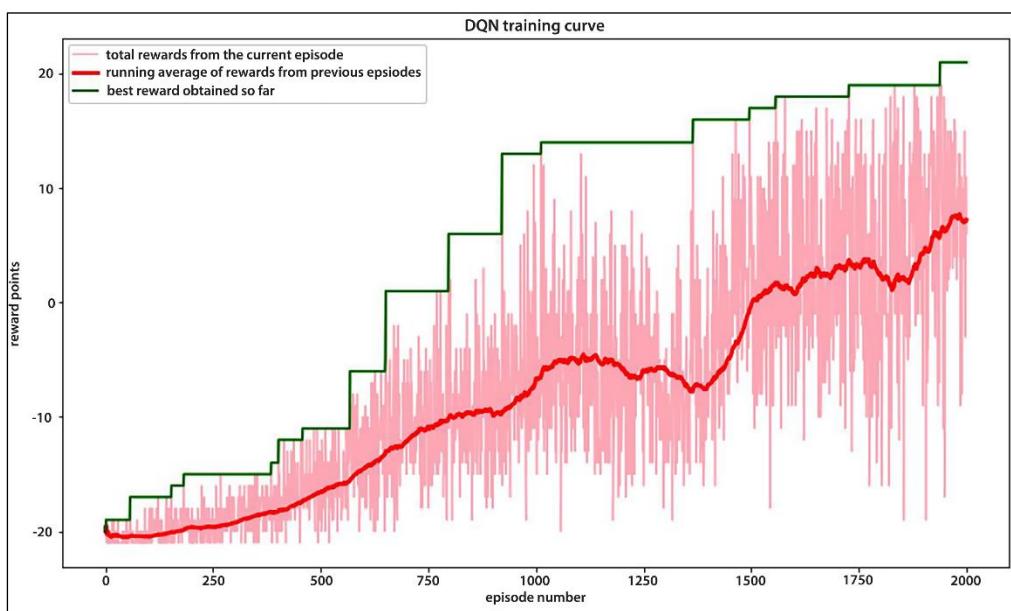
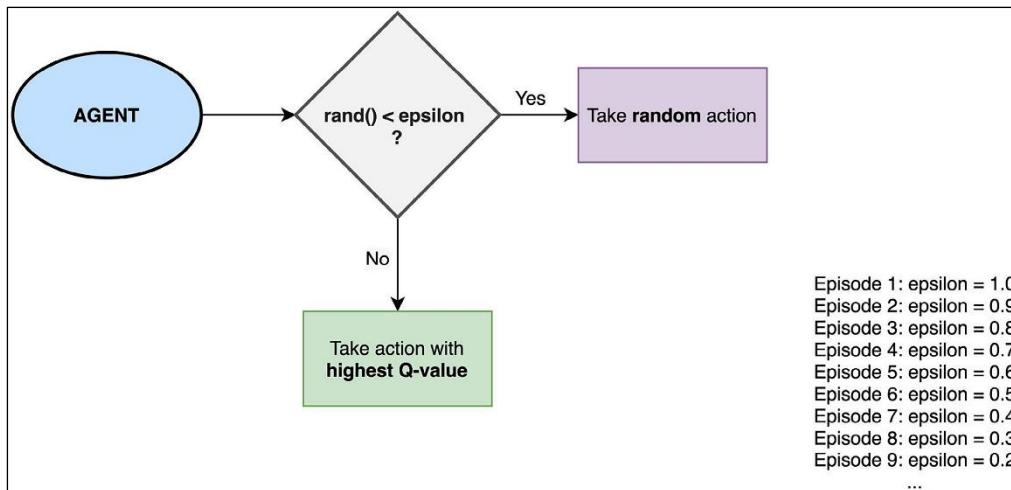


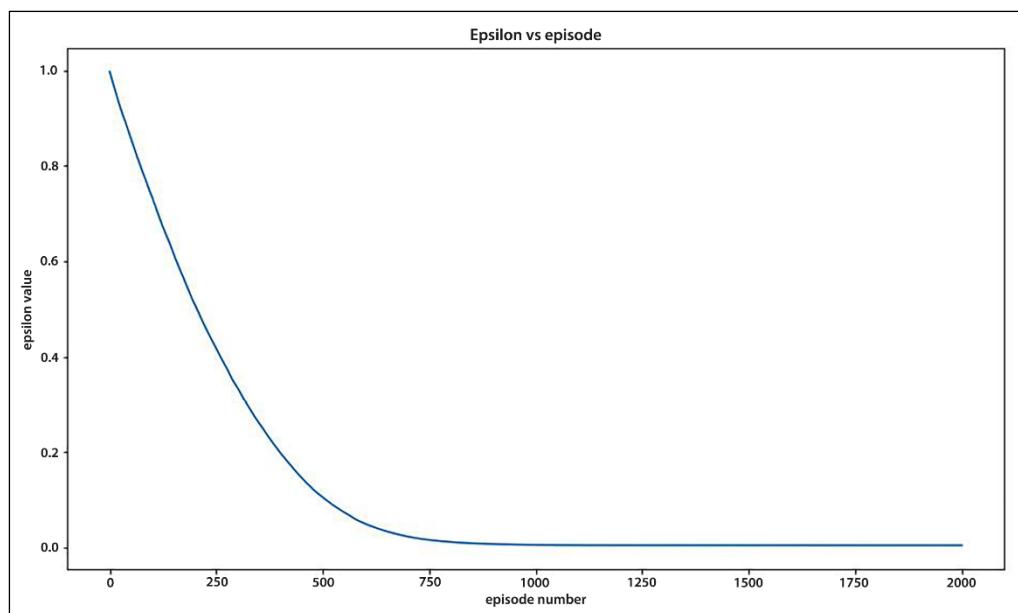


Chapter 11: Deep Reinforcement Learning

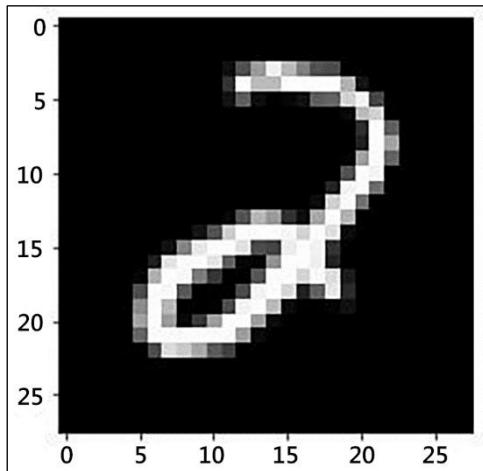




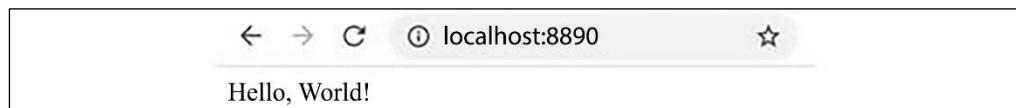




Chapter 13: Operationalizing PyTorch Models into Production



```
* Serving Flask app "example" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://localhost:8898/ (Press CTRL+C to quit)
```



```
* Serving Flask app "server" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://0.0.0.0:8890/ (Press CTRL+C to quit)
```

```

Sending build context to Docker daemon 7.283MB
Step 1/9 : FROM python:3.8-slim
--> 62297c9f4e5c
Step 2/9 : RUN apt-get -q update && apt-get -q install -y wget
--> 6142454652
Step 3/9 : COPY ./server.py .
--> Using cache
--> cb82fb5cb2e5
Step 4/9 : COPY ./requirements.txt .
--> Using cache
--> faa39c98f044
Step 5/9 : RUN wget -q https://github.com/PacktPublishing/Mastering-PyTorch/raw/master/Chapter10/convnet.pth
--> Running in 198679553bac
Removing intermediate container 198679553bac
--> 8b4dcd82ccfcfe
Step 6/9 : RUN wget -q https://github.com/PacktPublishing/Mastering-PyTorch/raw/master/Chapter10/digit_image.jpg
--> Running in 34836205c03d
Removing intermediate container 34836205c03d
--> 1e54479204
Step 7/9 : RUN pip install -r requirements.txt
--> Running in e4dbd692812b
Collecting torch==1.5.0
  Downloading torch-1.5.0-cp38-cp38-manylinux1_x86_64.whl (752.0 MB)
Collecting torchvision==0.6.0
  Downloading torchvision-0.6.0-cp38-cp38-manylinux1_x86_64.whl (6.6 MB)
Collecting Pillow==6.2.2
  Downloading Pillow-6.2.2-cp38-cp38-manylinux1_x86_64.whl (2.1 MB)
Collecting numpy==1.19.1
  Downloading numpy-1.19.1-py2.py3-none-any.whl (94 kB)
Collecting numpy_
  Downloading numpy-1.19.5-cp38-cp38-manylinux2_64_x86_64.whl (14.9 MB)
Collecting future
  Downloading future-0.18.2.tar.gz (829 kB)
Collecting Werkzeug==0.15
  Downloading Werkzeug-0.15.0-py2.py3-none-any.whl (298 kB)
Collecting itsdangerous<0.24
  Downloading itsdangerous-0.1.0-py2.py3-none-any.whl (16 kB)
Collecting MarkupSafe<1.1.1
  Downloading MarkupSafe-1.1.1-py2.py3-none-any.whl (125 kB)
Collecting click==5.1
  Downloading click-7.1.2-py2.py3-none-any.whl (82 kB)
Collecting MarkupSafe<1.1.1-cp38-cp38-manylinux1_x86_64.whl (32 kB)
Building wheels for collected packages: future
  Building wheel for future (setup.py): started
  Building wheel for future (setup.py): finished with status 'done'
    Created wheel for future: filename=future-0.18.2-py3-none-any.whl size=401050 sha256=723027831b4159f39497dee6280cf48539056b1b80d80fbff113330e7ed5af0d
  Stored in directory: /root/.cache/pip/wheels/8e/70/2b/3d6ccde515f65f245da85482e2c7d14b9003bf239e2c4
Successfully built future
Installing collected packages: numpy, future, torch, Pillow, torchvision, Werkzeug, itsdangerous, MarkupSafe, Jinja2, click, Flask
Successfully installed Flask-1.1.1 Jinja2>2.11.2 MarkupSafe-1.1.2 Pillow-6.2.2 Werkzeug-1.0.1 click-7.1.2 future-0.18.2 itsdangerous-1.1.0 numpy-1.19.5 torch-1.5.0 torchvision-0.6.0
Marking /usr/local/bin/python as executable; however, version 20.3.3 is available.
You should consider upgrading via the '/usr/local/bin/python -m pip install --upgrade pip' command.
Removing intermediate container e4dbd692812b
--> 47ef4f4681868
Step 8/9 : USER root
--> 90dc164d1dbc
Removing intermediate container 90dc164d1dbc
--> e448fbab0982
Step 9/9 : ENTRYPOINT ["python", "server.py"]
--> Running in d1d55f33d9ff
Removing intermediate container d1d55f33d9ff
--> 2de7d75ae6b
Successfully built 2de7d75ae6b
Successfully tagged digit_recognizer:latest

```

```

* Serving Flask app "server" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://0.0.0.0:8890/ (Press CTRL+C to quit)

```

```

-- The C compiler identification is AppleClang 13.1.6.13166021
-- The CXX compiler identification is AppleClang 13.1.6.13166021
-- Detecting C compiler ABI info
-- Detecting C compiler ABI info - done
-- Check for working C compiler: /Library/Developer/CommandLineTools/usr/bin/cc - skipped
-- Detecting C compile features
-- Detecting C compile features - done
-- Detecting CXX compiler ABI info
-- Detecting CXX compiler ABI info - done
-- Check for working CXX compiler: /Library/Developer/CommandLineTools/usr/bin/c++ - skipped
-- Detecting CXX compile features
-- Detecting CXX compile features - done
-- Performing Test CMAKE_HAVE_LIBC_PTHREAD
-- Performing Test CMAKE_HAVE_LIBC_PTHREAD - Success
-- Found Threads: TRUE
-- MKL_ARCH: None, set to 'intel_i4' by default
-- MKL_CBLAS: None, set to 'mkl' by default
-- MKL_LINPACK: None, set to 'mkl' by default
-- MKL_INTERFACE_FULL: None, set to 'intel_ilp64' by default
-- MKL_THREADING: None, set to 'intel_thread' by default
-- MKL_MPI: None, set to 'mpich' by default
CMake Warning at /Users/ashish.jha/opt/anaconda3/envs/mastering_pytorch_7_chaps/lib/python3.9/site-packages/torch/share/cmake/Torch/TorchConfig.cmake:22 (message):
  static library kineto_LIBRARY-NOTFOUND not found.
Call Stack (most recent call first):
  /Users/ashish.jha/opt/anaconda3/envs/mastering_pytorch_7_chaps/lib/python3.9/site-packages/torch/share/cmake/Torch/TorchConfig.cmake:127 (append_torchlib_if_found)
  CMakeLists.txt:3 (find_package)

-- Found Torch: /Users/ashish.jha/opt/anaconda3/envs/mastering_pytorch_7_chaps/lib/python3.9/site-packages/torch/lib/libtorch.dylib
-- Found OpenCV: /Users/ashish.jha/code/personal/Mastering-PyTorch/Chapter13/cpp_convnet/build_opencv (found version "4.6.0")
-- Configuring done
-- Generating done
-- Build files have been written to: /Users/ashish.jha/code/personal/Mastering-PyTorch/Chapter13/cpp_convnet/build

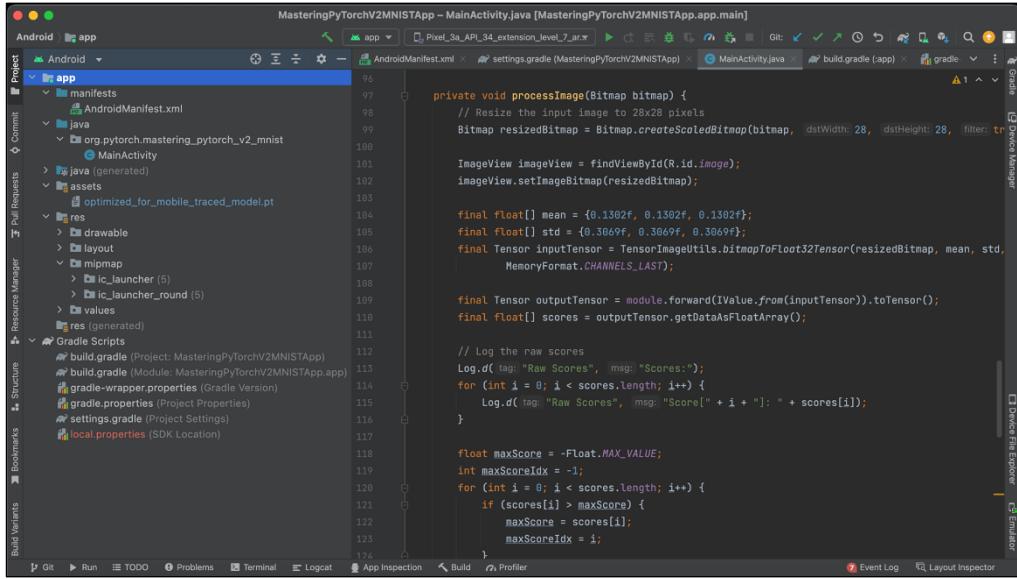
```

```

Scanning dependencies of target cpp_convnet
[ 50%] Building CXX object CMakeFiles/cpp_convnet.dir/cpp_convnet.cpp.o
[100%] Linking CXX executable cpp_convnet
[100%] Built target cpp_convnet

```

Chapter 14: PyTorch on Mobile Devices



The screenshot shows the Android Studio interface with the project 'MasteringPyTorchV2MNISTApp' open. The 'MainActivity.java' file is the active editor. The code implements a neural network to process images and log raw scores.

```
private void processImage(Bitmap bitmap) {
    // Resize the input image to 28x28 pixels
    Bitmap resizedBitmap = Bitmap.createScaledBitmap(bitmap, dstWidth: 28, dstHeight: 28, filter: true);

    ImageView imageView = findViewById(R.id.image);
    imageView.setImageBitmap(resizedBitmap);

    final float[] mean = {0.1302f, 0.1302f, 0.1302f};
    final float[] std = {0.3069f, 0.3069f, 0.3069f};
    final Tensor inputTensor = TensorImageUtils.bitmapToFloat32Tensor(resizedBitmap, mean, std,
        MemoryFormat.CHANNELS_LAST);

    final Tensor outputTensor = module.forward(Value.from(inputTensor)).toTensor();
    final float[] scores = outputTensor.getDataAsFloatArray();

    // Log the raw scores
    Log.d("Raw Scores", msg: "Scores:");
    for (int i = 0; i < scores.length; i++) {
        Log.d("Raw Scores", msg: "Score[" + i + "]": " + scores[i]);
    }

    float maxScore = -Float.MAX_VALUE;
    int maxScoreIdx = -1;
    for (int i = 0; i < scores.length; i++) {
        if (scores[i] > maxScore) {
            maxScore = scores[i];
            maxScoreIdx = i;
        }
    }
}
```

7:09



Camera permission

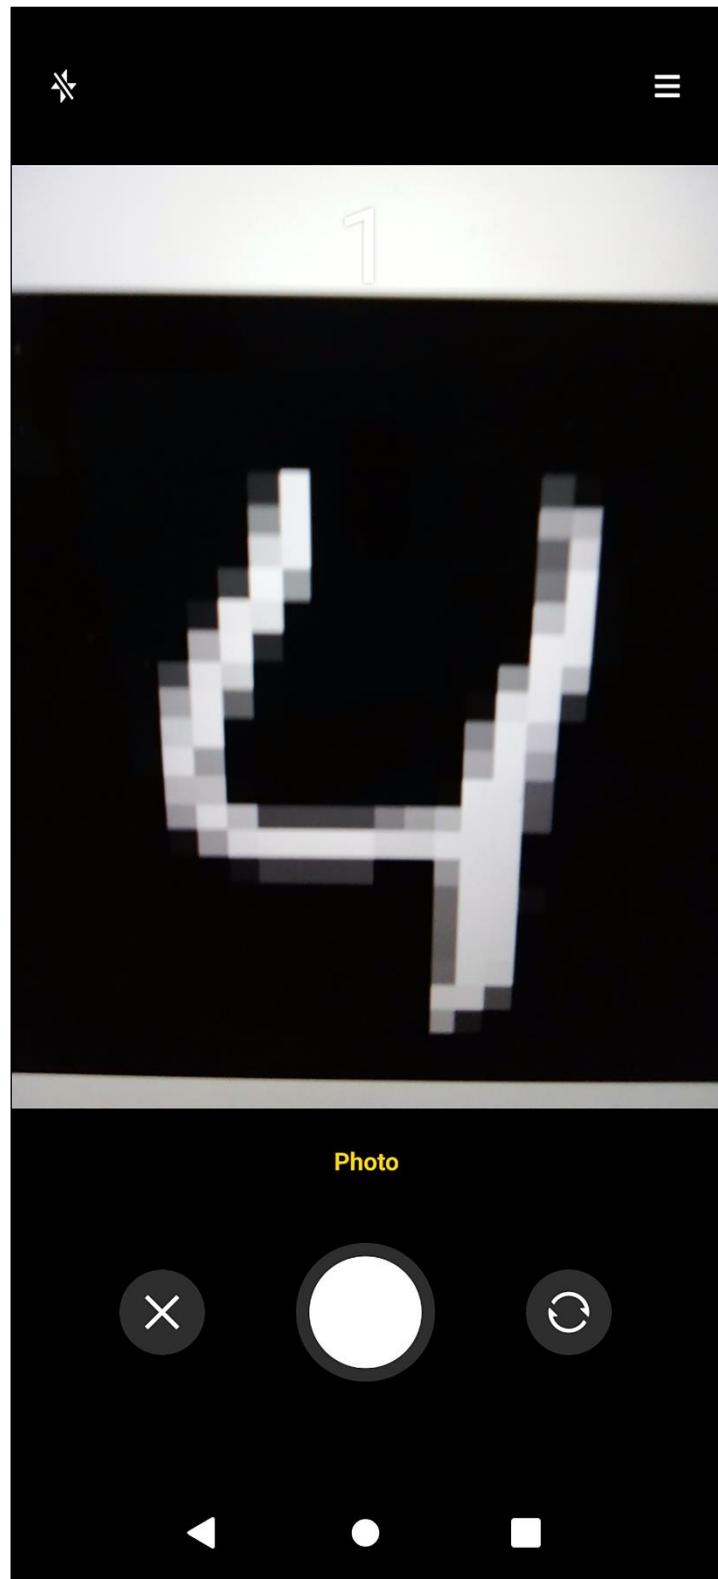


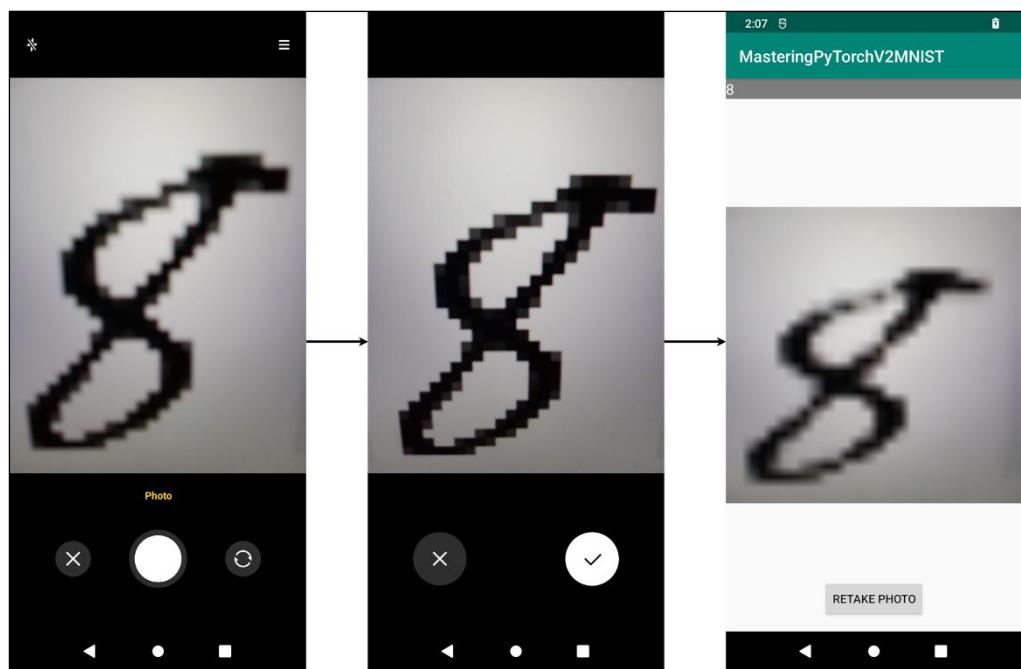
CAMERA ACCESS FOR THIS APP

- Allow only while using the app
- Ask every time
- Deny

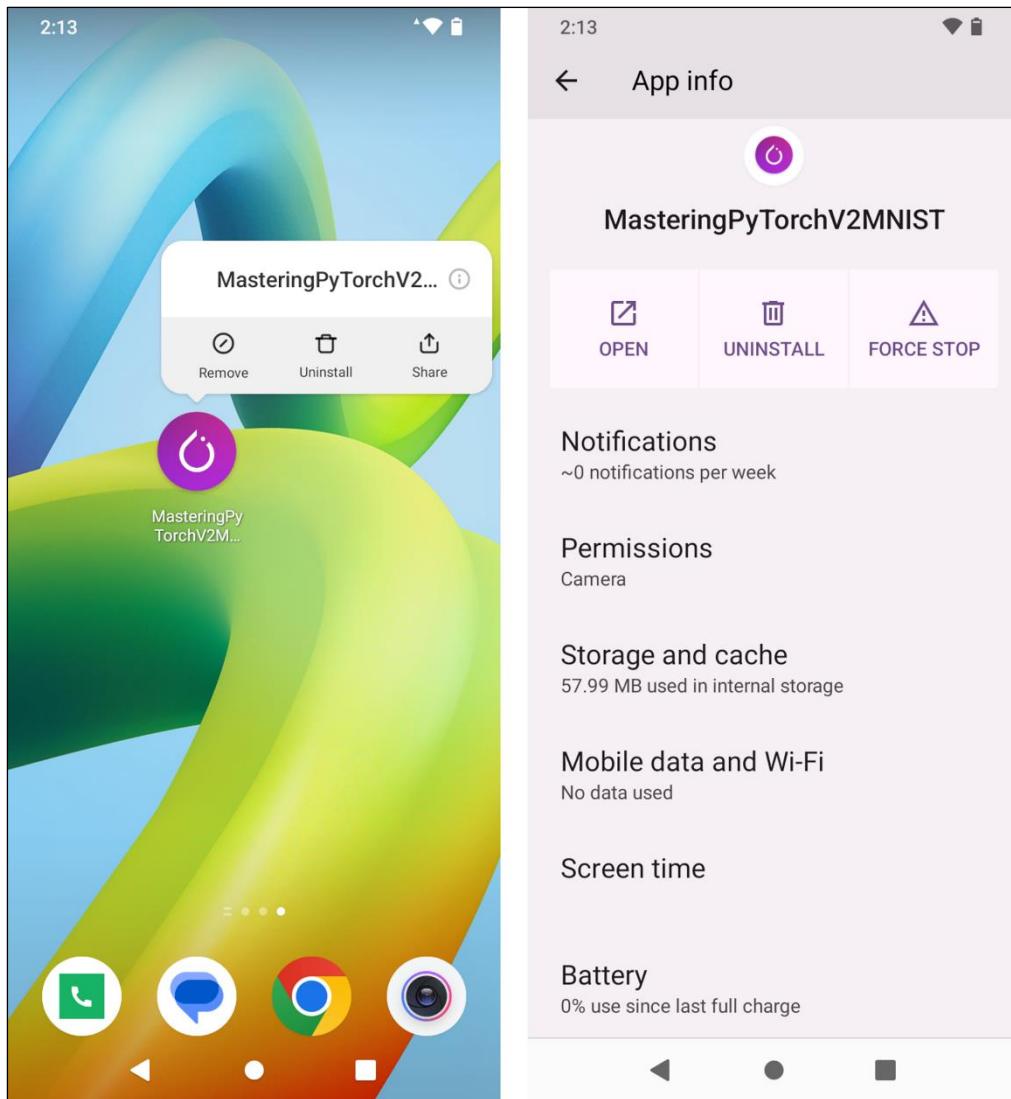
[See all apps with this permission](#)

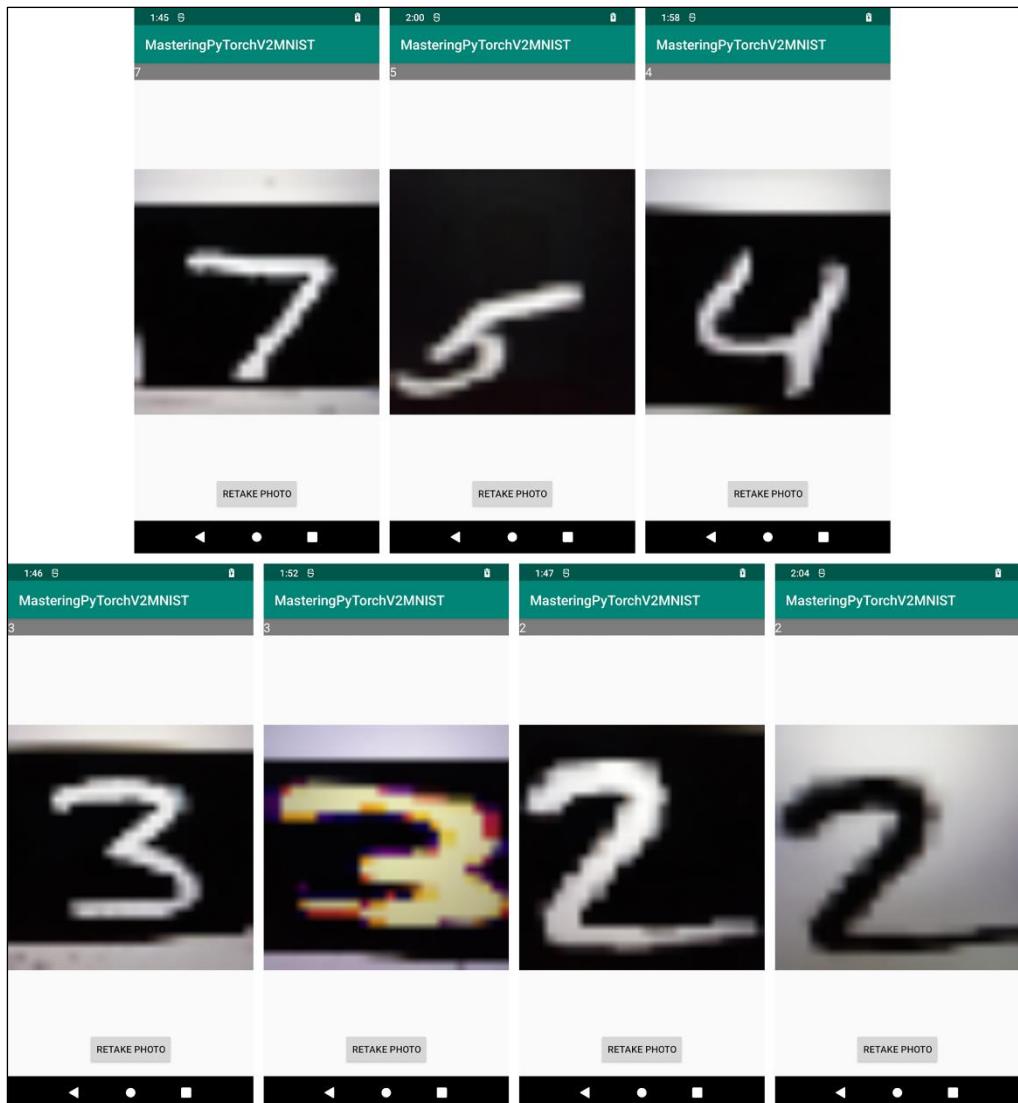


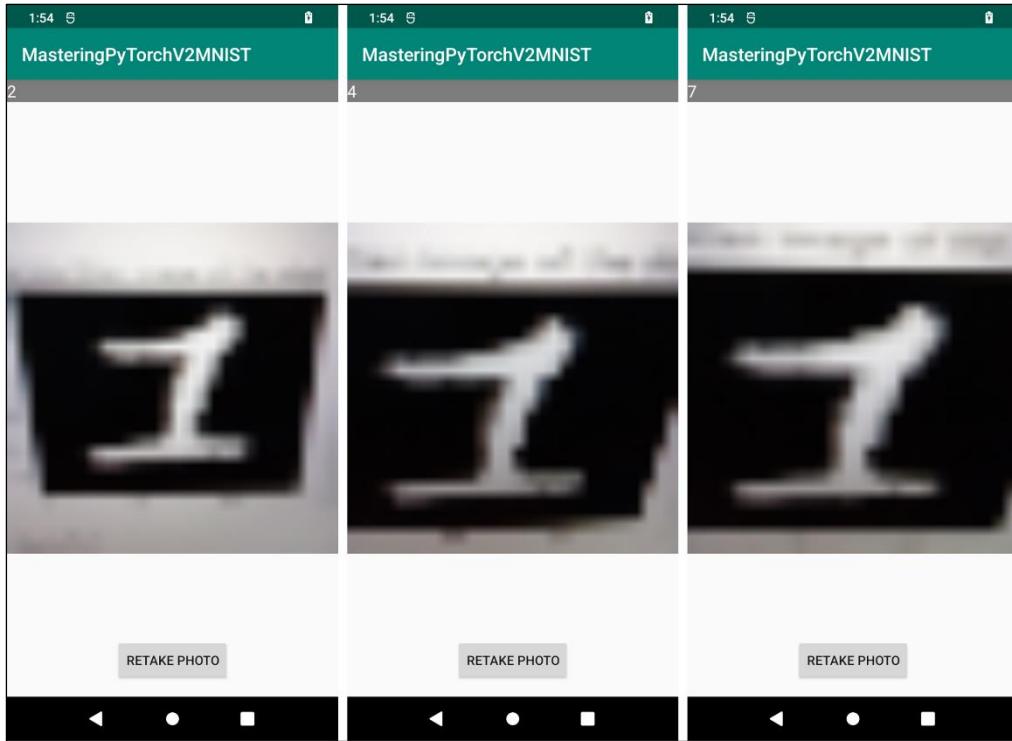












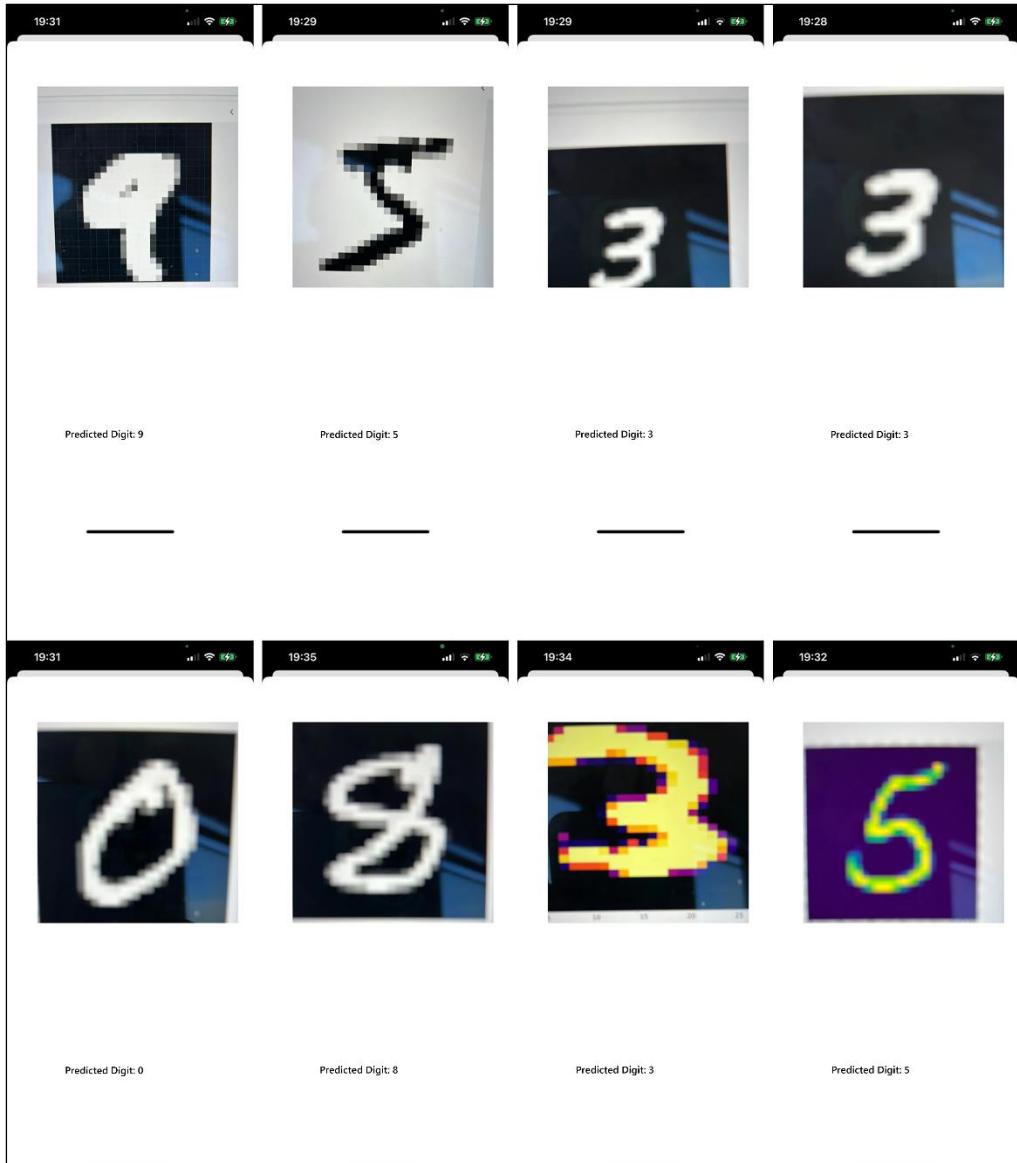
My Mac (Designed for iPhone) Build Succeeded | 23/08/2023 at 23:07

PreviewViewController.swift

```
1 // MARK: - PreviewViewController
2
3 class PreviewViewController: UIViewController {
4     private lazy var labels: [String] = {
5         fatalError("Can't find the text file!")
6     }
7
8     override func viewDidLoad() {
9         super.viewDidLoad()
10        imageView.image = capturedImage
11        guard let resizedImage = capturedImage?.resized(to: CGSize(width: 28, height: 28)),
12            var pixelBuffer = resizedImage.grayscaleNormalized() else {
13            return
14        }
15        imageView.image = resizedImage
16
17        guard let outputs = module.predict(image: UnsafeMutableRawPointer(&pixelBuffer))
18            else {
19                return
20            }
21
22        print("Raw Predictions: \(outputs)") // Print the raw predictions array
23        if let maxIndex = outputs.indices.max(by: { outputs[$0].floatValue < outputs[$1].floatValue }) {
24            let predictedDigit = maxIndex // This is the predicted digit
25            print("Predicted Digit: \(predictedDigit)")
26            resultView.text = "Predicted Digit: \(predictedDigit)"
27        }
28    }
29
30    // MARK: - Properties
31    private let module: MLModel
32    private let resultView: UILabel
33    private let capturedImage: UIImage
34
35    // MARK: - View
36    private let imageView: UIImageView
37
38    // MARK: - Lifecycle
39    private var timer: Timer?
40
41    // MARK: - Initializers
42    init(capturedImage: UIImage, module: MLModel) {
43        self.capturedImage = capturedImage
44        self.module = module
45        super.init(nibName: nil, bundle: nil)
46    }
47
48    required init?(coder aDecoder: NSCoder) {
49        fatalError("init(coder:) has not been implemented")
50    }
51
52    // MARK: - ViewDidLoad
53    override func viewDidLoad() {
54        super.viewDidLoad()
55        setupView()
56    }
57
58    // MARK: - ViewDidLayoutSubviews
59    override func viewDidLayoutSubviews() {
60        super.viewDidLayoutSubviews()
61        resultView.frame = CGRect(x: 0, y: 0, width: 100, height: 100)
62    }
63
64    // MARK: - Setup View
65    private func setupView() {
66        resultView = UILabel()
67        resultView.translatesAutoresizingMaskIntoConstraints = false
68        resultView.textAlignment = .center
69        resultView.textColor = .black
70        resultView.font = UIFont.systemFont(ofSize: 16)
71
72        imageView = UIImageView()
73        imageView.translatesAutoresizingMaskIntoConstraints = false
74        imageView.contentMode = .scaleAspectFit
75
76        let stackView = UIStackView()
77        stackView.translatesAutoresizingMaskIntoConstraints = false
78        stackView.axis = .vertical
79        stackView.spacing = 10
80
81        stackView.addArrangedSubview(imageView)
82        stackView.addArrangedSubview(resultView)
83
84        view.addSubview(stackView)
85
86        NSLayoutConstraint.activate([
87            stackView.centerXAnchor.constraint(equalTo: view.centerXAnchor),
88            stackView.centerYAnchor.constraint(equalTo: view.centerYAnchor)
89        ])
90    }
91
92    // MARK: - Capture
93    private func captureImage() {
94        guard let image = cameraOutputImage else {
95            return
96        }
97
98        capturedImage = image
99
100       let output = module.predict(image: UnsafeMutableRawPointer(&image))
101      let maxIndex = output.indices.max(by: { output[$0].floatValue < output[$1].floatValue })
102      let predictedDigit = maxIndex // This is the predicted digit
103      print("Predicted Digit: \(predictedDigit)")
104      resultView.text = "Predicted Digit: \(predictedDigit)"
105
106      timer?.invalidate()
107      timer = Timer.scheduledTimer(
108          timeInterval: 0.5,
109          target: self,
110          selector: #selector(captureImage),
111          userInfo: nil,
112          repeats: true
113      )
114
115      imageView.image = capturedImage
116  }
```

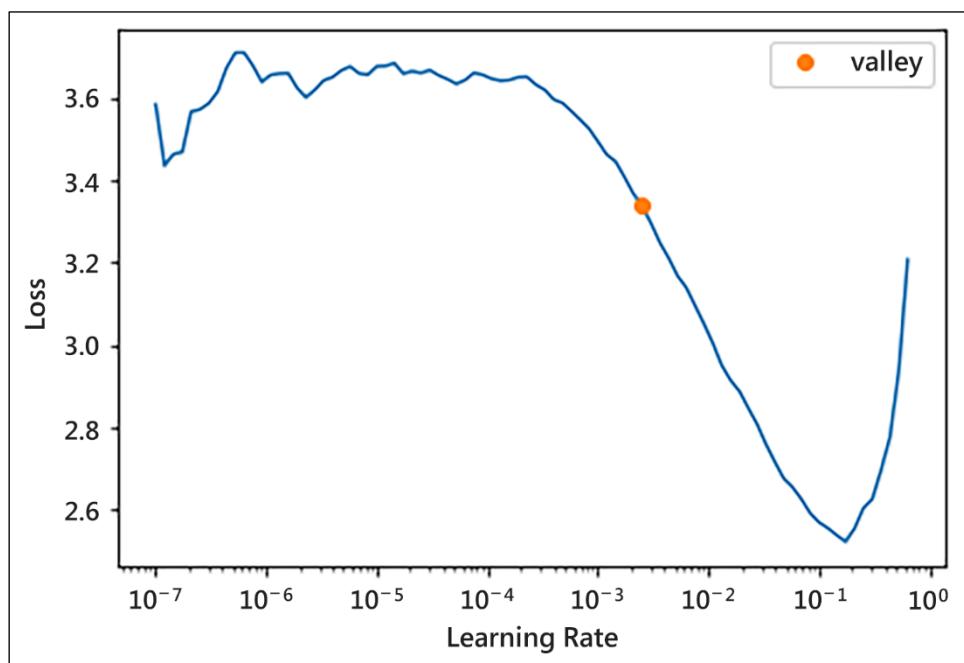
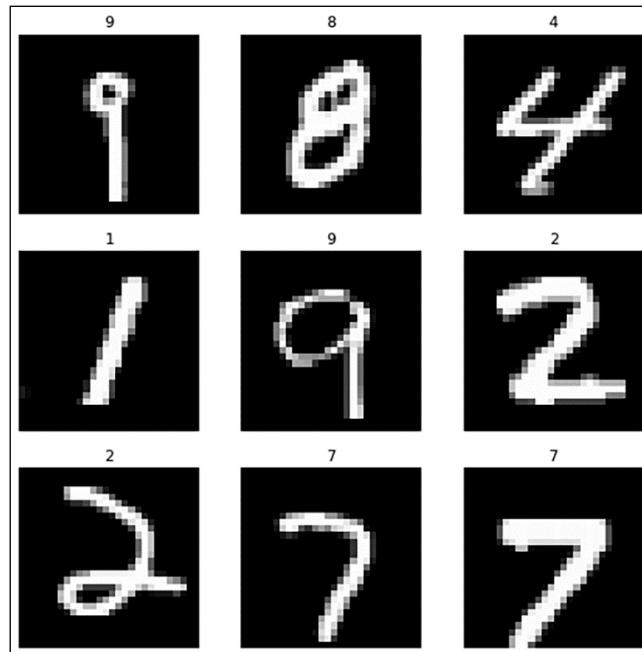
The screenshot shows the Xcode interface with the 'Info' tab selected. The table lists various key-value pairs for the application's Info.plist file. One specific entry, 'Privacy - Camera Usage Description', is highlighted with a red border.

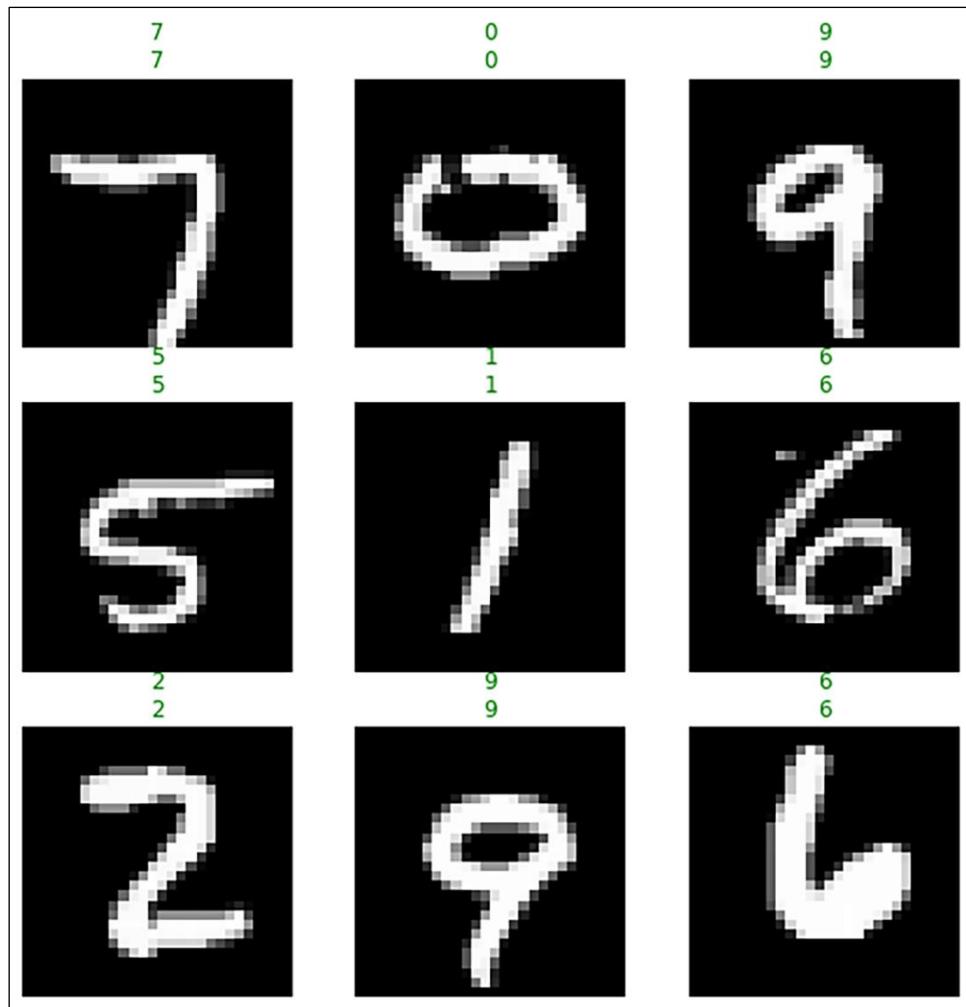
Key	Type	Value
Information Property List	Dictionary	(15 items)
Default localization	String	\$(DEVELOPMENT_LANGUAGE)
Executable file	String	\$(EXECUTABLE_NAME)
Bundle identifier	String	\$(PRODUCT_BUNDLE_IDENTIFIER)
InfoDictionary version	String	6.0
Bundle name	String	\$(PRODUCT_NAME)
Privacy - Camera Usage Description	String	We need access to your camera to capture images for analysis.
Bundle OS Type code	String	\$(PRODUCT_BUNDLE_PACKAGE_TYPE)
Bundle version string (short)	String	1.0
Bundle version	String	1
Application requires iPhone environment	Boolean	YES
Launch screen interface file base name	String	LaunchScreen
Main storyboard file base name	String	Main
Required device capabilities	Array	(1 item)
Supported interface orientations	Array	(1 item)
Supported interface orientations (iPad)	Array	(4 items)





Chapter 15: Rapid Prototyping with PyTorch



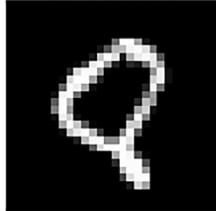


Prediction/Actual/Loss/Probability

3/5 / 11.47 / 1.00



9/8 / 9.58 / 1.00



7/3 / 9.47 / 1.00



9/3 / 8.37 / 1.00



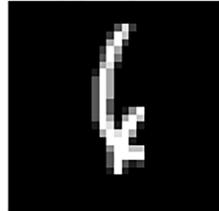
2/7 / 8.17 / 1.00



7/2 / 8.14 / 1.00



6/4 / 7.84 / 0.96

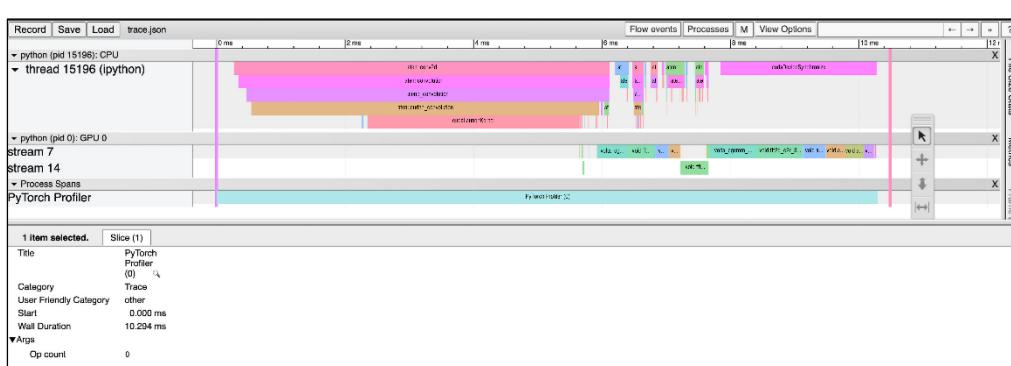
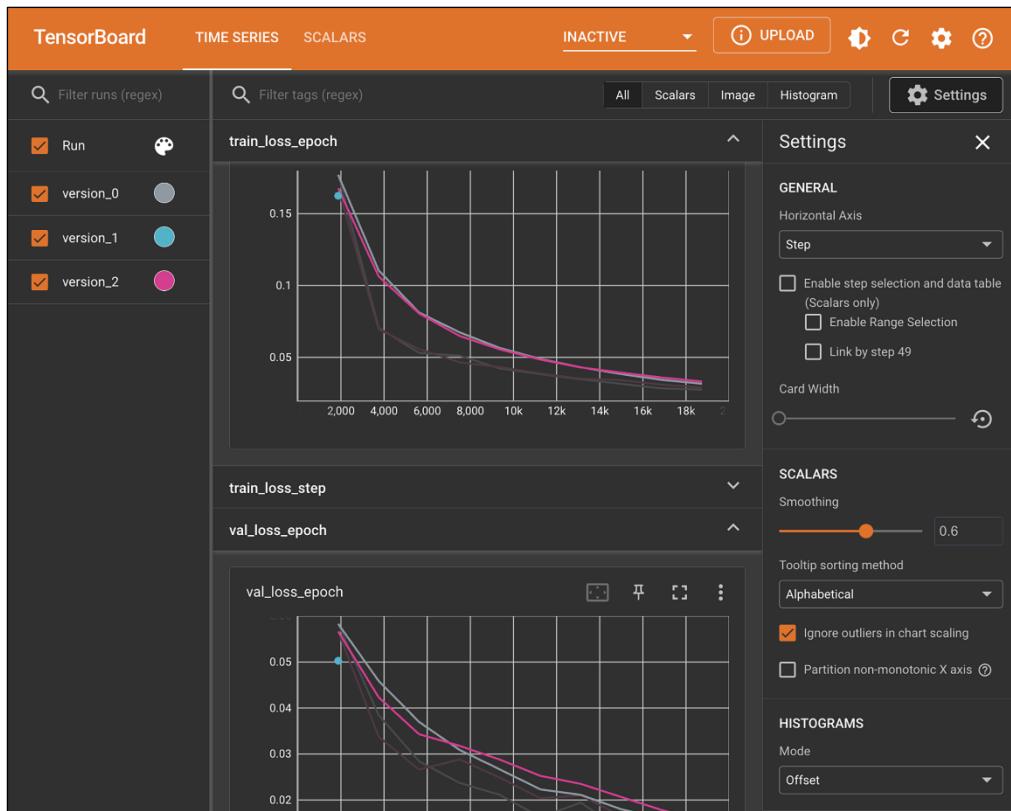


8/2 / 7.63 / 1.00

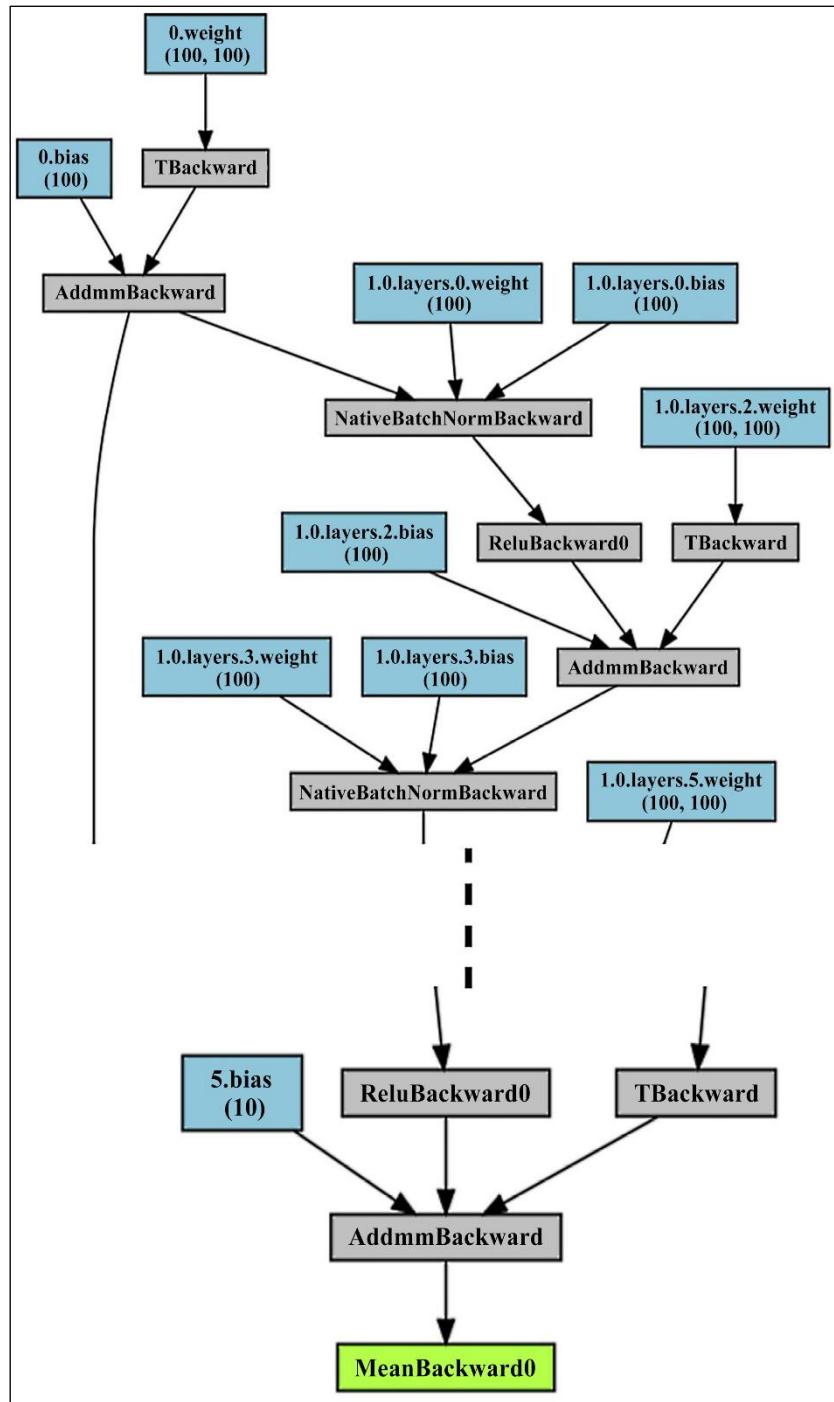


1/7 / 7.62 / 1.00





Chapter 16: PyTorch and AutoML

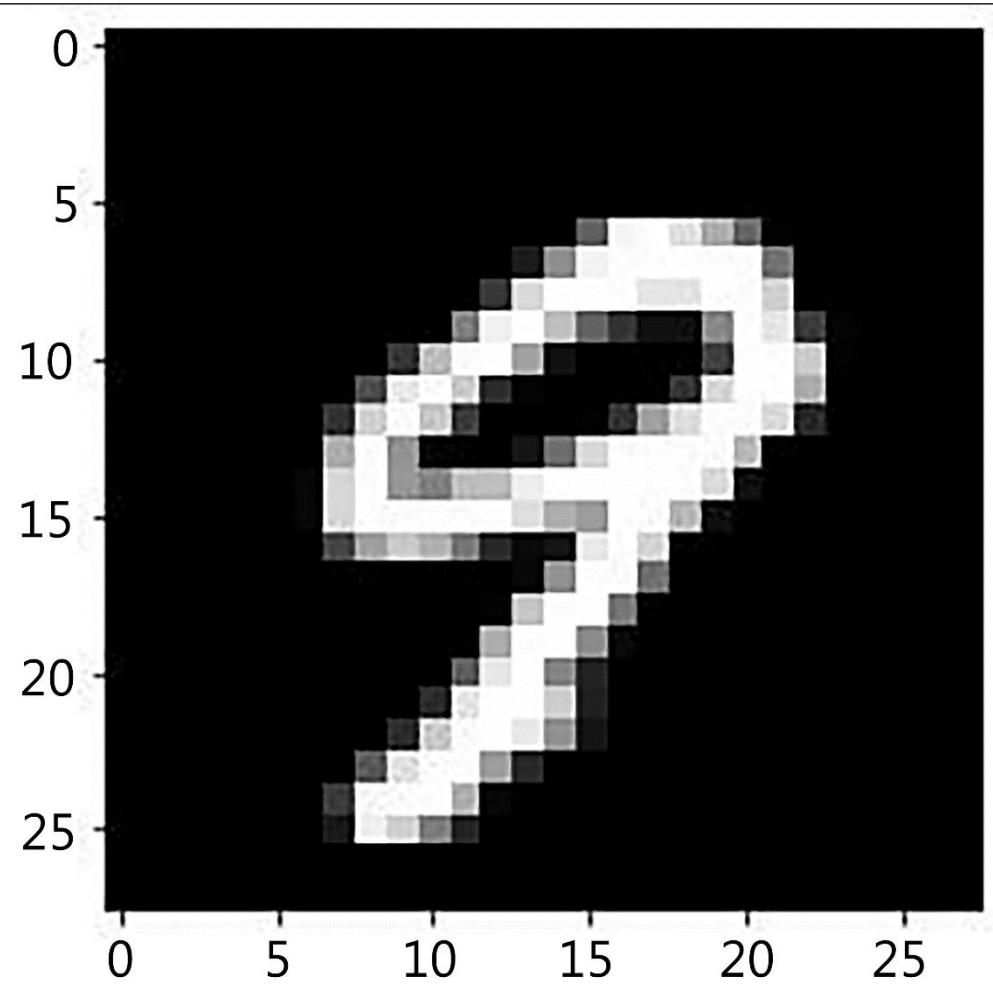


Chapter 17: PyTorch and Explainable AI

```
epoch: 1 [0/60000 (0%)] training loss: 2.324445
epoch: 1 [320/60000 (1%)] training loss: 1.727462
epoch: 1 [640/60000 (1%)] training loss: 1.428922
epoch: 1 [960/60000 (2%)] training loss: 0.717944
epoch: 1 [1280/60000 (2%)] training loss: 0.572199

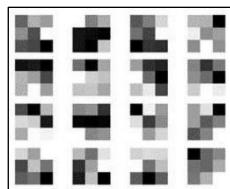
epoch: 19 [58880/60000 (98%)] training loss: 0.016509
epoch: 19 [59200/60000 (99%)] training loss: 0.118218
epoch: 19 [59520/60000 (99%)] training loss: 0.000097
epoch: 19 [59840/60000 (100%)] training loss: 0.000271

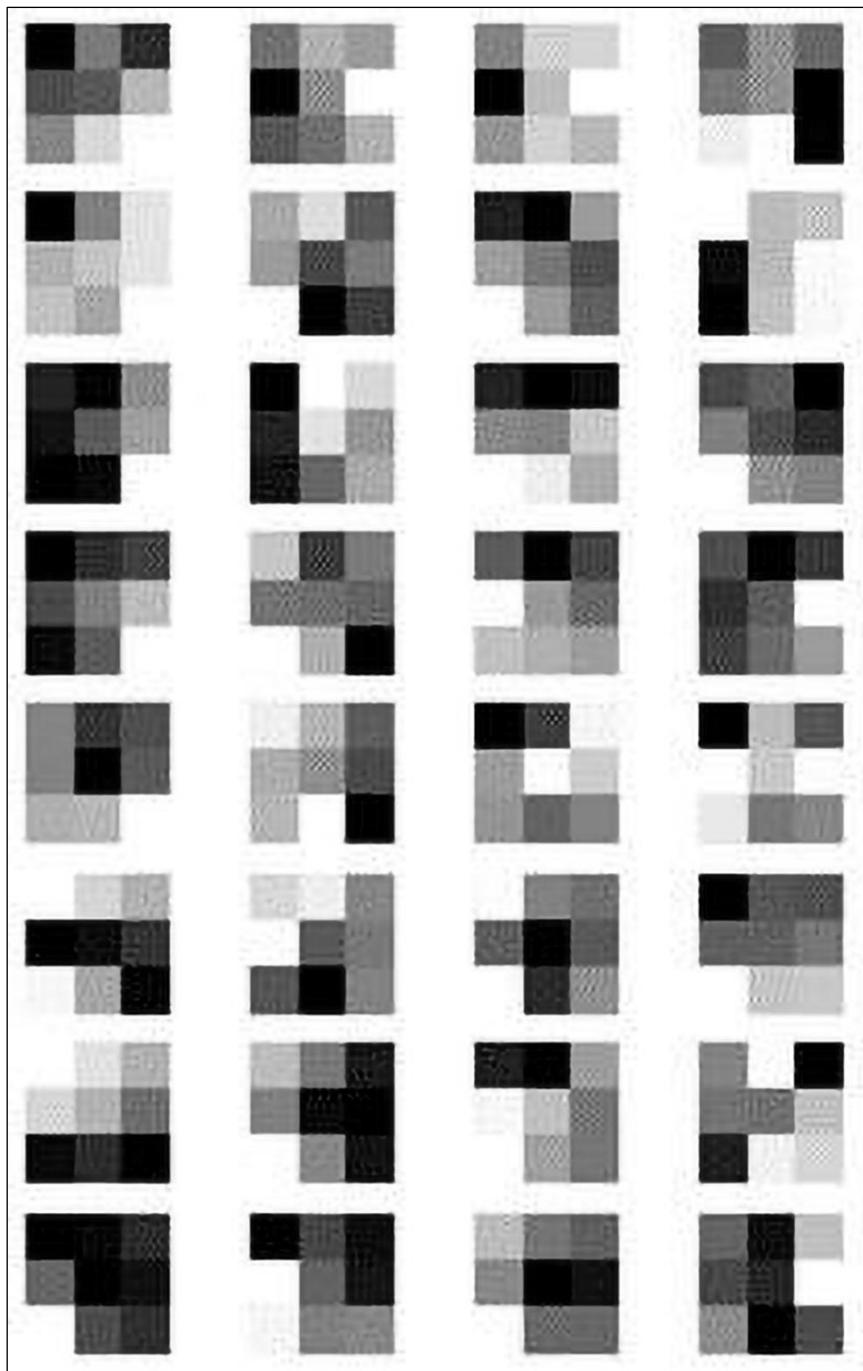
Test dataset: Overall Loss: 0.0387, Overall Accuracy: 9910/10000 (99%)
```



```
Model prediction is : 9  
Ground truth is : 9
```

```
[Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1)),  
 Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1)),  
 Dropout2d(p=0.1, inplace=False),  
 Dropout2d(p=0.25, inplace=False),  
 Linear(in_features=4608, out_features=64, bias=True),  
 Linear(in_features=64, out_features=10, bias=True) ]
```





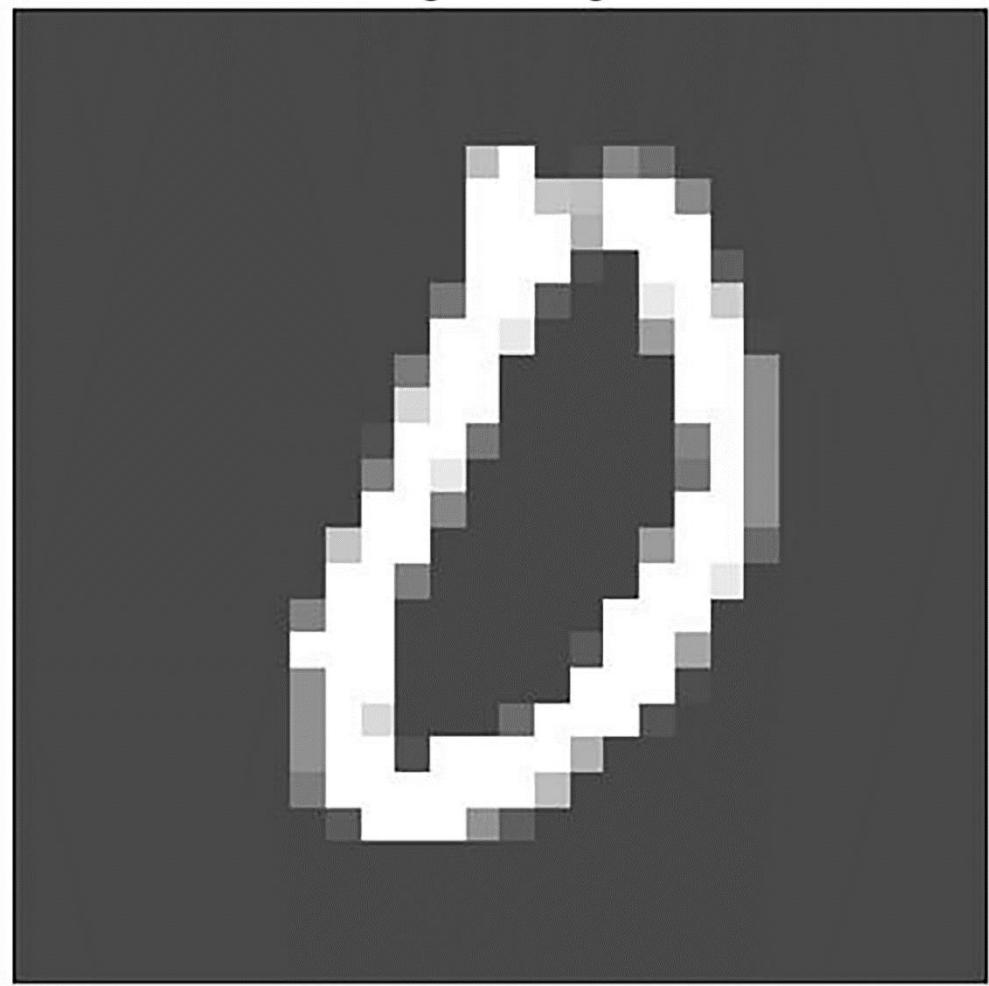
```
torch.Size([16, 26, 26])
```



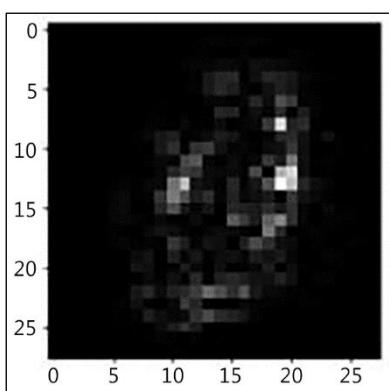
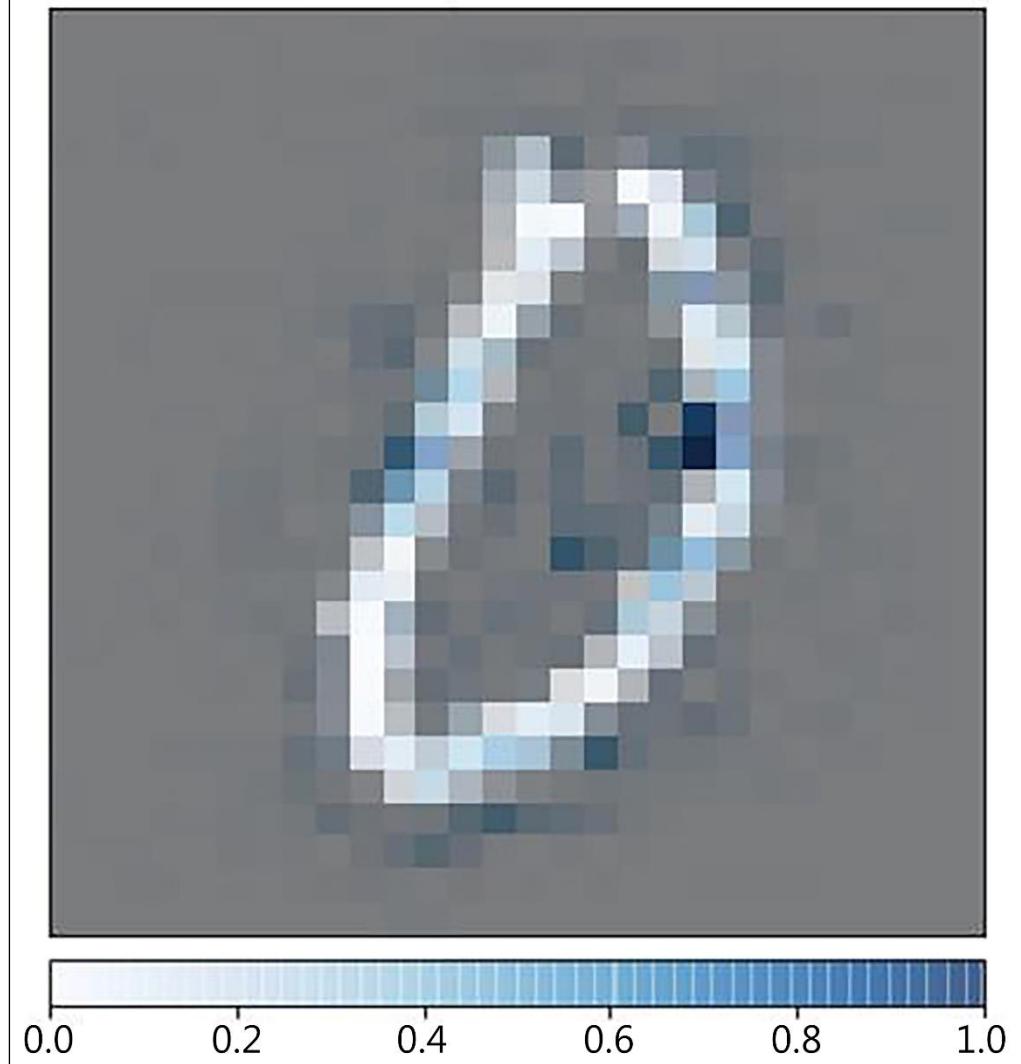
```
torch.Size([32, 24, 24])
```



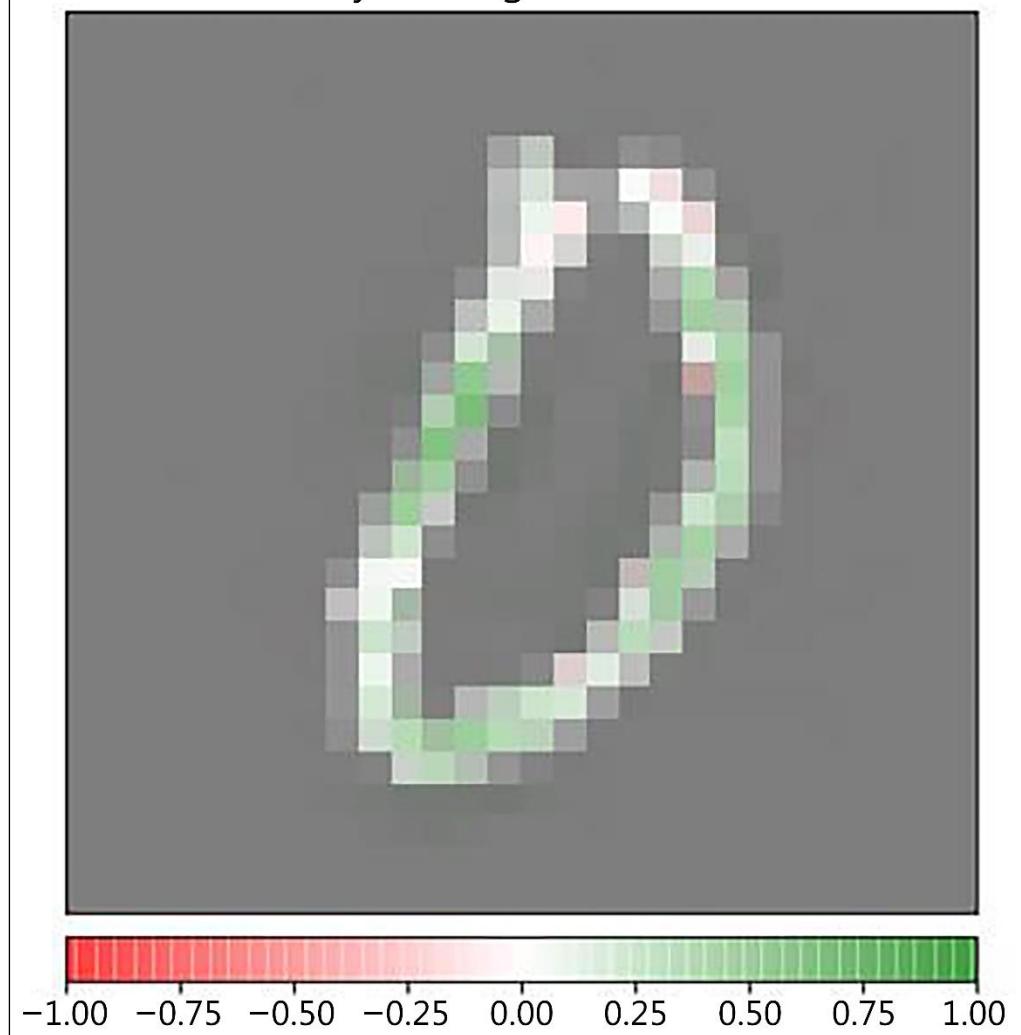
Original Image



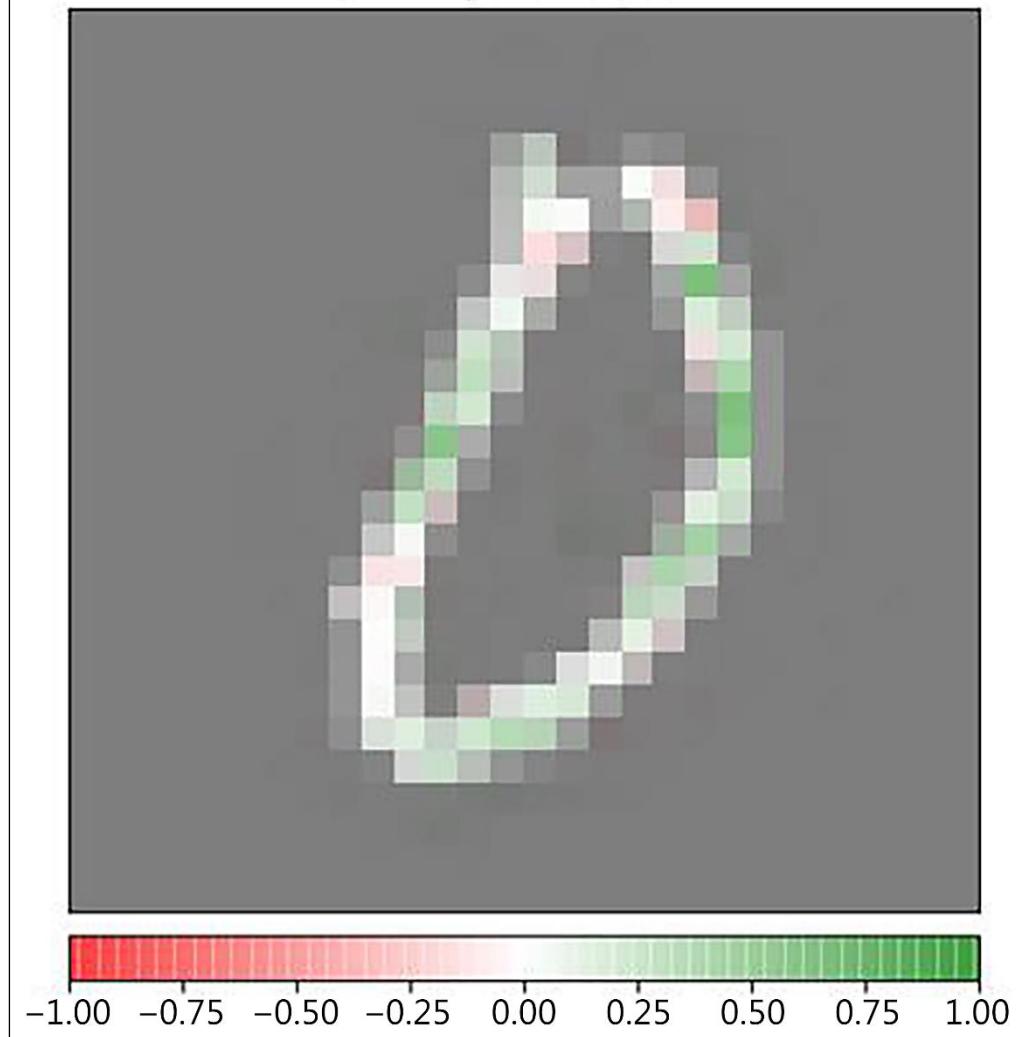
Overlaid Gradients



Overlaid Integrated Gradients



Overlaid DeepLift



Chapter 18: Recommendation Systems with PyTorch

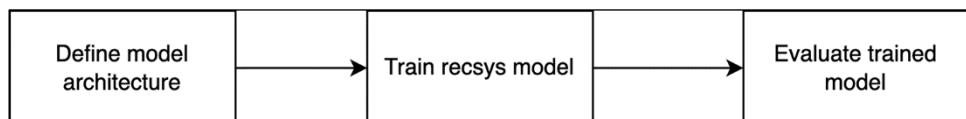
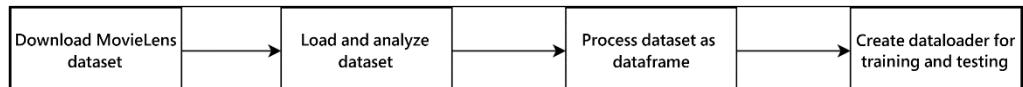
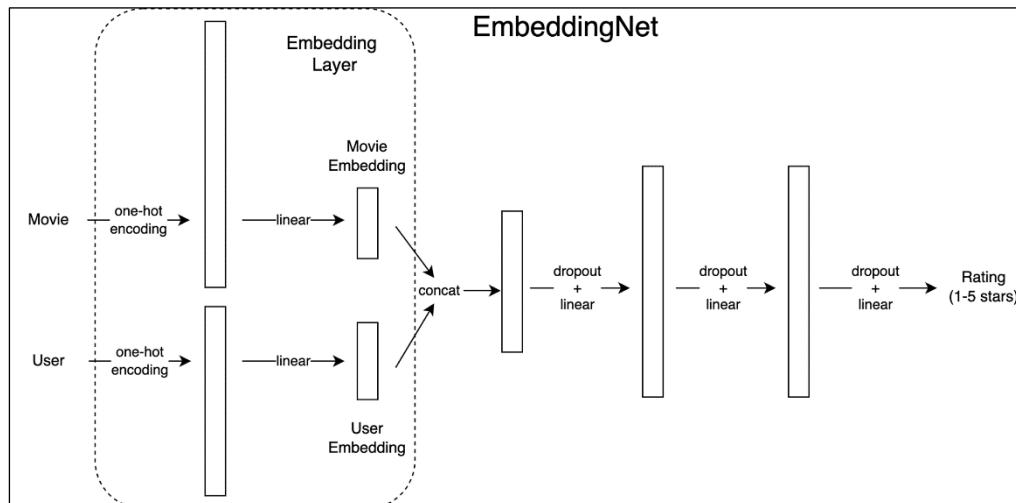
Popular recommendation systems

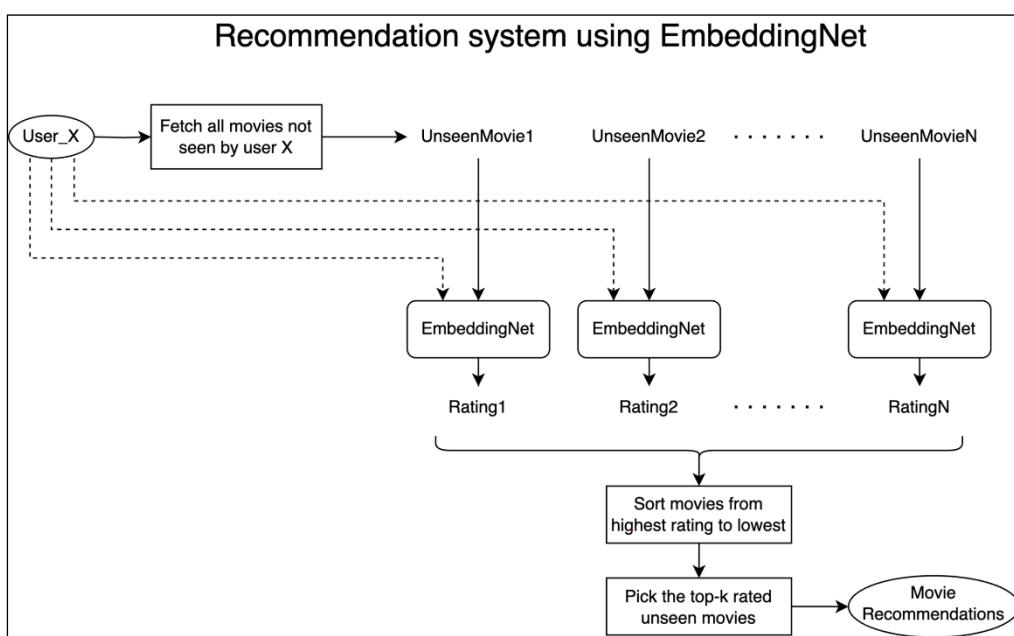
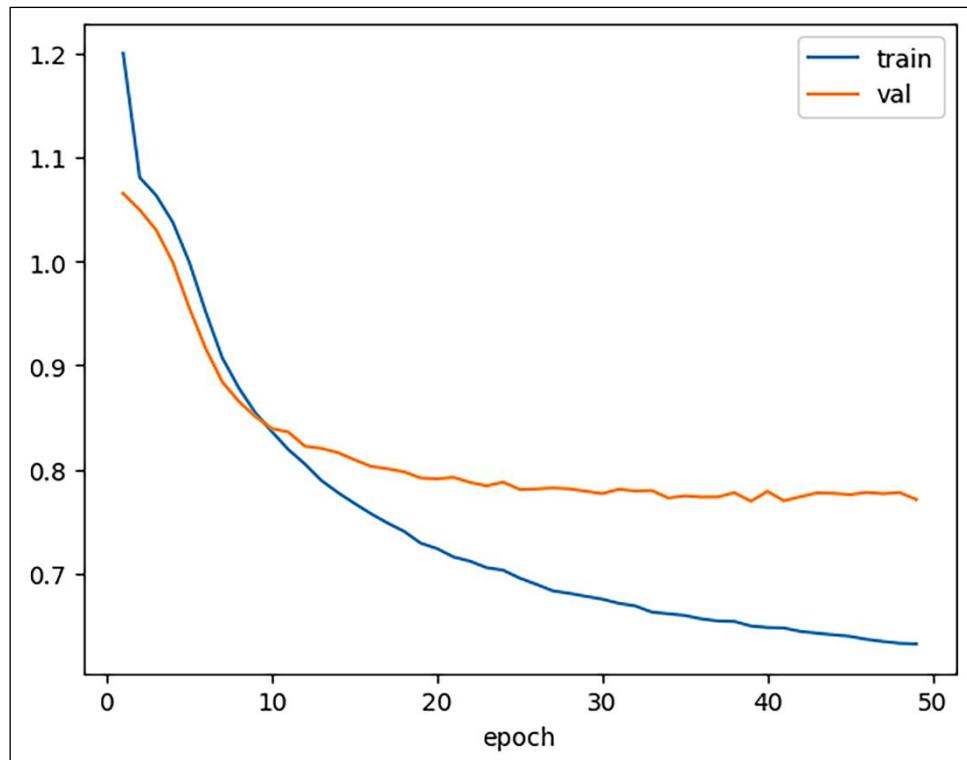
The image consists of four vertically stacked screenshots, each demonstrating a different type of recommendation system:

- Top Screenshot:** A movie recommendation interface. It shows three movie posters: "GRIZZY & THE LEMMINGS", "SING 2", and "Unicorn Academy". Above the first poster, the text "Because you watched Sing" is displayed, with the word "Sing" highlighted by a yellow oval.
- Second Screenshot:** A podcast recommendation interface. It shows four podcast covers: "Who Killed JFK?", "Political Beatdown", "The Obi One Podcast", and "The Rest Is Football". Above the first cover, the text "Shows that you might like" is displayed, with "Shows that you might like" highlighted by a yellow oval.
- Third Screenshot:** A group recommendation interface. It shows a grid of group covers: "Python Developers: Machine Learning,..", "DevOps", "Data Structures and Algorithms", and "GeeksforGeeks". Above the first group, the text "Groups you may be interested in" is displayed, with "Groups you may be interested in" highlighted by a blue oval.
- Bottom Screenshot:** A book recommendation interface. It shows four book covers: "Generative Deep Learning", "Deep Learning with PyTorch Step-by-Step: A Beginner's Guide: Volume III: Sequences ...", "Mastering PyTorch: Build powerful neural network architectures using advanced PyTorch 1.x...", and "Make Your Own Neural Network". Above the first book, the text "Customers who read this book also read" is displayed, with "Customers who read this book also read" highlighted by a blue oval.

Each screenshot is framed by a red border. In the top-right corner of the third and bottom screenshots, there is a navigation bar with buttons for "All", "From your search", "From Packt", and "Related".

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6	Movie 7	Movie 8	New Movie
User 1	★★★★★	?	?	☆☆	★★★★	★★★★	★★★★	★★★★★	?
User 2	?	☆☆☆	☆☆☆☆	☆	?	☆☆☆	?	?	?
User 3	★★★★★	?	☆☆☆	?	★★★★★	?	☆☆☆☆	?	?
User 4	★★★★★	☆☆☆	?	☆☆	?	?	☆☆☆☆	?	?
User 5	?	☆☆☆☆	☆☆	☆☆☆	?	☆☆☆	☆☆	☆☆☆☆☆	?
New User	?	?	?	?	?	?	?	?	?





Chapter 19: PyTorch and Hugging Face

The screenshot shows the Hugging Face Documentation homepage. At the top, there is a navigation bar with links for Models, Datasets, Spaces, Posts, Docs, Pricing, Log In, and Sign Up. Below the navigation bar, there is a search bar labeled "Search across all docs". The main content area is titled "Documentations" and contains a grid of numbered items (1-7) each describing a different tool or service:

- 1. Transformers**: State-of-the-art ML for Pytorch, TensorFlow, and JAX.
- 2. Diffusers**: State-of-the-art diffusion models for image and audio generation in PyTorch.
- 3. Hub**: Host Git-based models, datasets and Spaces on the Hugging Face Hub.
- 4. Datasets**: Access and share datasets for computer vision, audio, and NLP tasks.
- 5. Hub Python Library**: Client library for the HF Hub: manage repositories from your Python runtime.
 - Inference API (serverless)**: Experiment with over 200k models easily using the serverless tier of Inference Endpoints.
 - Inference Endpoints (dedicated)**: Easily deploy models to production on dedicated, fully managed infrastructure.
- 6. Accelerate**: Easily train and use PyTorch models with multi-GPU, TPU, mixed-precision.
 - Tokenizers**: Fast tokenizers, optimized for both research and production.
 - Dataset viewer**
- 7. Optimum**: Fast training and inference of HF Transformers with easy to use hardware optimization tools.
 - Evaluate**: Evaluate and report model performance easier and more standardized.
 - TRL**
- PEFT**: Parameter efficient finetuning methods for large models.
- AWS Trainium & Inferentia**: Train and Deploy Transformers & Diffusers with AWS Trainium and AWS Inferentia via Optimum.
- Tasks**: All things about ML tasks: demos, use cases, models, datasets, and more!
- Amazon SageMaker**

The screenshot shows the Hugging Face Model Hub. At the top, there's a search bar with the placeholder "Search models, datasets, users...". Below it are navigation tabs for "Models" and "Datasets". On the left, there's a sidebar titled "Other" with a badge "1" containing filters for "text-generation-inference": "AutoTrain Compatible", "Inference Endpoints", "Has a Space", "4-bit precision", "8-bit precision", and "Carbon Emissions". The main area shows a list of 56,819 models. Each model entry includes the model name, its type (e.g., "Text Generation"), last update date, size, and popularity metrics (e.g., 1.57k stars).

- gpt2**: Text Generation • Updated Jun 30, 2023 • 18.1M • 1.57k
- distilgpt2**: Text Generation • Updated Apr 29, 2023 • 12.2M • 300
- t5-small**: Translation • Updated Jun 30, 2023 • 3.03M • 185
- davidkim205/komt-mistral-7b-v1**: Text Generation • Updated Oct 24, 2023 • 2.1M • 13
- t5-base**: Translation • Updated Apr 6, 2023 • 1.87M • 391
- Rostlab/prot_t5_xl_uniref50**: Text2Text Generation • Updated Jan 31, 2023 • 1.7M • 31
- google/flan-t5-base**: Text2Text Generation • Updated Jul 17, 2023 • 1.55M • 489

The screenshot shows the Transformers documentation for the T5 model. On the left, there's a sidebar with a dropdown menu for "Transformers" and a search bar. The "T5" model is selected. The main content area discusses the T5Tokenizer and provides a code example for generating text from input sentences.

tokens) and the output sequence is formed as a concatenation of the same sentinel tokens and the *real* masked tokens. Each sentinel token represents a unique mask token for this sentence and should start with <extra_id_0>, <extra_id_1>, ... up to <extra_id_99>. As a default, 100 sentinel tokens are available in `T5Tokenizer`.

For instance, the sentence “The cute dog walks in the park” with the masks put on “cute dog” and “the” should be processed as follows:

```
>>> from transformers import T5Tokenizer, T5ForConditionalGeneration
>>> tokenizer = T5Tokenizer.from_pretrained("t5-small")
>>> model = T5ForConditionalGeneration.from_pretrained("t5-small")

>>> input_ids = tokenizer("The <extra_id_0> walks in <extra_id_1> park", return_tensors="pt")
>>> labels = tokenizer("<extra_id_0> cute dog <extra_id_1> the <extra_id_2>", return_tensors="pt")

>>> # the forward function automatically creates the correct decoder_input_ids
>>> loss = model(input_ids=input_ids, labels=labels).loss
>>> loss.item()
3.7837
```

Hugging Face Models Datasets Spaces Docs Solutions Pricing

Tasks Sizes Sub-tasks Languages Licenses Other Filter Tasks by name Reset Tasks

Datasets 2,162 Filter by name new Full-text search ↑ Sort: Most downloads

glue Viewer - Updated Jun 1, 2023 - 980k - 265	super_glue Viewer - Updated Apr 5, 2023 - 226k - 122
imdb Viewer - Updated Apr 5, 2023 - 204k - 136	ceval/ceval-exam Viewer - Updated Aug 31, 2023 - 92.3k - 171
tasksource/bigbench Viewer - Updated May 11, 2023 - 86.7k - 43	sst2 Viewer - Updated May 2, 2023 - 54.2k - 37
rotten_tomatoes Viewer - Updated Apr 5, 2023 - 51.4k - 32	shunk031/JGLUE Viewer - Updated Sep 26, 2023 - 50.3k - 33
ag_news Viewer - Updated Apr 5, 2023 - 47.8k - 88	blimp Viewer - Updated Apr 5, 2023 - 40.9k - 30
tweet_eval Viewer - Updated Jun 1, 2023 - 38.9k - 87	lex_glue Viewer - Updated Jun 1, 2023 - 35.1k - 36
Open-Orca/OpenOrca Viewer - Updated Oct 21, 2023 - 34.2k - 971	THUDM/LongBench Viewer - Updated Aug 29, 2023 - 32.7k - 45

Hugging Face Models Datasets

Datasets: rotten_tomatoes like 32

Tasks: Text Classification Sub-tasks: sentiment-classification Languages: English Multilinguality: monolingual Size Categories: 1K< n <10K Lang

Source Datasets: original License: unknown

Dataset card Files and versions Community 3

Dataset Viewer Auto-converted to Parquet API Go to dataset viewer

Split

train (8.53k rows)

text string · **lengths**

label class label

label	Count
1 pos	8.53k
2 classes	~1.5k

the rock is destined to be the 21st century's new "conan" and that he's going to make a splash even greater than arnold schwarzenegger , jean-claud van damme or steven segal .

the gorgeously elaborate continuation of " the lord of the rings " trilogy is so huge that a column of words cannot adequately describe co-writer/director peter jackson's expanded vision of...

effective but too-tepid biopic

if you sometimes like to go to the movies to have fun , wasabi is a good place to start .

emerges as something rare , an issue movie that's so honest and keenly observed that it doesn't feel like one .

the film provides some great insight into the neurotic mindset of all comics -- even those who have reached the absolute top of the game .

offers that rare combination of entertainment and education .

< Previous 1 2 3 ... 86 Next >

